Raise Ready

Raise Ready

The Founder's Guide to Financial ModelsThat Close Rounds

Yanni Papoutsi

Raise Ready: The Founder's Guide to Financial Models That Close Rounds

© 2026 Yanni Papoutsi. All rights reserved.

No part of this publication may be reproduced, distributed, or transmitted in any form or by any means without the prior written permission of the author, except in the case of brief quotations embodied in critical reviews.

First Edition, 2026

Contents

Part: Fundraising

001 You Have 93 Seconds. Most Founders Waste 80 of Them.

002 Stop Polishing Your Deck. Investors Are Not Buying the Font.

003 Running Out of Cash During a Fundraise Is More Common Than You Think. Here Is How to Not Die.

004 Raise When You Don't Need To, Sell When You Don't Have To

005 Venture Debt: The Startup Financing Tool Nobody Explains Properly

006 Convertible Notes Are Not Free Money. Here Is What Founders Get Wrong.

007 Not All Money Is Equal. How to Pick the Right VC for Your Stage.

008 What Creandum, Profounders, and B2Ventures Actually Look For in Your Model

009 The 30 Investor Diligence Questions You Will Get. Here Are the Answers.

010 Seed Round vs. Series A: How Your Financial Model Must Change Between Stages

011 How to Run a Competitive Fundraising Process Without Burning Every Bridge in Town

012 Term Sheet Red Flags: What to Watch For Before You Sign Away Control of Your Company

013 Bridge Rounds: When They Save Your Company and When They Bury It

014 How to Negotiate Valuation Without Killing the Deal

015 What Happens When Investors Challenge Your Numbers (And How to Win the Conversation)

016 How to Build a Data Room That Closes Rounds Instead of Stalling Them

017 The Fundraising Timeline: What Actually Takes How Long (And Why It Always Takes Longer Than You Think)

018 Post-Raise: The First 90 Days That Determine Whether the Money Actually Works

019 How to Write an Investor Update That Keeps Your Round Warm

020 The Cap Table: What Every Founder Must Understand Before Taking a Single Dollar

Part: Financial Modeling

021 How to Build a Financial Model for Your Startup (2026 Guide)

022 Driver-Based Modeling: Build Revenue From Reality, Not From Hope

023 The Cohort Method for Revenue Forecasting: The Most Accurate Way to Predict Startup Revenue

024 The Assumptions Tab: The Most Important Sheet in Your Entire Financial Model

025 How to Build a Headcount Plan That Investors Trust and Founders Can Actually Use

026 COGS vs. OpEx: Why Getting This Wrong Destroys Your Gross Margin Story

027 The Cash Flow Statement Founders Always Get Wrong

028 How to Build Three Scenarios That Prove You Have Actually Thought About What Could Go Wrong

029 How to Model a Two-Sided Marketplace: Lessons From Building the Indeed Flex Financial Model

030 SaaS vs. Marketplace Financial Models: The Key Differences That Change Everything

031 Multi-Currency Financial Models: How to Run Finance Across UK, US, and Europe

032 How to Model Gross Margin for Different Business Models

033 Sensitivity Analysis: Pre-Building the Questions Investors Will Ask

034 How to Audit Your Financial Model Before Sharing It With Investors

035 Excel vs. Google Sheets for Startup Financial Modeling: The Honest Answer

036 The Balance Sheet Founders Skip (And Why That's a Mistake)

037 How to Model Pre-Revenue Startups

038 Financial Model Red Flags: What Breaks Investor Confidence

039 How to Update Your Model Mid-Fundraise Without Losing Consistency

040 From Model to Board Report: Keeping the Spreadsheet Alive Post-Raise

Part: Unit Economics

041 CAC, LTV, and Payback Period: Calculated Correctly for Once

042 Why Blended CAC Is a Lie (And How to Model It by Channel)

043 Net Revenue Retention: The Number That Predicts Everything

044 How to Model LTV for a Marketplace Business

045 Unit Economics Across Multiple Markets: UK vs US vs UAE

046 The LTV:CAC Benchmarks by Stage and Business Model

047 Gross Margin by Business Model: What's Good, What's Bad, What's a Problem

048 How to Build a Cohort Analysis That Actually Tells You Something

049 The Metrics That Matter at Pre-Seed vs. Seed vs. Series A

050 How to Model Network Effects in a Marketplace

051 Churn Rate: Every Way to Calculate It (And Why the Method You Choose Changes the Story You Tell)

Part: Finance Operations

067 The 10 KPIs Every Founder Should Track (And the 20 They Should Ignore)

068 How to Build a Metrics Dashboard That You Will Actually Use Every Week

069 Leading vs. Lagging Indicators: How to See Problems Three Months Before They Hit Your Revenue

Part: AI in Finance

076 You Do Not Need to Be a Genius With Excel Anymore. AI Does That Now.

Part: War Stories

091 How I Saved a more than 9 figures Exit on a Saturday Morning (While My Girlfriend Waited)

Appendix A: Chapter Exercises

Appendix B: Recommended Reading

About the Author

Introduction

This book started as a problem I kept running into.

Founders were losing deals not because their businesses were weak, but because they could not explain their numbers. Not in the way that closes rounds. They had spreadsheets. They had a deck. They had a story they believed in. What they were missing was the ability to make the financial model do the talking, to have it answer the questions before the investor asked them, to sit in a partner meeting and defend every assumption without checking the screen.

I spent years on the finance side of that table. Through a seed round, a Series A, venture debt, convertible notes, and eventually a nine-figure acquisition, I watched the same patterns repeat. The founders who raised were not always the ones with the best products. They were almost always the ones who understood the math underneath their business well enough to make it legible to the people holding the money.

The founders who did not raise were often closer than they knew. One had a genuinely impressive marketplace with strong cohort retention but could not explain why their LTV calculation was more conservative than a comparable. One had a SaaS product growing 15% month-over-month but presented the financials in a way that made the growth look like an accident rather than a mechanic. One had a realistic shot at a Series A but walked into the partner meeting with a model that had three different versions of the same revenue number, none of them traceable to a source.

These are fixable problems. That is what this book is for.

What This Book Covers

The book is organised into five parts. Part One covers the mechanics of venture capital and the fundraising process itself: how VCs think, what the diligence process actually looks like, which terms to fight for and which ones do not matter, and what happens after you close. Part

Two covers financial modelling from the ground up: how to build a model that an investor can stress-test, how to structure your assumptions so that every number is defensible, and how to avoid the specific errors that destroy credibility in the room. Part Three covers unit economics: the metrics that tell a VC whether your business model actually works. Part Four covers finance operations: the dashboards, KPIs, and analytics infrastructure that makes the model credible once you have closed the round. Part Five covers AI and tooling, and ends with a war story from one of the most stressful Saturday mornings of my career.

The appendices include chapter exercises if you want to test the frameworks against your own numbers, a reading list for going deeper on specific topics, templates for investor updates and financial models, and a glossary of the terms that come up most often in investor conversations.

How to Use This Book

If you are actively fundraising, start with Part One. Read it end to end before your next investor meeting. The VC psychology chapter alone tends to change how founders think about who to pitch and why.

If you are building or rebuilding your financial model, go directly to Part Two. The chapter on the Assumptions tab is the most important in that section. If your model does not have a clean assumptions tab, nothing else will hold up in diligence.

If an investor has specifically pushed back on your unit economics, go to Part Three. The chapters on LTV, CAC, and NRR cover every common calculation error and the benchmarks that matter by stage and business model.

If you just closed a round and are trying to build the financial infrastructure to support your next one, Parts Four and Five are for you.

Each chapter can be read independently if you need a fast answer to a specific question. But the book is designed to be read in order at least once. The concepts in the later parts build on the language and frameworks from the earlier ones.

A Note on the Numbers

Every benchmark, percentage, and ratio in this book comes from a named source or from direct experience across the companies I have worked with and the investor conversations I have been part of. Where I have made a calculation, I have shown the method. Where I have cited a benchmark, I have named where it comes from.

Numbers change. The benchmarks that applied at Seed in 2021 did not apply in 2023. Use these as orientation, not doctrine. The frameworks do not change. The numbers need to be calibrated to your market, your stage, and the current environment.

The goal is not to hand you the right answers. It is to give you the tools to find them, and the language to make them legible to the people sitting across the table from you.

PART

Part One: Fundraising

Raising capital is a skill, not a lottery. This part covers everything that happens between the decision to raise and the wire hitting your account: how venture capital actually works, what investors are looking for in your model and your pitch, how to run a process that creates competitive tension without burning relationships, and how to read a term sheet before you sign one.

000 VC Psychology: Why the Math Forces Decisions That Seem Arbitrary

How VC Fund Economics Actually Work

Every venture capital fund has a legal obligation to return capital to its limited partners (LPs) with a meaningful multiple on top. The industry standard target is 3x net returns over the life of the fund. That means a $100M fund needs to return $300M in cash to its investors before fees and carry.

This is not aspirational. It is contractual. And it shapes everything.

A fund that returns 1.5x is not a success. It is a failure. The GP (general partner, i.e., the VC) earns carried interest, typically 20% of profits above the hurdle, only if they meaningfully beat that bar. Underperformance means no carry, no reputation, and usually no successor fund.

Here is the math that every founder should understand before walking into a pitch meeting:

Fund size: $150M. Target return multiple: 3x. Total capital to return: $450M.

Average ownership at entry (Seed): 10–15%. Dilution by exit: down to 5–8%. Exit value needed to return fund on a single company: $450M / 7% ownership = $6.4B.

That last number is why VCs say no to good companies. A company doing $20M ARR growing 40% year over year is a genuinely impressive business. It is also not a fund returner for a $150M fund. It might be a strong outcome for a smaller fund or an angel, but it does not solve the math problem for a growth-stage VC.

> *Key insight: When a VC passes on your company, they may not be saying it is a bad company. They may be*

> *saying it cannot return their fund at the valuation required for them to invest. These are different problems with different solutions.*

Power Law Returns: Why 10 Companies Do Not Equal One Fund Returner

Venture returns follow a power law distribution, not a normal distribution. In a typical portfolio of 30 companies, one or two drive the vast majority of returns. The rest return capital, underperform, or go to zero.

Cambridge Associates data shows that the top 10% of venture investments generate approximately 80% of returns across the asset class. This is not a quirk. It is the structural reality that defines every decision VCs make.

What this means for founders: VCs are not looking for good companies. They are looking for potential outliers. A company that can credibly grow to $1B in revenue is more investable than a company that will very reliably reach $50M, even if the $50M outcome has a higher probability.

This is counterintuitive until you understand the math. A company with a 70% chance of reaching $50M might be a better expected value calculation in isolation. But it does not solve the power law problem. The VC needs the $1B shots, even if most of them miss.

Ownership Targets: The Number Behind Every Valuation Negotiation

Every VC firm has an ownership target. This is typically 10–20% at the seed stage and 15–25% at Series A. It is not arbitrary. It is derived from the fund math above.

If a VC needs to own 10% of a company at exit to return their fund, and they expect dilution between initial investment and exit, they back-calculate the ownership required at the time of investment. That number determines the maximum valuation they can pay.

This is why valuation negotiations that seem purely about ego are often purely about math. A VC offering a $10M pre-money valuation at a $2M check is targeting 16.7% ownership. If you push to a $15M pre-money, their ownership drops to 11.8%. Their model may require 15%. They pass, not because the company is worse, but because the ownership does not work.

The practical implication: when you receive a term sheet, understand the ownership math before arguing the headline valuation number. Sometimes raising at a lower valuation with a better investor is better than a higher valuation that barely closes. Ownership in a functioning relationship beats valuation in a difficult one.

Fund Size and Stage: Why This Determines Who Should Pitch Whom

Fund size determines stage focus. This is arithmetic, not strategy.

A $25M micro-fund making $250K checks can build a meaningful portfolio across 30 companies and return the fund from multiple smaller exits. A $500M fund making $10M checks cannot. A $50M exit at 8% ownership returns $4M on a $10M investment at 0.4x. That is a write-off, not a portfolio contribution.

This produces a near-complete mismatch between what early-stage founders think they want and what actually helps them.

Many first-time founders target Tier-1 large funds because the brand recognition is appealing. But a $300M growth fund doing a seed deal is almost always a misalignment. They cannot reserve enough capital to follow on, they will not be patient with a long development cycle, and your company is not meaningful to their fund economics at the check size they can deploy. You become an afterthought.

Target investors by fund size as much as by reputation. A $30M founder-focused seed fund that leads your round and reserves $500K for follow-on is a better institutional partner than a brand-name firm that writes a small check and moves on.

Why Good Companies Get Rejected: The Real Reasons

Founders who receive polite rejections often assume they failed somewhere. The pitch was unclear. The numbers were not compelling. The team was not strong enough. Sometimes this is true. But a significant proportion of rejections have nothing to do with the company.

Timing: The fund just made a competing investment in the same category. Most funds have portfolio conflict policies. If they backed a company solving a similar problem six months ago, they will not invest in a direct competitor regardless of quality.

Fund cycle: Early-stage funds make most investments in the first two-thirds of the fund life, then reserve capital for follow-ons. A fund three years into a four-year deployment period is nearly done making new bets. You may be pitching someone who is not actually deploying.

Ownership math: As above. The valuation required to close your round does not produce the ownership percentage their fund model requires.

Conviction gap: Some VCs operate on full conviction or pass. They like the company but cannot build enough conviction to lead. Rather than say this, they pass with compliments. This is not dishonesty; it is how most investment decisions actually work.

Market timing mismatch: They believe in the problem but think it is two years early or two years late for their current fund. This is a real consideration. A fund making ten-year bets does care when the category becomes mainstream.

> *Key insight: The most useful question to ask after a rejection is not "what did we do wrong" but "was this a fit problem or a quality problem?" Fit problems are solvable by targeting different investors. Quality problems require actually fixing the business.*

What This Changes About How You Pitch

If you understand the fund math, you pitch differently.

You lead with market size not as a compliance slide but as a fund returner argument. You are not trying to prove the market exists. You are proving the market is large enough to justify the bet.

You reference the investor's portfolio and fund size when you can. Showing that you understand their ownership targets signals financial sophistication. Most founders never do this.

You do not argue against ownership dilution in ways that make your company uninvestable for that fund. If a VC needs 15% and your valuation produces 8%, either adjust the valuation or find a different investor. Arguing about philosophy rarely changes the math.

You qualify investors before pitching them. A 30-minute partner meeting is expensive. Spend twenty minutes researching whether their fund size, stage, sector focus, and fund cycle actually align before requesting one.

The founders who raise efficiently are not usually the ones with the best pitch decks. They are the ones who understand what they are asking for and who they are asking it from.

Venture capital decisions are driven by fund mathematics, not personal taste. Power law returns mean VCs need outliers, not averages. Ownership targets determine valuation maximums. Fund size determines stage fit. And most rejections have a structural cause that has nothing to do with the quality of your company. Understanding this does not guarantee a raise, but it focuses your energy on the investors who can actually write the check and gives you the right framework to interpret the ones who cannot.

001 You Have 93 Seconds. Most Founders Waste 80 of Them.

What Is the 93-Second Rule for Pitch Decks?

The 93-second rule is the observed average time venture capital investors spend reviewing a startup pitch deck before deciding whether to pursue the conversation. It comes from DocSend's analysis of thousands of investor interactions and reflects a structural reality: investors see hundreds of decks per month and are filtering for reasons to say yes, not reading for comprehension.

Feature	Convertible Note	SAFE
Accrues interest?	Yes	No
Has maturity date?	Yes	No
Cap and discount?	Both common	Both common
Investor protection	Slightly stronger	Slightly weaker
Founder-friendliness	Moderate	Higher

In markets where SAFEs are well understood (US, increasingly UK), they are generally preferable from a founder perspective. In markets where they are less common, note holders may prefer the debt structure as more familiar.

Convertible notes are the default pre-seed instrument for good reasons: they are fast, cheap to document, and defer the valuation debate to when there is more data. The mistakes happen when founders sign without modelling the dilution, set the cap without understanding what it means at different Series A valuations, and ignore interest and maturity provisions that seem small now but matter at conversion. Run the dilution model before agreeing to terms. Understand the cap vs. discount interaction. Know your maturity date and have a plan for what happens if the qualifying round is delayed.

007 Not All Money Is Equal. How to Pick the Right VC for Your Stage.

Why VC Selection Matters Beyond the Check

Most first-time founders approach the VC selection question from the wrong direction. The question they ask is: who will invest in us? The better question is: of the investors who would fund us, which ones should we take money from?

This distinction matters because a VC relationship typically lasts seven to ten years. The investor who leads your round often has a board seat, information rights, and significant influence over future financing decisions. Taking the wrong partner's money is not like a bad hire you can undo in three months.

Typical VC relationship duration	7-10 years
Average board seats per partner	8-12 companies
Most common founder regret post-investment	Partner responsiveness
Best source of VC reference checks	Portfolio founders, not VC websites
Stage mismatch risk	Lead investor focused on later stage

The Stage Fit Question

Every VC firm has a stage sweet spot where they do their best work. Some funds are exceptional at pre-seed, where the primary skill is founder evaluation without much business data. Others are exceptional at Series B and beyond, where the skill is scaling a proven model. Many funds invest across stages but have internal preferences.

Raising from a fund that is primarily a growth-stage investor for your seed round creates specific risks:

The partner on your deal may be less engaged because your company is not the firm's priority investment type. The fund's diligence process may be designed for later-stage metrics you do not yet have. And the partner's time is divided across 10+ portfolio companies where larger bets are at stake.

Ask directly: "What percentage of your portfolio is at our stage? What is the typical check size for investments at this stage relative to your fund size?"

A $200M fund writing a $500k seed check is doing a favour, not a primary investment. That changes the relationship.

Sector and Business Model Expertise

A VC who has invested in fifteen marketplace businesses will add different value to a marketplace founder than one who has invested primarily in SaaS. They understand the unit economics, the scaling levers, the common failure modes, and the likely buyers at exit. That pattern recognition is worth more than the capital for many founders. At Creandum, the portfolio expertise in European SaaS and marketplace businesses was directly useful in conversations about growth model assumptions and comparable benchmarks. That is the conversation that only happens when the investor has seen the pattern before. Questions to ask: "Which companies in your portfolio are most similar to us by business model? Can I speak to two or three of those founders?"

Partner vs. Firm

You are not taking money from a VC firm. You are taking money from a partner at that firm. The person who leads your deal, sits on your board, and answers the phone on a difficult Friday is the individual partner, not the brand.

The same firm can have partners who are deeply engaged and genuinely helpful and partners who are distracted and difficult. The brand does not tell you which you are getting.

The only reliable way to assess this is by talking to founders in the partner's current portfolio. Not the ones the partner recommends. Ones you find yourself through LinkedIn, through the portfolio page, through your network.

Ask the founders: "How available is this partner when things are going poorly? Have they helped you in ways that were not just about connections to other investors? Have there been situations where you disagreed and how did they handle it?"

Portfolio Conflicts

Most VC funds have conflict policies that prevent them from investing in direct competitors to existing portfolio companies. But "direct competitor" is often interpreted narrowly, and indirect competition or market adjacency is common.

Before accepting a term sheet, understand whether any other portfolio company at the fund operates in your market. Even if not a direct competitor today, a portfolio company targeting the same customer segment creates awkward dynamics for the partner sitting in your board meeting.

The VC Selection Matrix

Criteria	Green flag	Red flag
Stage fit	Majority of portfolio at your stage	Primarily later-stage focus
Sector knowledge	Multiple portfolio companies in your space	First investment in your category
Partner availability	Partner currently has capacity for board seats	Partner already has 12+ board seats
Portfolio reference	Founders describe partner as engaged in hard times	“Great when things were going well”

Fund size vs. check	Check is 2–5% of fund	Check is under 0.5% of fund
Conviction signal	Partner led the deal from first meeting	Was passed a deal by an associate

What to Ask Every VC Before Accepting a Term Sheet

1. What is the stage focus of this fund and how many investments are at our stage?
2. Which portfolio companies are most similar to us by business model? 3. Can you connect us with three portfolio founders at companies you have backed through a difficult period?
4. How many board seats does the partner leading our deal currently have?
5. What does your typical involvement look like in the first 12 months post-investment?
6. Are there any portfolio companies we should be aware of that operate in adjacent markets?

The answers to these questions reveal more about the likely investor relationship than any pitch meeting.

VC selection is a long-term partnership decision, not a capital transaction. Stage fit, sector expertise, partner availability, and portfolio founder references are the variables that predict the quality of the relationship, not the brand name on the term sheet. Run reference checks on the partner, not the firm. Speak to portfolio founders who experienced something going wrong, because that is when the quality of the investor relationship actually becomes visible. The right investor adds value beyond the check and is a genuine asset when things are difficult. The wrong investor is a fixture at your board meetings for years.

008 What Creandum, Profounders, and B2Ventures Actually Look For in Your Model

Why the Model Is a Credibility Test

Before a VC writes a check, they are evaluating three things about the model: whether the numbers are internally consistent, whether the assumptions are defensible, and whether the founder presenting the model actually understands what drives it.

A beautiful model with circular logic fails the first test. A model with "industry average" as the only source for every assumption fails the second. A founder who cannot explain why the CAC assumption is $320 and not $500 fails the third.

The model is not just a financial document. It is a proxy for how the founder thinks about their business.

What experienced investors check first:

Check	What they are looking for
Revenue drivers	Is revenue an output of real mechanics or an assumed growth rate?
Gross margin	Is it benchmarked and are COGS correctly separated from OpEx?
Headcount plan	Is it role-by-role with loaded costs or a lump-sum?
Scenario structure	Are there at least three scenarios and does conservative show survival?
Assumption sourcing	Does every key input have a cited benchmark or logic?

6-12 months history. Unit economics calculated from real customer data by segment and channel. Cohort analysis with 6-12 monthly cohorts showing retention trends. Headcount plan for 20-40+ hires with department structure. Sophisticated scenario modeling with 3+ cases, driver-level variation, and decision triggers. 80%+ assumptions sourced from internal data. 3-5 year horizon with monthly detail for Year 1-2, quarterly for Years 3-5. Variance analysis showing forecast accuracy. Gross margin from actual COGS data with trajectory modeling. Multi-variable sensitivity with tornado chart.

> *Key insight: The shift from seed to Series A is not about making the model bigger. It is about replacing estimates with evidence. Every benchmark you used at seed should now be replaced with your own data. Every hypothesis should now be either validated or invalidated. Investors can tell instantly whether a founder has updated their model with real learnings or simply scaled up the seed projections.*

The Actuals Layer: What Changes Everything

The single most important addition to a Series A model is the actuals layer: a row or tab that shows what actually happened alongside what was projected.

At Indeed Flex, when we moved from seed to Series A fundraising, we added an actuals column next to every forecast line. Revenue forecast vs. actual. CAC forecast vs. actual. Churn forecast vs. actual. This transparency was initially uncomfortable because some numbers missed the forecast. But the investor response was consistently positive, and here is why:

Variance between forecast and actuals is normal. What matters is whether you can explain it. A founder who shows a 15% revenue miss and can say "we missed because enterprise sales cycles were 30 days longer than projected, so we have adjusted the model to reflect a 52-day

average cycle based on actual data" demonstrates learning. A founder who shows perfect forecast accuracy either has a two-month-old model or is adjusting the forecast to match actuals (which investors will check).

Metrics That Only Matter at Series A

Several metrics that are irrelevant or unmeasurable at seed become central to the Series A conversation:

Metric	Why It Matters at Series A
Net Revenue Retention (NRR)	Shows whether existing customers are growing or shrinking
Magic Number	Revenue efficiency: net new ARR / prior quarter S&M spend
Burn Multiple	Net burn / net new ARR (lower is better, <2x is strong)
Gross margin trend	Must show improvement trajectory, not just current level
CAC payback by channel	Identifies which channels scale and which to cut
Revenue per employee	Operating efficiency signal Forecast accuracy

seed-stage model could not have provided them because the data did not exist. A Series A model must.

How to Transition Your Model Between Stages

Step 1: Preserve the structure, replace the inputs

Do not throw away your seed model. The architecture is likely sound if you built it driver-based. Replace benchmark assumptions with actual

data. Add granularity where you now have information you did not have before (channel-level CAC, segment-level retention, geographic variation).

Step 2: Add the actuals layer

For every historical month, add a row showing what actually happened. Calculate variance (actual minus forecast, as both absolute and percentage). Add a notes column explaining material variances. This layer turns your model from a projection tool into a management tool. **Step 3: Extend the horizon**

Seed models typically cover 18-24 months. Series A models need 3-5 years with monthly detail for Years 1-2 and quarterly for Years 3-5. The extended horizon shows the path to break-even or next round, which Series A investors require to model their return.

Step 4: Add departmental detail

Seed models can get away with aggregate cost categories. Series A models need department-level P&L: Engineering, Sales, Marketing, Operations, G&A. This granularity shows investors where money is being allocated and whether the allocation is sensible for a company at this stage. **Step 5: Build the narrative bridge**

Add a summary tab that tells the story: where were we at seed? What did we learn? Where are we now? What do we need to do next? This tab should reference the actuals data to show the journey. Investors at Series A are evaluating not just the company's current state but the team's ability to learn and adapt, and the model is where that ability is demonstrated.

A seed financial model is a hypothesis supported by benchmarks and early signals. A Series A model is a data-driven operating plan supported by months of actual performance. The transition requires replacing estimates with evidence, adding cohort-level detail, incorporating an actuals layer with variance analysis, extending the horizon to 3-5 years, and adding metrics that only become measurable with scale. Founders who make this transition earn investor confidence by demonstrating that they learn from their own data and use it to make better decisions. Founders who show up to Series A with a

scaled-up seed model signal that the last 12-18 months of execution taught them nothing about their own business.

011 How to Run a Competitive Fundraising Process Without Burning Every Bridge in Town

Why a Competitive Process Matters

Fundraising is a market. Like every market, the price (your valuation) and terms (your term sheet) are determined by supply and demand. If only one investor is considering your deal, they set the terms. If three investors are considering it simultaneously, you set the terms. This is not theoretical. Across the fundraising processes I have been part of, the rounds where we had competitive dynamics closed at 15-30% higher valuations than the rounds where we had a single interested party. The difference was not in the quality of the company or the model. It was in the process.

Investors know this dynamic as well as you do. A partner at Profounders once told me directly: "We always assume there is competition. If there is not, we wonder why." The absence of a competitive process signals one of two things to an investor: either the founder does not know how to run one, or other investors have already passed. Neither interpretation helps you.

> **Key insight: A competitive process is not about deceiving investors. It is about respecting your own time and theirs by ensuring that decision timelines converge. The goal is not to create artificial urgency. It is to prevent the fundraise from becoming a 6-month serial conversation where every investor waits to see what every other investor does.**

The Structure of a Competitive Process

Phase 1: Preparation (2-4 weeks before outreach)

Before you contact a single investor, everything must be ready. The model, the deck, the data room, the narrative. Nothing kills momentum

like an investor expressing interest and you saying "I will send the financials next week." Next week is too late. Interest has a half-life. Build your target list: 30-50 investors for seed, 20-30 for Series A. Categorize them into three tiers. Tier 1: your dream investors, the ones whose brand and network would transform the company. Tier 2: strong investors who would be excellent partners. Tier 3: solid investors who you would accept if Tier 1 and 2 do not convert. You will approach Tier 2 first (more on why below).

Phase 2: Warm-up (1-2 weeks)

Start with Tier 2 investors. This is counterintuitive but critical. The first 3-5 meetings of any fundraise are your worst meetings. Your pitch is not tight. Your answers to hard questions are not polished. Your model presentation is clunky. Burn those reps on investors you would be happy with but who are not your top choice.

By the time you approach Tier 1 investors, your pitch has been battle-tested. You know which questions come up, where the model gets challenged, and which parts of the story resonate. This sequencing alone can be the difference between a Tier 1 term sheet and a polite pass. **Phase 3: Parallel outreach (1-2 weeks)**

Contact all Tier 1 investors within the same 5-7 day window. The goal is that first meetings happen within 2 weeks of each other. If Investor A meets you on Monday and Investor B meets you three weeks later, there is no competitive dynamic. They are not making decisions in the same window.

Use warm introductions wherever possible. A warm intro converts to a meeting at 5-10x the rate of a cold email. Your existing investors, advisors, and founder network are the distribution channel. **Phase 4: Acceleration (2-4 weeks)**

As investors progress from first meeting to partner meeting to due diligence, keep the timelines aligned. If Investor A requests the data room on Tuesday, send it the same day and notify Investor B that you are entering the diligence phase. You do not need to name the other investors. The signal is sufficient.

The language: "We are moving into due diligence with a couple of funds this week, and I wanted to make sure you have the materials you need to move at the same pace." This is honest, professional, and creates the time pressure that drives decisions.

Phase 5: Term sheet convergence (1-2 weeks)

The goal is multiple term sheets arriving within a narrow window. This gives you the ability to compare terms, negotiate improvements, and choose the best partner. If only one term sheet arrives, you still benefited from the competitive process because the investor priced the deal knowing competition existed.

How to Communicate Without Burning Bridges

The biggest risk of a competitive process is coming across as manipulative. Investors talk to each other. Your reputation follows you. Here are the rules:

Never lie about interest. If you have one term sheet, do not say you have three. If you are in early conversations, do not claim you are in late-stage diligence. Investors verify. Getting caught in a lie ends not just this raise but your ability to raise from that fund permanently. **Be transparent about timelines.** "We are aiming to close by the end of next month" is helpful. It lets investors self-select. Those who cannot move that fast will tell you, and you avoid wasting each other's time.

Respond quickly and consistently. If one investor gets the data room in 24 hours and another waits a week, the second investor knows they are not a priority. Treat every investor in your process with equal responsiveness.

Never use one investor's name to pressure another. "Creandum is looking at this" is off-limits unless Creandum has explicitly authorized you to share that information. The startup community is small. This kind of behavior ends careers.

What the Financial Model Must Enable

A competitive process requires a model that can handle rapid diligence. When an investor asks for sensitivity on three variables, you need to deliver within 24 hours, not a week. When a partner meeting surfaces a new question, you need to open the model and answer it live. This means: a working scenario toggle, a complete assumptions tab, and a model that does not break when someone changes an input. Every hour of model preparation before the fundraise pays back tenfold during the process when speed determines whether investors stay engaged or move on.

A competitive fundraising process is structured parallel outreach that creates natural urgency and improves your terms. Prepare everything before first contact. Start with Tier 2 investors to refine your pitch. Launch Tier 1 outreach in a tight window. Keep timelines aligned as investors progress through diligence. Communicate honestly, respond quickly, and never manufacture fake interest. The financial model is your primary diligence tool during this process, so it must be ready for real-time interrogation. Founders who run competitive processes consistently raise faster, at better valuations, and with investors who respect the professionalism of the approach.

012 Term Sheet Red Flags: What to Watch For Before You Sign Away Control of Your Company

Why Term Sheets Matter More Than Valuation

A $10M pre-money valuation with aggressive terms can be worth less to a founder than a $7M pre-money with clean terms. This is not intuitive, which is why investors with more experience consistently negotiate terms while founders fixate on headline numbers.

During one of the Indeed Flex rounds, we received two term sheets with similar valuations. The differences were in the fine print: liquidation preference structure, anti-dilution mechanics, board composition, and founder vesting. Our legal counsel estimated that the "better valuation" offer was actually worth 20-30% less to the founders in any exit scenario below $50M because of the participating preferred structure. We took the lower-valuation offer with simpler terms, and it was the right call.

> **Key insight: Valuation is what gets announced. Terms are what determine your actual outcome. A sophisticated investor knows this, which is why they will often give you a slightly higher valuation in exchange for terms that protect their downside at your expense. If you do not understand the terms, you cannot evaluate the trade-off.**

The Red Flags, Ranked by Impact

1. Participating Preferred ("Double Dip")

In a standard (non-participating) liquidation preference, the investor chooses: take their preference amount back, or convert to common and take their ownership percentage. In a participating preferred structure, they get both: their preference amount first, plus their pro-rata share of

whatever remains. This is the single most economically significant term in a term sheet.

Example: Investor puts in $3M for 30% at a $10M post-money. Company sells for $15M. Non-participating: investor takes 30% of $15M = $4.5M. Participating: investor takes $3M preference + 30% of remaining $12M = $3M + $3.6M = $6.6M. The founder's share drops from $10.5M to $8.4M. At lower exit values, the impact is even more dramatic. Some participating preferred structures have a cap (e.g., 3x participation cap), which limits the double dip. Uncapped participating preferred is the most aggressive version and should be a dealbreaker for most founders.

2. Multiple Liquidation Preference

A 1x liquidation preference means the investor gets their money back before common shareholders receive anything. This is standard and fair. A 2x or 3x preference means they get 2-3 times their money back first. On a $5M investment with a 2x preference, $10M must be returned to the investor before a single dollar flows to founders and employees. In a $20M exit with a 2x preference on $5M invested: investor gets $10M first, then their pro-rata share of the remaining $10M. Founders and employees split what remains. At lower exit values, they may get nothing. Greater than 1x preferences are a red flag and are typically a sign that the investor is concerned about the company's prospects. **3. Full Ratchet Anti-Dilution**

Anti-dilution protection adjusts the investor's conversion price if you raise a future round at a lower valuation (a "down round"). Broad-based weighted average anti-dilution is standard and reasonable. It adjusts the price proportionally based on how much new money is raised at the lower price.

Full ratchet anti-dilution reprices the investor's entire investment to the new lower price, regardless of how much money is raised in the down round. If you raised at $10M and then raise $100K at a $5M valuation, full ratchet reprices the entire original investment as if it were done at $5M. This can result in massive dilution to founders. Full

ratchet should be rejected in almost every circumstance. If an investor insists, it signals they expect the company to struggle.

4. Board Control

The standard seed-stage board is 3 seats: 2 founders, 1 investor. Series A often moves to 5 seats: 2 founders, 2 investors, 1 independent. The red flag is any structure where investors control a majority of board seats before the company has raised a Series B or later.

Board control means the investor can fire the CEO, approve or block budgets, and make strategic decisions without founder consent. Some investors request board observer seats (no voting power, but present at all meetings) as a compromise. This is generally acceptable. Pay close attention to who selects the independent board member. If the investor has approval rights over the independent seat, they effectively control 3 of 5 seats. Negotiate for mutual agreement on independent directors.

5. Founder Vesting Reset

Some term sheets include a provision that resets founder vesting, meaning your shares go back to a 4-year vesting schedule even if you have been working on the company for 2 years. This is a leverage play: it ties you to the company for another 4 years and means you forfeit unvested shares if you leave or are removed.

Reasonable terms: accelerated vesting for existing shares (recognizing time already served), and a new vesting schedule that starts from the closing date for any new equity grants. A full reset is aggressive and should be pushed back on.

6. Broad Protective Provisions

Protective provisions give investors veto rights over specific company actions: raising new capital, selling the company, changing the business, issuing new shares, taking on debt. Some protection is standard and fair. The red flag is when the list is so broad that the founder cannot make any material decision without investor approval. Watch for: veto on any expenditure above a low threshold (e.g., $10K), veto on any new hire above a certain salary, veto on pivoting the business model. These provisions can effectively transfer operational

control to the investor without them holding a board majority. **7. Redemption Rights**

Redemption rights allow the investor to force the company to buy back their shares after a certain period (typically 5-7 years) if no exit has occurred. For a startup, this can be a death sentence: the company may not have the cash to redeem, which gives the investor leverage to force a sale or take control.

Redemption rights are more common in later-stage deals but occasionally appear in seed and Series A term sheets. If present, ensure the redemption is at the option of the company, not the investor, or negotiate them out entirely.

The Negotiation Approach

You do not negotiate a term sheet by rejecting everything. You negotiate by understanding which terms are standard, which are aggressive, and which are dealbreakers. Then you focus your energy on the 2-3 terms that matter most.

Term	Standard	Red Flag
1x non-participating preferred	✓ Standard	—
Participating preferred	—	Red flag (negotiate to non-participating)
> *1x liquidation preference*	—	Red flag (negotiate to 1x)
Broad-based weighted average anti-dilution	✓ Standard	—
Full ratchet anti-dilution	—	Red flag (negotiate to weighted average)
3-person board (2 founder, 1 investor)	✓ Standard for seed	—

Investor board majority before Series B	—	Red flag
Founder vesting reset	—	Aggressive (negotiate for credit)
Standard protective provisions	✓ Standard	—

Term sheets contain the economic and governance terms that determine your actual outcome as a founder. Valuation gets the attention, but liquidation preferences, anti-dilution provisions, board structure, and protective provisions determine how much of that valuation you actually receive. Participating preferred, multiple liquidation preferences, full ratchet anti-dilution, investor board control, founder vesting resets, and broad veto rights are the red flags that experienced founders and lawyers watch for. Hire a specialist lawyer. Negotiate the 2-3 terms that matter most. And remember that the best negotiating position is having more than one offer on the table.

013 Bridge Rounds: When They Save Your Company and When They Bury It

What a Bridge Round Actually Is

A bridge is not a standard funding round. It is a temporary financing instrument, typically a convertible note or a SAFE, that provides cash to cover operations for 3-9 months. It "bridges" the gap between where you are now and a future event: the close of a priced round, the achievement of a milestone that makes you fundable, or a revenue threshold that changes your negotiating position.

Bridge rounds are most commonly provided by existing investors because they have the most to gain from keeping the company alive (they are already invested) and the most information to evaluate the risk. New investors rarely participate in bridges because the information asymmetry is too high, and the signal is ambiguous.

When a Bridge Makes Strategic Sense

You are 60-90 days from closing a priced round. The term sheet is in negotiation, legal is working on definitive documents, but the cash will not arrive for 2-3 months. A bridge from an existing investor or the incoming lead covers the gap. This is the cleanest use of a bridge and is common in rounds that take longer than expected to close. **You are 3-6 months from a clear milestone that changes your fundraising story.** Your ARR is $800K, and at your current growth rate, you will hit $1.2M in 5 months, which puts you in a much stronger Series A position. A bridge buys that time at a lower dilutive cost than raising a full round at today's weaker metrics.

Market conditions are temporarily unfavorable. The fundraising market tightened (as it did in 2022-2023), and you need 6-12 months for conditions to normalize. A bridge from existing

investors preserves the company until the market reopens. This only works if the underlying business is performing.

When a Bridge Is a Red Flag

You missed your plan and are out of options. If the bridge exists because your seed metrics did not hit Series A thresholds and no new investor is interested, the bridge delays the problem but does not solve it. Unless the bridge funds a specific plan to fix the underlying issue, it is life support, not a strategy.

Your existing investors will not participate. If the people who know your business best will not put more money in, that is the strongest possible signal that the risk is too high. Raising a bridge from strangers in this situation is extremely difficult and usually comes with predatory terms.

The bridge amount is too small to reach the milestone. A $200K bridge that buys 2 months when you need 6 months to hit the milestone just creates the same problem again in 2 months, but now with additional debt or dilution. If the bridge does not fully fund the path to the milestone, it is not a bridge; it is a postponement.

> **Key insight: The question to ask before taking a bridge: "What specifically will be different in 6 months that is not true today?" If you have a clear, credible answer, the bridge is strategic. If the answer is "hopefully the market improves" or "hopefully we figure it out," the bridge is a way to avoid making hard decisions now at the cost of harder decisions later.**

How to Structure a Bridge That Does Not Hurt Your Next Round

Instrument: Convertible note or SAFE

Most bridges are structured as convertible instruments that convert into the next priced round at a discount (typically 15-25%) and/or with a valuation cap. The discount compensates bridge investors for the risk of investing earlier. The cap sets a maximum valuation at which the

bridge converts, protecting the bridge investor from excessive dilution if the next round is at a much higher valuation.

Discount range: 15-25%

Below 15% is unusual because it does not adequately compensate for bridge risk. Above 25% signals distress and can create signaling problems with new investors in the next round. The standard is 20%.

Valuation cap: set it carefully

The cap should reflect a reasonable expectation of the next round's valuation. Set it too low and you give away excessive dilution. Set it too high and bridge investors feel unprotected. A common approach: set the cap at 80-90% of what you realistically expect the next round's pre-money valuation to be.

Amount: enough to reach the milestone, no more

A bridge should be the minimum amount needed to reach the specific milestone it is designed to fund. Raising more than necessary on bridge terms (which are more dilutive than priced round terms) is giving away equity cheaply. Calculate the exact runway needed, add a 2-month buffer, and raise that amount.

What Bridge Rounds Signal to Future Investors

Signal	Investor Interpretation
Bridge from existing investors,	Positive: insiders are confident, clear milestone ahead
Bridge to close a gap while priced	Neutral: standard logistics, no round finalizes
Bridge after missing milestones	Negative: company is struggling, insiders buying time
Bridge from new investors (not	Curious: why are insiders not existing)
Multiple sequential bridges	Very negative: company cannot raise a real round

Bridge with >25% discount or	Negative: priced as distress aggressive terms

Bridge rounds are a tool, not a strategy. Used proactively to reach a clear milestone or close a timing gap, they are smart and common. Used reactively to delay a reckoning, they create worse problems later. Structure bridges as convertible notes or SAFEs with 15-25% discounts and reasonable valuation caps. Raise exactly enough to reach the milestone. Ensure existing investors participate, because their absence is the strongest negative signal. And always answer the question: what will be different on the other side of this bridge that is not true today?

013B When the Raise Fails: Down Rounds, Washout Rounds, and Navigating Financial Distress

Every book on fundraising is written assuming the raise succeeds. This chapter is about what happens when it does not.

Down rounds happen. Bridge rounds collapse. Term sheets expire. Investors pass on companies that later become significant businesses, and investors fund companies that fail to reach their next milestone. The financial distress scenarios are less visible in the startup literature not because they are rare but because they are uncomfortable and because the people who experienced them often prefer not to write about them publicly.

This chapter covers three scenarios: a down round, where you raise at a lower valuation than the previous round; a bridge round that buys time but creates structural problems if it goes wrong; and a failed raise, where no term sheet arrives and the company must restructure or wind down. Understanding these scenarios before they happen is what separates founders who navigate them successfully from those who are surprised by them.

Down Rounds: What They Are and Why They Happen

A down round is an equity financing where the pre-money valuation is lower than the post-money valuation of the previous round. If a company raised its Series A at a 20 million euro post-money valuation and raises its Series B at a 15 million euro pre-money valuation, the Series B is a down round.

Down rounds happen for several reasons. The company missed its milestones and the growth trajectory no longer supports the previous valuation. The market for comparable companies has contracted and multiples have compressed (this happened broadly across venture in

2022 and 2023 as interest rates rose and public market comparables fell 60-70%). A key customer was lost, a product launch failed, or a key hire left. Or the company raised at an inflated valuation in a frothy market and is now raising in a rational one.

The economic consequences of a down round are real and felt immediately by founders and early employees. Anti-dilution provisions, which appear in most preferred stock term sheets, adjust the conversion ratio of earlier preferred shares to compensate investors for the lower price. In a full-ratchet anti-dilution structure, earlier investors effectively receive additional shares as if they had invested at the new lower price. In a weighted-average anti-dilution structure (more common and more founder-friendly), the adjustment is proportional to the size of the down round relative to the existing share base. Both mechanisms transfer value from common shareholders to preferred shareholders.

A worked example of weighted-average anti-dilution: Series A investors paid 10 euros per share. The new Series B price is 6 euros per share. Using a broad-based weighted-average formula, the conversion price for Series A preferred shares adjusts from 10 euros to approximately 8.5 euros, meaning Series A investors receive more common shares upon conversion. The founder and early employee equity is diluted further beyond the dilution from the new shares issued in the down round itself.

Option pool refresh is another painful element of a down round. The new investor will almost certainly require a refreshed option pool, which is created before the round (diluting founders and existing investors) rather than after (diluting all including the new investor). At a lower valuation, the option pool represents a larger percentage of the company than the same number of options would have at the previous valuation.

> *The founders who navigate down rounds best are the ones who communicate early. An investor who has been receiving transparent monthly updates and understands the variance against plan is in a different*

emotional position than one who receives a down round term sheet as their first indication that something is wrong. Transparency before the down round does not eliminate the down round, but it preserves the relationship through it.

How to Handle a Down Round

The first decision in a down round is whether to accept it. This sounds obvious but it is not. Some founders are better served by a smaller bridge to a cleaner equity story than by a down round that permanently resets the cap table economics. The analysis requires modelling two paths: the down round with its anti-dilution consequences, option pool refresh, and psychological impact on employees whose options are now underwater; and a bridge that extends runway to a stronger position, with the risk that the stronger position does not materialise.

If the down round is the right path, the negotiation priorities change from a normal equity round. Pay-to-play provisions are common in down rounds: existing investors who do not participate in the new round have their preferred stock converted to common, losing their liquidation preferences. This mechanism forces existing investors to support the company by investing pro rata or accepting economic punishment. From a founder perspective, pay-to-play clarifies quickly which existing investors are genuinely supportive versus passive.

Founder vesting acceleration is worth negotiating in a down round context. If the down round resets the company's clock in some sense, a negotiated single-trigger acceleration or a vesting reset (starting a new four-year schedule from the down round date) protects founder retention without additional equity grants.

Communication with employees whose options are underwater is one of the most difficult management tasks in a down round. Options with an exercise price above the current share price have no intrinsic economic value at that moment. The honest conversation acknowledges this while contextualising the path forward: if the company reaches its

next milestones and raises at a recovery valuation, the options regain value. Repricing options (formally reducing the exercise price to the new per-share price) is an option in some jurisdictions but has tax, accounting, and legal implications that require specialist advice.

Bridge Rounds: When They Help and When They Don't

A bridge round is an interim financing designed to extend runway to a specific milestone, after which a full priced equity round becomes possible. The structural form is almost always a convertible note or SAFE, converting into the next equity round at a discount, sometimes with a cap.

Bridge rounds have a bad reputation that they do not fully deserve. Many excellent companies have used bridges to navigate between rounds without damage to the business. The bridge is a problem when it becomes a habit: when one bridge leads to another, when the milestone the bridge was meant to fund keeps shifting, and when the cap table accumulates a series of convertible instruments at different caps and discounts that cannot be resolved cleanly in any future financing.

The financial model for a bridge round must answer three questions: How much runway does this bridge actually buy, net of the legal and transaction costs? What specific milestone will the bridge fund, and how confident am I in that milestone given the current trajectory? What does the cap table look like after the bridge converts into the next equity round, including the discount, the cap, and any interest that has accrued?

A bridge that creates a clean path to a stronger equity story is a legitimate and often sensible tool. A bridge that merely defers a difficult conversation about the business adds 6-12 months of runway at the cost of a messier cap table and the same difficult conversation, now with additional accumulated dilution and lower negotiating leverage.

The warning signs that a bridge is the wrong decision: the founding team disagrees about whether the underlying business is working; existing investors are offering the bridge but not actively working to find a lead for the next round; the company is burning more than

80,000 euros per month and the bridge will provide less than 12 months of runway; the milestone being funded has already slipped once.

The Failed Raise: Options When No Term Sheet Arrives

A failed raise is not a failed company. It is a failed round at a specific moment in time. The distinction matters because the decisions available after a failed raise depend on whether the underlying business has value that the current financing structure cannot surface.

The first option is a structured wind-down. If the business has not achieved meaningful traction and the runway is below four months, an organised wind-down that returns whatever value remains to stakeholders is sometimes the most responsible path. This means stopping hiring, renegotiating or terminating contracts, providing employees maximum notice, and distributing any remaining cash after liabilities are settled. A clean wind-down preserves the founder's reputation for future ventures in a way that a chaotic one does not.

The second option is a sale. A company that has not achieved the growth required for a venture return may still have genuine value: a product that a larger company wants, a customer base worth acquiring, a team with specific technical skills in demand, or intellectual property. Strategic acquirers evaluate acquisitions on criteria that are different from VC return requirements. A 5 million euro acquisition that returns 1.5x to seed investors is not a venture success, but it is a successful exit compared to zero.

The third option is a pivot to profitability. If the company has meaningful revenue and can reach break-even by cutting costs significantly, it may be possible to survive without new external capital. This requires reducing the team (the most painful decision in startup finance), renegotiating supplier terms, and sometimes giving up geographic or product scope to focus on the profitable core. Companies that successfully pivot to profitability after a failed round often find that the constraint forced a discipline that made the business stronger.

The fourth option is finding a different capital source. A failed VC round does not foreclose RBF, government grants, strategic investment from a corporate partner, or a restructured convertible bridge from a single committed angel. Understanding the full menu of capital sources means a failed VC process is the end of one path, not the end of all paths.

What the Financial Model Must Show in a Distress Scenario

When a company is in financial distress, the financial model serves a different purpose than it does in normal fundraising. It is no longer a growth story. It is a survival analysis.

The distress model has three mandatory components: a week-by-week cash flow for the next 90 days (not monthly), showing exactly when the company runs out of cash under current operations; a cost reduction scenario that shows the runway extension achievable by making specific cuts to headcount, contractors, software subscriptions, and discretionary spending; and a revenue upside scenario that models the cash impact of specific deals or contracts that are in the pipeline.

Every assumption in the distress model must be conservative. In a growth model, you can defend an aggressive assumption with reference to the trend. In a distress model, any optimistic assumption that fails to materialise can mean payroll is missed. Build the distress model as if every deal in the pipeline takes twice as long to close, every cost reduction takes a month longer than planned, and cash collections arrive two weeks later than contracted. If the company can survive that model, it is in a position to manage through the crisis.

> *The founders who survive financial distress are almost always the ones who made the difficult decisions 60 days earlier than they felt necessary. Cutting costs feels premature when there is still three months of runway. The decision feels obvious when there is six weeks left. The outcome is the same but the options are very*

different. Three months of runway allows for a structured process. Six weeks forces a panicked one.

014 How to Negotiate Valuation Without Killing the Deal

How Investors Actually Determine Valuation

Founders think valuation is arbitrary. It is not. Investors use a combination of methods, and understanding them gives you the ability to negotiate from an informed position.

Method 1: Comparable transactions

What did similar companies raise at, and at what stage? If three B2B SaaS companies at your stage raised at $10-15M pre-money in the last 6 months, that is the band. Your job is to argue why you are at the top of the band, not outside it.

Method 2: Revenue multiples

For companies with revenue, investors apply a multiple to ARR or forward revenue. SaaS companies at seed might raise at 20-40x ARR. At Series A, 15-25x ARR is common. Marketplaces are valued on net revenue, not GMV, and the multiples are different. The multiple itself depends on growth rate, retention, and margin profile.

Method 3: Return math

Every VC has a fund return target. If a fund needs 3x net returns and invests at a $12M post-money valuation, they need the company to be worth at least $36M at their ownership percentage to make the math work for that deal alone. In practice, because most portfolio companies fail, they need the winners to return 10-50x. This means the investor is evaluating whether your company can plausibly reach a $100M+ outcome, and the valuation must give them enough ownership to make that return meaningful.

What Gives You Negotiating Leverage

Leverage Factor	Impact on Valuation

Multiple term sheets	Highest impact. Competition sets the price.
Strong growth metrics	High. Revenue growth >3x YoY commands premium.
High NRR (>120%)	High. Proves expansion without new acquisition.
Well-known co-investors or angels	Moderate. Social proof validates the opportunity.
Founder track record	Moderate. Prior exits or domain expertise.
Favorable market timing	Moderate. Hot sectors command higher multiples.
12+ months of runway remaining	Moderate. You can walk away from bad offers.
Proprietary technology or data	Moderate. Defensibility commands premium.

$3M" is not a negotiating position. "Three funds are in our process, our metrics put us in the top quartile of comparable raises, and we are targeting a round that gives us 24 months of runway" is a negotiating position.

The Negotiation Conversation

Here is how the valuation conversation typically unfolds and how to handle each phase:

Phase 1: The investor anchors

The investor proposes a valuation range or specific number. This is their opening position, not their final offer. If they say $8M pre-money

and your research suggests comparable companies raised at $10-12M, you have room to negotiate.

Do not accept the first number. Do not reject it either. Acknowledge it and provide your perspective: "Thank you for the proposal. Based on comparable transactions we have been tracking and our current metrics, we have been thinking about a range of $10-12M pre-money. Let me walk you through the data behind that."

Phase 2: The data exchange

This is where your financial model earns its keep. Show the investor your unit economics trajectory, your cohort data, your growth rate relative to comparables. The model is the evidence that supports your valuation position. If your numbers are strong and well-documented, the investor's internal model will independently validate a higher number. At Indeed Flex, during the Series A negotiation, we used the model to demonstrate that our unit economics by cohort were improving faster than the comparable set. That data point alone moved the conversation by $2M because it demonstrated a trajectory that justified a higher multiple. **Phase 3: The negotiation on specific terms**

If the valuation gap is small (10-15%), it often closes through a combination of headline price adjustment and term modifications. An investor might accept a higher valuation if they get a board seat, stronger protective provisions, or a larger option pool carved out before the investment. These trade-offs are where the real negotiation happens.

Know in advance which terms you will trade. A slightly larger option pool (which dilutes you, not the investor) in exchange for a $1.5M higher valuation might be a good trade if you planned to create that pool anyway. A board seat you were going to offer regardless can be positioned as a concession.

> **Key insight: The most effective negotiators I have seen in startup fundraising are not the ones who demand the highest price. They are the ones who make the investor feel that the agreed valuation is fair and evidence-based. An investor who*

> *believes the valuation is justified by the data becomes an advocate for the deal inside their partnership. An investor who feels squeezed becomes a reluctant participant who will be difficult to work with for the next 5-10 years.**

What Not to Do

Do not name a number first if you can avoid it. Let the investor anchor. If they ask "what valuation are you looking for?", redirect: "We have been focused on finding the right partner and letting the market inform the pricing. What range do you see based on our metrics?" **Do not negotiate against yourself.** If you say $12M and they say $10M, do not immediately say $11M. Ask what drives the $10M number. Understand their reasoning before adjusting.

Do not use ultimatums. "We will not take less than $15M" is a negotiation killer unless you genuinely have an alternative at that price. If they call your bluff, the conversation is over.

Do not optimize for valuation at the expense of everything else. A $12M pre-money with a good investor, clean terms, and supportive board is worth more than $15M pre-money with aggressive terms, a controlling board, and an investor who resents the deal.

Valuation negotiation is a structured conversation, not a confrontation. Understand how investors derive valuations (comparables, revenue multiples, return math). Build leverage through competitive processes and strong metrics. Let the investor anchor, then respond with data. Trade terms strategically where it benefits both sides. Prioritize a fair deal with an enthusiastic investor over a maximized number with a resentful one. The financial model is your primary negotiation tool: clean, defensible numbers that independently validate your position are worth more than any negotiation tactic.

015 What Happens When Investors Challenge Your Numbers (And How to Win the Conversation)

Why Investors Challenge Models

An investor challenging your numbers is not trying to embarrass you. They are doing three things simultaneously:

Testing your understanding. Can you explain why the number is what it is? A founder who says "my CTO estimated that" when asked about infrastructure costs has outsourced their understanding. A founder who says "our AWS bill is $0.42 per user per month based on the last 6 months of billing data, and it scales sub-linearly because of caching efficiencies" owns their numbers.

Testing your honesty. When they find a weak assumption, do you defend it stubbornly, or do you acknowledge the uncertainty? Investors are not looking for founders who are always right. They are looking for founders who know what they do not know and are honest about it. **Testing your model's integrity.** When they change an assumption, does the model hold together? Do the outputs make sense? Or does something break? A model that survives investor stress-testing earns a level of trust that accelerates the entire diligence process.

The Ten Challenges That Come Up Every Time

1. "Your revenue projections are too aggressive."

What they are really asking: show me the drivers behind these numbers. Winning response: "Let me walk you through the construction. Revenue is built from [number] leads per month at a [X%] conversion rate, producing [Y] new customers at $[Z] ACV. Here is the historical trend for each of those inputs. Which specific driver looks aggressive to you?"

This shifts the conversation from "your number is wrong" to "let us discuss which input you would change." That is a collaborative conversation, not a defensive one.

2. "Your churn assumption is too low."

What they are really asking: do you have real retention data or is this a guess?

Winning response: show cohort data. If you have 6+ months of cohort retention data, present it. If your observed churn is 4.2% monthly and the investor thinks it should be 6%, you can say: "Our observed churn across 8 monthly cohorts is 4.2%. The earliest cohort has 14 months of data. Here is the retention curve. Where in this data do you see a reason to model 6%?"

If your data is thin, say so: "We have 4 months of retention data showing 3.8% monthly churn. That is early and may not be stable. We have modeled the conservative scenario at 6% to account for that uncertainty."

3. "How do you get CAC this low?"

What they are really asking: is this sustainable or is it early-stage organic traffic that will not scale?

Winning response: break CAC by channel. Show which channels are producing at what cost. Acknowledge that organic and referral channels produce lower CAC that may not scale linearly. Show the paid channel CAC separately and explain how it trends as spend increases.

4. "Your headcount plan is too lean / too heavy."

What they are really asking: can this team execute the plan, and is the burn rate sensible?

Winning response: connect each hire to a business outcome. "The second AE starts in Month 5 because our model shows we need 16 new deals per month by Q3 to hit the revenue target, and each AE at ramp capacity closes 8. Here is the ramp assumption and the revenue it enables." **5. "What happens if you miss by 30%?"**

What they are really asking: do you have a plan for the downside? Winning response: open the conservative scenario. "In our conservative

case, where we model 30% below base on new customer acquisition and higher churn, here is what happens. Runway extends to Month X. We defer these hires. We reduce marketing spend here. We still reach [milestone] by [date], though on a longer timeline."

6. "These margins do not match industry benchmarks."

What they are really asking: is your cost classification correct? Winning response: walk through COGS line by line. If your gross margin is lower than benchmarks, explain the specific costs that drive it and whether they improve with scale. If higher than benchmarks, explain what you are doing differently. Reference the assumptions tab sourcing. **7. "Your TAM is too large / unrealistic."**

What they are really asking: do you understand your actual addressable market?

Winning response: show the SAM and SOM, not just TAM. "Our TAM is $4B. Our SAM, filtered for our geography, segment, and pricing, is $400M. We are targeting 2% of SAM in 3 years, which is $8M ARR. Here is why that is achievable based on our current trajectory."

8. "Your model does not account for seasonality."

Winning response: if you have seasonal data, add it. If not: "You are right. We have not yet modeled seasonality because we have fewer than 12 months of data. Once we have a full year, we will layer in seasonal adjustments. In the meantime, our conservative scenario provides a buffer for potential seasonal dips."

9. "When do you run out of money?"

What they are really asking: how much time pressure is driving this raise?

Winning response: answer directly. "At current burn, our zero-revenue runway is X months. In the base case, we reach cash flow break-even in Month Y. In the conservative case, we would need to raise again in Month Z or implement cost reductions starting in Month W." Evasiveness on runway destroys trust faster than any other question.

10. "Can I take the model and run my own assumptions?"

This is not a challenge. This is the best possible outcome. It means the investor is seriously engaged and wants to stress-test the model themselves.

Winning response: "Absolutely. The assumptions tab has all key inputs in one place with a scenario toggle. Here is a brief guide to the tab structure. Let me know which assumptions you would like to explore and I am happy to walk through the logic." *Key insight: The pattern across all ten challenges is the same: the investor is not looking for a founder who has the perfect answer. They are looking for a founder who understands their own model well enough to have an honest, data-informed conversation about any part of it. That understanding is what they are funding.*

Investor challenges are not obstacles. They are opportunities to demonstrate operational command. The founders who handle challenges best are the ones who built driver-based models with sourced assumptions, cohort data, and working scenarios before the first meeting. When a challenge comes, they do not defend. They explore. They open the model, show the data, acknowledge uncertainty where it exists, and engage the investor in a collaborative analysis. That is what earns term sheets.

016 How to Build a Data Room That Closes Rounds Instead of Stalling Them

Why the Data Room Matters

Due diligence is the phase between the investor's verbal commitment and the wire transfer. It is where deals die. Not because of dramatic revelations, but because of friction. Documents that take days to produce. Numbers that do not match between the model and the accounting records. Cap tables with errors. Employment agreements that cannot be located.

At Indeed Flex, we maintained a live data room from Series A onward. When the exit to Indeed began, the data room was 90% ready on day one. This saved weeks during the six-month diligence process and sent a clear signal to the acquirer: this company is well-managed. The financial benefit of that signal is impossible to quantify precisely, but the acquirer's legal team commented multiple times on how organized the process was. That kind of operational credibility translates directly into smoother negotiations.

> **Key insight: Build the data room before you need it. The worst time to organize your corporate documents is during a live fundraise when every day counts. The best time is on a quiet Tuesday three months before you plan to raise.**

The Data Room Structure

Organize your data room into these sections. This structure is based on what institutional VCs (Creandum, Profounders, B2Ventures) and corporate acquirers expect to see.

1. Corporate and Legal

Articles of incorporation, certificate of incorporation, all amendments. Shareholder agreements and any side letters. Board meeting minutes (all of them). Any IP assignment agreements. Material contracts (customer agreements, vendor agreements, partnership agreements). Privacy policy and terms of service. Any regulatory filings or licenses relevant to your business.

2. Capitalization

Cap table (fully diluted, current). Stock option plan documents and individual grant agreements. SAFE or convertible note agreements. Any warrants outstanding. A waterfall analysis showing returns at 2-3 exit valuations. Pro-forma cap table showing post-round ownership. **3. Financial**

Financial model (the full working spreadsheet, not a PDF). Historical financial statements (P&L, balance sheet, cash flow) from inception or at least last 2-3 years. Current month's management accounts. Bank statements for the last 6-12 months. Tax returns for all years the company has filed. Any debt agreements (venture debt, loans, lines of credit). Accounts receivable and payable aging reports.

4. Metrics and KPIs

Monthly KPI dashboard (MRR, ARR, churn, CAC, LTV, NRR). Cohort analysis (revenue retention by monthly cohort). Customer list with contract values and start dates (anonymized if needed for early-stage diligence). Pipeline data from CRM. Unit economics by channel and segment. **5. Team and HR**

Org chart. Key employee bios and LinkedIn profiles. Employment agreements for founders and key executives. Stock option vesting schedules. Any outstanding HR issues, disputes, or litigation. Contractor agreements. Hiring plan (from the financial model). **6. Product and Technology**

Product roadmap. Technical architecture overview (not detailed code, just the system design). IP documentation (patents, trademarks, copyrights). Third-party technology dependencies and licenses. Security and compliance documentation (SOC 2, GDPR, relevant certifications). **7. Market and Customers**

Customer case studies or testimonials (with permission). NPS or satisfaction data if available. Market sizing analysis. Competitive landscape overview. Any published research or reports about your market.

Data Room Best Practices

Use a professional data room provider or a well-organized Google Drive. For seed rounds, a clearly structured Google Drive folder is sufficient. For Series A and M&A, consider dedicated platforms (Datasite, Ansarada, DocSend) that provide access tracking, permissions control, and audit trails.

Name files consistently. Use a naming convention:

[Section]-[Document Type]-[Date]. Example:

"03-Financial-Model-2025-03.xlsx" or "02-Cap-Table-2025-Q1.pdf". Never use names like "Final_v3_FINAL_updated.xlsx".

Keep documents current. Nothing undermines credibility faster than a financial model from 6 months ago sitting next to current bank statements. Update the model, KPI dashboard, and management accounts monthly.

Include a data room index. A single document at the top level that lists every document in the room, its location, and its date. This saves the investor's legal team hours and signals professionalism. **Control access.** Know who has access to what. Track when documents are viewed. This is not paranoia. It is standard practice and helps you understand which parts of your business the investor is focused on.

The Documents That Delay Rounds Most Often

Document	Why It Delays
Cap table with errors	Investors cannot model ownership; legal halts until fixed

Missing IP assignment agreements	Creates ownership uncertainty; legal risk flag
Financial model that does not match Revenue in model does not match accounts	P&L; trust breaks down Incomplete employment agreements
No board minutes	Governance red flag; suggests informality
Stale financial data	Investor cannot assess current state; asks for updates

A data room is not a storage folder. It is a structured communication tool that tells investors: this company is well-managed, transparent, and ready for institutional capital. Organize it into clear sections (corporate, cap table, financial, metrics, team, product, market). Name files consistently and keep them current. Build it 2-3 months before you raise. Control access and track engagement. The data room does not close the round by itself, but a bad one can certainly prevent a round from closing.

017 The Fundraising Timeline: What Actually Takes How Long (And Why It Always Takes Longer Than You Think)

The End-to-End Timeline

Here is what a typical seed round looks like, end to end. Series A timelines are similar but the diligence and legal phases are longer.

Phase	Duration	What Happens
Preparation	2-6 weeks	Model, deck, data room, target list
Outreach and first	2-4 weeks	Intro emails, first meetings
Due diligence	2-6 weeks	Data room, reference calls, legal review
Term sheet negotiation	1-2 weeks	Valuation, terms, legal redlines
Legal and closing	2-4 weeks	Definitive docs, signatures, wire
Total	12-26 weeks	3-6 months is typical

longer because of investor scheduling conflicts, partner meeting delays, holidays, market hesitancy, or the need to restart the process after initial investors pass.

Phase-by-Phase Breakdown

Phase 1: Preparation (2-6 weeks)

This phase takes the longest relative to what founders expect. Building a genuinely investor-ready financial model (not a back-of-envelope version) takes 2-4 weeks for a founder doing it themselves, or 1-2 weeks with professional support. The pitch deck takes another week to get right. The data room takes 1-2 weeks to assemble if you have not been maintaining corporate records (which most early-stage startups have not).

The most common mistake: starting outreach before materials are ready. An investor who expresses interest and then receives a half-finished model or a disorganized data room 10 days later has already lost confidence. Prepare everything before you send the first email.

Phase 2: Outreach and first meetings (2-4 weeks)

Getting meetings takes 1-3 weeks from first outreach. Warm introductions convert at 30-50%. Cold emails convert at 2-5%. You need 15-20 first meetings for a seed round and 8-12 for Series A. At a conversion rate of 30%, that means 50+ warm intro requests for seed and 30+ for Series A. Scheduling is the silent killer. Investors are busy. Getting on a calendar takes 5-10 business days in most cases. If your top investor is traveling for 2 weeks, your timeline just slipped. Account for this.

Phase 3: Partner meetings (2-4 weeks)

After a positive first meeting, the investor invites you to present to their partnership. This is the decision-making body. Getting a partner meeting scheduled takes 1-2 weeks. The meeting itself might last 45-90 minutes, and the partnership will deliberate afterward.

Some funds decide within 48 hours of a partner meeting. Others take 2-3 weeks and may ask for additional materials or a follow-up session. During this period, maintain momentum with other investors in your process. A competitive dynamic is your best tool for accelerating partner decisions.

Phase 4: Due diligence (2-6 weeks)

Once an investor signals serious interest, diligence begins. For seed rounds, this is relatively light: model review, reference calls (3-5 calls

with customers, advisors, or former colleagues), and basic legal checks. For Series A, diligence is comprehensive: detailed model audit, extensive reference calls (8-12), market analysis, technical review, and full legal and financial audit.

The diligence phase is where data rooms win or lose rounds. If every document is organized and current, diligence takes 2-3 weeks. If documents are missing or inconsistent, it takes 4-6 weeks, and every week of delay increases the risk of the deal falling apart. **Phase 5: Term sheet negotiation (1-2 weeks)**

Once diligence is substantially complete, the investor issues a term sheet. Negotiation typically takes 3-7 business days for straightforward terms. If there are competing offers, the negotiation may take slightly longer as you compare and potentially play offers against each other. If term sheet negotiation drags past 2 weeks, something is wrong. Either the investor is having second thoughts, their partnership approval is stalling, or there is a fundamental disagreement on terms that a few extra days will not resolve.

Phase 6: Legal and closing (2-4 weeks)

This is the phase everyone forgets to account for. After the term sheet is signed, lawyers on both sides draft and negotiate definitive agreements (stock purchase agreement, investor rights agreement, right of first refusal, voting agreement). This takes 2-4 weeks even when both sides are motivated.

The wire transfer itself takes 3-5 business days. International wires can take longer. Do not count the money as raised until it is in your bank account.

How to Plan Runway Around the Fundraise

The single most important planning decision: start fundraising with at least 6-9 months of runway remaining. Here is why:

Runway When Fundraising Starts	Likely Outcome

12+ months	Maximum leverage. You can walk away from bad offers.
9-12 months	Comfortable. Enough time for a normal process.
6-9 months	Tight. One major delay puts you in distress.
3-6 months	Distressed. Investors sense urgency and price it.
<3 months	Crisis. Bridge or bust. Terms will be painful.

the date you plan to start outreach until the money arrives. If the resulting cash position goes negative before the projected wire date, you need to either start earlier, raise a bridge, or cut costs to extend runway.

> *Key insight: The fundraise timeline should appear in your financial model as a specific assumption with a date range. Show the investor: "We start raising in Month X, expect to close in Month X+4, and have Y months of runway at that point." This demonstrates that you have planned the process, not just hoped for the best.*

What Causes Delays (And How to Mitigate Each One)

Cause of Delay	Mitigation
Materials not ready at outreach	Build everything before first email Slow investor scheduling
Missing data room documents	Assemble data room 2-3 months before raise

Model inconsistencies found in	Audit model against accounting diligence
Legal negotiation on terms	Use NVCA standard documents as baseline
Holiday periods	Avoid launching during December or August

Fundraising takes 3-6 months for seed and 4-8 months for Series A. Plan for the long end of these ranges. Start raising with 6-9 months of runway minimum. Prepare all materials before first outreach. Build and maintain a data room proactively. Model the fundraise timeline into your cash flow projections so you can see exactly when the money needs to arrive. Every week of avoidable delay (missing documents, unready model, slow responses) is a week of runway burned without progress. The founders who raise efficiently are the ones who treat the fundraise as a project with a timeline, milestones, and deliverables, not as a series of conversations that will eventually produce money.

018 Post-Raise: The First 90 Days That Determine Whether the Money Actually Works

Why the First 90 Days Define the Round's Success

The financial model you raised on was a plan. The first 90 days are where that plan meets reality. Every assumption you defended in investor meetings, your CAC, your hiring timeline, your churn rate, is now being tested with real execution and real money.

The transition from fundraising mode to execution mode is harder than founders expect. For 3-6 months, your primary job was selling the story. Now your primary job is delivering it. The mindset shift is significant, and the operational demands are different.

At Indeed Flex, the 90-day period after each round followed a specific structure. It was not something we invented on the first round. It was something we refined through experience, and by the later rounds, it was a documented playbook. What I am sharing here is that playbook.

The 90-Day Framework

Days 1-30: Set the Foundation

Set up board reporting. Your new investors expect regular updates. Establish the format now: monthly investor email (short, factual, structured) and quarterly board meeting. Define the metrics you will report: MRR, burn, runway, KPIs, hiring progress, key wins, key risks. Use the same format every month so investors can track trends, not just snapshots.

Calibrate the model to day-one reality. The model was built before the money arrived. Now update it: actual cash received (net of legal fees), actual current burn rate, revised hiring start dates based on recruiter conversations, updated pipeline data. This recalibrated model

becomes your operating plan. Every variance from this point forward is tracked against this version, not the fundraising version. **Open the first 2-3 roles.** Do not try to fill all planned hires simultaneously. Start with the 2-3 most critical roles from the headcount plan. Job descriptions should have been drafted during the fundraise (a good use of the CEO's non-meeting time). Post them in week one.

Set up the actuals tracking system. If you do not already have it, establish the process for recording actuals alongside your forecast. This is where the model transitions from a fundraising tool to a management tool. Monthly, you will compare actual revenue, costs, and cash flow to the plan, identify variances, and explain them. **Days 31-60: Execute the Hiring Plan**

First hires start. Your first 1-2 hires from the plan should have accepted offers by now. Their onboarding quality sets the standard for every subsequent hire. Invest real time in their first two weeks. A poorly onboarded first hire cascades into underperformance for months. **Run the first variance analysis.** At the end of Month 2, compare actual performance to the plan. Where are you ahead? Behind? Why? This first variance report is the most important one because it tells you how accurate your model actually is. A 5% variance is normal. A 30% variance means something in your assumptions was fundamentally wrong, and you need to investigate.

First investor update. Send your first formal monthly update. Keep it to one page. Structure: key metrics, progress against milestones, hiring update, cash position and runway, one big win, one big challenge. Investors who receive consistent, honest updates are significantly more likely to support you in future rounds or difficult moments. **Days 61-90: Lock in the Operating Rhythm**

First board meeting. Prepare a proper board deck: financial performance vs. plan, KPI trends, strategic priorities, risks and mitigations, and an ask for the board (specific help you need). A well-run first board meeting sets the tone for the entire board relationship. A disorganized one creates a dynamic where the board loses confidence and starts micromanaging.

Assess model accuracy and adjust. After 3 months of actuals, you have enough data to assess whether your key assumptions are holding. If CAC is 20% higher than planned, adjust the model and the budget. If hiring is two months behind schedule, revise the revenue timeline that depended on those hires. The model is a living document. This is the first real update cycle.

Set quarterly OKRs or milestones. With 3 months of data, set specific quarterly targets tied to the milestones you committed to investors. These should come directly from the model: revenue targets, hiring milestones, product milestones, and unit economics thresholds.

The Reporting Structure That Builds Investor Trust

Reporting Type	Cadence and Content
Monthly investor email	1 page: metrics, progress, cash, wins, challenges
Monthly internal review	Actuals vs. plan, variance analysis, action items
Quarterly board meeting	Board deck: financials, KPIs, strategy, asks
Quarterly model update	Recalibrate forecast based on trailing 3 months
Annual planning	Full model refresh with updated assumptions

> *Key insight: The investors who backed you made a bet on your ability to execute. Monthly reporting is not a burden. It is the mechanism through which you prove them right. The founders who report consistently, honestly, and on time build the kind of investor relationship that produces follow-on checks, warm introductions to later-stage investors, and support during difficult periods. The founders who go silent for 3 months and then send an emergency email requesting a bridge get neither.*

Common Mistakes in the First 90 Days

Spending too fast. The temptation to fill every planned hire immediately is strong. Resist it. Hire in the sequence the model dictates, not all at once. Front-loading all hires accelerates burn before revenue can follow.

Not tracking actuals against the plan. If you do not start comparing actual performance to the forecast from Month 1, you will not notice a 15% variance until Month 6, by which point 6 months of compounding error have consumed a meaningful chunk of runway.

Changing the strategy immediately after closing. You just spent 3-6 months selling investors on Plan A. Pivoting to Plan B in Week 3 sends a terrible signal. If you genuinely believe the strategy needs to change, discuss it with your board first and present data.

Neglecting investor communication. Silence is never interpreted positively by investors. If you are too busy to send a monthly update, that tells your investor either the company is in trouble or the founder does not value the relationship. Both are bad.

Ignoring early warning signs. If Month 1 and Month 2 both show CAC 25% above plan, that is not noise. That is a signal. The founders who course-correct in Month 3 protect their runway. The founders who wait until Month 9 hoping it will fix itself are the ones who end up raising an emergency bridge.

The first 90 days after closing a round are where the money either starts working for you or starts draining. Set up board reporting and investor communication in the first week. Calibrate the model to actual cash and revised timelines. Start the top-priority hires immediately. Run your first variance analysis at 60 days. Hold your first board meeting at 90 days. Track actuals against the plan from day one. The discipline you establish in these 90 days is the discipline that determines whether you are raising from a position of strength or desperation 18 months from now.

018B Managing Your Board and Your Investors After the Close

Closing the round is the end of the fundraising process and the beginning of a governance relationship that will last for years. The board you just populated and the investors you just took money from now have legal rights, information rights, and economic interests that are partially aligned with yours and partially in tension with yours. Understanding how this relationship works is one of the most underestimated aspects of being a VC-backed founder.

This chapter covers the practical mechanics of board governance: how to run board meetings, how to manage the relationship between your lead investor and your independent directors, how to use the board effectively as a resource rather than experiencing it as a constraint, and how to navigate the moments when the board's interests and the founder's interests diverge. It also covers investor relationship management beyond the board: the informal communications, the strategic conversations, and the moments when you need your investors to act.

Understanding Your Board's Legal Role

A board of directors has a fiduciary duty to the company, not to any individual class of shareholders. In practice, this duty is interpreted through the lens of Delaware corporate law (for US-incorporated companies) or the relevant national company law for European entities. The practical implication is that board members must consider the interests of all shareholders, including common shareholders (founders and employees) and preferred shareholders (investors), and cannot simply act in the interest of the investors who appointed them.

This matters most in two scenarios. The first is a down round or distressed financing, where the interests of preferred shareholders (who have liquidation preferences) and common shareholders (whose equity may be nearly worthless in a low-exit scenario) can diverge sharply. The

second is an acquisition offer, where the board must evaluate the deal on its merits for all shareholders rather than simply on whether it returns capital to preferred investors.

The Trados case (In re Trados Inc. Shareholder Litigation, Delaware Court of Chancery, 2013) is the clearest illustration of this tension. The Trados board approved a merger that returned capital to preferred shareholders and their VCs but left common shareholders with nothing, despite the company having genuine enterprise value at the time of the merger. The Delaware court found that the board had failed its duty to common shareholders. For founders, the lesson is concrete: if you are in a board meeting approving a transaction that pays out preferred but leaves common with zero, you should ensure the decision is documented as having been made in the interests of all shareholders, not just the preferred.

For European founders, the equivalent concept applies through national corporate law. In the UK, directors must act in the interests of the company as a whole, which includes all shareholders. In Germany, the Geschäftsführer of a GmbH has similar obligations under the GmbHG. The specific case law varies, but the principle is consistent: board decisions that favour one class of shareholders at the expense of another are legally vulnerable.

How to Structure Board Meetings

A board meeting that runs poorly is worse than no board meeting. It signals that the management team lacks operational discipline, gives board members no useful information, and creates an adversarial dynamic between the board and the management team. A board meeting that runs well is one of the most valuable resources a VC-backed company has access to: experienced investors with broad networks, pattern recognition across dozens of companies, and a genuine interest in the company's success.

The materials for a board meeting should be sent at least 72 hours in advance. The standard format includes a financial update (prior month actuals versus plan, current month forecast, cash balance and runway,

year-to-date performance against the annual operating plan), a metrics dashboard (the five to ten KPIs the board and management have agreed are the leading indicators of business health), a functional update from each department head (product, sales, engineering, people), strategic discussion topics that genuinely require board input, and a consent agenda (routine approvals that do not require discussion).

The consent agenda is a tool for preserving board meeting time for strategic discussion. Routine approvals such as minutes from the previous meeting, option grants below a threshold, non-material contract approvals, and travel policy updates can be included in the consent agenda and approved as a block without discussion unless a board member specifically requests to pull an item for discussion.

The ratio of management presentation to board discussion should target 40:60. Most boards get this backwards: management spends 80% of the meeting presenting and 20% is left for discussion. The board's value is in the discussion, not the presentation. Board members can read materials in advance. What they cannot do in advance is challenge assumptions, surface pattern-match insights from other portfolio companies, or engage in the strategic reasoning that management needs from them.

After each board meeting, circulate minutes within 5 business days. Minutes should record decisions made, actions taken (with named owner and due date), and the substance of strategic discussions without attributing specific statements to specific board members. Accurate minutes protect the company legally and create accountability for follow-up items.

Working Effectively With Your Lead Investor

Your lead investor is your most important external relationship for the duration of their investment in the company. This is true whether the relationship is excellent or difficult. The lead investor controls the most board seats, leads future rounds or creates the conditions that make future rounds possible, and sets the tone for how co-investors engage with the company.

The best lead investor relationships are characterised by honest, frequent, informal communication that makes the formal board meeting feel like a review of decisions already made rather than a forum for surprises. This means talking to your lead investor before the board meeting about anything significant, sharing good news promptly but sharing bad news even more promptly, and treating the lead investor as an informed internal partner rather than an external observer.

Proactive bad news sharing is the single most important discipline in investor relationship management. Investors understand that startups encounter problems. What they cannot forgive is discovering problems from sources other than the management team, or discovering problems at a stage where the options for addressing them are already constrained. A lead investor who learns about a major customer churn from the investor update email, rather than a personal call two days earlier, begins to lose confidence in management's judgment before the facts of the situation are even considered.

The corollary is that you should never surprise your lead investor in a board meeting. If you have bad news to share with the full board, call your lead investor the day before. If you have a strategic proposal that requires board approval, walk your lead investor through it informally first. Board meetings where the lead investor is hearing material information for the first time are board meetings where the lead investor is reactive and sometimes defensive rather than engaged and constructive.

Independent Directors: Why They Matter and How to Use Them

Most venture-backed boards include one or two independent directors: experienced operators or executives who join the board without an investor affiliation. Independent directors can serve as counterweights to investor-appointed board members in moments of tension, bring specific functional expertise that neither the management team nor the investors have, and provide a perspective grounded in operational experience rather than portfolio returns.

The most effective independent directors are those who are recruited proactively by the management team rather than nominated by investors. An independent director who is the management team's choice and who has built a relationship with the founder before joining the board is a fundamentally different resource than one who was installed as a condition of a financing.

Recruit independent directors as you would recruit a senior executive: with a clear brief of what you need (specific functional expertise, relevant market experience, specific network connections), a structured interview process, and reference checks. The most common mistake is recruiting independent directors too late, after the board has already developed entrenched dynamics, and finding that the new voice has insufficient context to be effective.

When the Board and Founder Disagree

Conflict between founders and their board is not unusual. It is the natural consequence of people with different information, different incentives, and different time horizons making decisions together about a high-stakes company. The conflicts that destroy companies are not the ones that arise, but the ones that are mismanaged.

The most common board-founder conflict is about pace and risk. Investors have portfolio companies; founders have one company. An investor's portfolio mathematics might rationally support a swing-for-the-fences strategy that the founder, who lives inside the company, considers reckless. Conversely, an investor optimising for a large exit might push for pace and risk-taking that the founder believes could destroy the company they have built. Neither position is obviously wrong. Both reflect rational perspectives from different vantage points.

The resolution mechanism for board-founder conflict begins before the conflict arises: in the term sheet, where voting thresholds and protective provisions determine what decisions the board can make without founder consent, and what decisions require a supermajority. Protective provisions typically require investor approval for significant transactions (acquisitions, issuance of new equity, debt above a

threshold, changes to the certificate of incorporation), but routine operational decisions remain within management's authority. Understanding exactly which decisions require board approval versus which are management decisions is the foundation of productive governance.

When conflict does arise, the structured approach is: make the disagreement explicit and specific rather than allowing it to become a generalised tension; identify the underlying assumptions each party is making and test them explicitly against available data; and bring in a trusted third party (an independent director, an experienced advisor, or a neutral facilitator) if the disagreement is not resolving through normal dialogue.

CEO removal by a board is the most extreme outcome of board-founder conflict. It happens more often than founders expect: a 2016 Noam Wasserman study found that by the time companies reached Series C, more than two-thirds of founders were no longer CEO. The triggers are varied: missing multiple consecutive quarters of plan, loss of key executives, strategic disagreements about the direction of the company, or personal conduct issues. Understanding that board members have the legal authority to remove a CEO and understanding the conditions under which they are likely to exercise it changes how founders engage with the board throughout the company's life.

> *The founders who maintain long tenure and strong board relationships are almost universally the ones who treat board governance as a management discipline rather than a compliance obligation. They prepare thoughtfully, communicate proactively, invite challenge, and treat the board as a resource. The founders who struggle with their boards are almost universally the ones who treat board meetings as a performance to be delivered and board members as constraints to be managed.*

Observer Rights and Information Rights

Investors who do not hold board seats often hold observer rights: the right to attend board meetings without voting, and information rights: the right to receive the same materials that board members receive. These rights are typically granted to smaller investors or early-stage angels who do not warrant full board representation.

Observer rights and information rights create obligations. If your board materials go to 12 people instead of 4, the information contained in those materials is at greater risk of leaking to competitors or to the broader market. Consider what information goes into the standard board package and what information is restricted to voting board members only. Competitive customer information, personnel decisions, and M&A discussions are examples of topics that should be restricted to the voting board rather than shared with all observers.

As the company matures and the investor roster grows, proactively managing the distribution of information rights becomes important. Term sheets for later rounds often include provisions that limit information rights to investors above a certain ownership threshold. This is standard practice and should be flagged to smaller investors as a condition of their investment at the time of the round, not as a surprise in a subsequent financing.

019 How to Write an Investor Update That Keeps Your Round Warm

Why Investor Updates Are Part of Your Fundraising Strategy

There is a common misconception that investor updates are something you do after you raise, to keep your existing investors in the loop. That is partially true. But the more powerful use case is during the fundraise itself. When a VC is tracking your company across multiple conversations over several months, your update is the thread that keeps them connected. It shows trajectory. It shows execution discipline. It shows that you know your numbers and you are honest about your challenges. DocSend's 2024 Fundraising Report found that founders who maintained consistent investor communication during their raise closed rounds on average 30 days faster than those who did not. That is not a coincidence. Investors who already trust your communication style spend less time in diligence asking questions your updates have already answered. I have seen this play out on both sides of the table. At one company, a founder sent a monthly update to a VC they had met but who had not yet committed. Six months later, when the round launched formally, that VC moved faster than anyone in the process because they had been watching the numbers grow for half a year. The update had done the diligence work for us.

The Anatomy of a Good Investor Update

A good update is short, honest, and follows a consistent structure. Aim for 400 to 600 words. Investors receive dozens of these. Long updates do not get read fully. Consistent short updates do. **Section 1: The Headline Number**

Open with your single most important metric this month. For most early-stage companies this is revenue or ARR. If you are pre-revenue, it might be active users, pilot customers, or a key product milestone. Put

it in the first two lines. Do not bury it. Example: "MRR: GBP 47,000 (up 18% month on month). ARR run rate: GBP 564,000." **Section 2: What Went Well**

Two or three specific wins. Not vague positivity. Specific outcomes. Signed a contract with a named customer type. Hit your hiring target. Launched the feature that was blocking three enterprise deals. **Section 3: What Did Not Go Well**

This is the section most founders skip or soften. Do not. Investors know things go wrong. What they are evaluating is whether you know things went wrong and whether you have a plan. A founder who writes "we missed our sales target this month because our outbound sequence was underperforming, so we have rebuilt it and will retest in Q3" is far more fundable than one who writes "it was a slower month but we are feeling positive about Q3." Honesty in the bad section builds more trust than anything you can say in the good section. **Section 4: Key Metrics Table**

Keep this consistent every month. Do not change the metrics you report unless you explain why. Changing what you measure is the fastest way to make an investor think you are hiding something. **Section 5: The Ask**

Every update should have a clear, specific ask. Not "let us know if you can help." Something like: "We are hiring a Head of Sales. Do you know anyone who has scaled a B2B SaaS team from 0 to 10 reps in the UK?" "We are exploring venture debt as a bridge option. Has anyone in your portfolio used Lighter Capital or Clearco?" "We are opening our Series A in Q4. If you are interested in being part of the round, I would love to schedule a 30-minute conversation this month." This trains investors to be useful. It also surfaces who is engaged and who is not.

The Frequency Question

Monthly is the right cadence for most early-stage startups. It is frequent enough to show momentum, infrequent enough that you are not creating noise. Quarterly updates are fine post-Series A when you have a board and a more formalized reporting structure. They are too

infrequent when you are still building relationships with potential investors. Weekly is too often for an external update. That is an internal ops cadence.

What Not to Do

Do not only send updates when things are good. This is the most common mistake. An investor who receives six months of radio silence and then a glowing update right before a fundraise is immediately suspicious. They know they are being sold to. **Do not write updates in prose only.** Numbers without context are confusing. Context without numbers is marketing. Put both in every update. **Do not use updates to manage perception.** If you missed targets, say so. If a key hire left, mention it. Investors who eventually do diligence will find out. Finding out during diligence that things were worse than the updates suggested is one of the fastest ways to kill a deal. **Do not forget your existing investors.** They are your best reference checks, your warmest intros, and often the first to follow on. Keep them as informed as potential new investors.

The Compound Effect

The founders who raise the best rounds are not always the ones with the best metrics. They are often the ones investors trust the most. Trust is built across repeated interactions over time. A monthly investor update is a low-cost, high-return mechanism for building that trust systematically. Start sending updates before you need them. By the time you launch a formal round, you will have a list of investors who have been watching you execute for six to twelve months. That list is worth more than any warm intro. **Recommended format:** Email, plain text, sent on the same day every month. No PDF attachments. No designed templates. Just honest numbers and clear writing.

Chapter Exercise

Write the first version of your investor update using the five-section format above. Use your real numbers from last month. Include one

thing that did not go well. Send it to at least three investors or advisors who have expressed interest in your company, even if you are not yet fundraising. Collect feedback on what they found most useful.

020 The Cap Table: What Every Founder Must Understand Before Taking a Single Dollar

What Is a Cap Table and Why Should You Care Right Now

A capitalization table (cap table) is a spreadsheet that tracks the ownership of your company. It lists every shareholder, how many shares they hold, what class of shares those are, and what percentage of the total they represent on both a basic and fully diluted basis. In the early days, your cap table is simple: the founders own 100% of the company. Once you take money from investors, issue options to employees, or sign any convertible instruments, the cap table becomes more complex. And complexity, if not managed carefully, becomes a problem. The reason cap table literacy matters before you raise is straightforward: the decisions you make in your first round create a structural baseline that compounds through every future round. A cap table that is poorly structured at seed will create friction at Series A and real problems at Series B and exit. Here is a real pattern I have seen repeatedly. A founder raises a small pre-seed round, issues shares informally, never sets up a proper option pool, and moves fast. Two years later, a serious institutional investor runs a fully diluted analysis and finds that the founder's ownership stake has already been diluted to a level that creates misaligned incentives for the founder to keep building. The institutional investor hesitates. The deal slows. In some cases it dies. Get this right early.

The Core Concepts

Authorized Shares vs. Issued Shares

Authorized shares are the total number of shares your company is legally permitted to issue, as set in your articles of incorporation. Issued

shares are the ones that have actually been allocated. You do not need to issue all authorized shares immediately. Most companies authorize more than they issue to give themselves flexibility. **Basic vs. Fully Diluted Ownership**

Basic ownership counts only issued shares. Fully diluted ownership counts everything: issued shares, options granted, options reserved but ungranted, warrants, and any convertible instruments that could convert into shares. Investors always quote ownership on a fully diluted basis. When a VC says they want 20% of the company for a given check, they mean 20% of the fully diluted share count post-investment.

Pre-Money vs. Post-Money Valuation

Pre-money valuation is what your company is worth before the investment. Post-money valuation is what it is worth immediately after the investment is received. Post-money = Pre-money + Investment amount

Investor ownership = Investment amount / Post-money valuation Example: You raise GBP 1.5M at a GBP 6M pre-money valuation. Post-money is GBP 7.5M. The investor owns 1.5 / 7.5 = 20%. This is the single most important calculation in any term sheet. Make sure you are comparing pre-money to pre-money, or post-money to post-money. Mixing them up is a common source of confusion when reading term sheets.

How Dilution Works in Practice

Dilution means your ownership percentage decreases as new shares are issued. Your absolute number of shares stays the same. Your percentage of the total decreases. Example over three rounds:

At Series A, a founder who started at 100% now owns approximately 51%. This is healthy dilution. It means the company raised meaningful capital from reputable investors, and the founder still has a majority stake before the employee option pool is factored in. A founder who started at 60% because of an early co-founder split and took a large pre-seed round on a low valuation might be at 30% by the time they reach Series A. That is where problems start.

The Option Pool: What It Is and How It Affects You

An employee stock option pool (ESOP) is a block of shares reserved for employee grants. Options give employees the right to buy shares at a fixed price (the strike price) at a future date, subject to a vesting schedule. Most institutional investors will require you to have an option pool in place before they invest, typically 10 to 20% of the fully diluted cap table. The critical point is when the pool is created relative to the investment. If the pool is created before the investment (pre-money), it dilutes the founders, not the investor. This is called the option pool shuffle, and it is one of the most common and least-discussed ways that term sheets reduce effective founder economics. Example: You negotiate a GBP 6M pre-money valuation. The investor requires a 15% option pool to be created before closing. The effective pre-money valuation of the shares the founders actually hold is not GBP 6M. It is GBP 5.1M. The 15% pool was carved out before the valuation was applied. This does not mean you should refuse to create an option pool. You need one to hire good people. But you should negotiate the size carefully and understand the impact before you sign. A reasonable option pool at seed is 10 to 12% for most companies. If an investor asks for 20% upfront, ask them to model the hiring plan that justifies it. The pool size should be tied to a specific 18 to 24 month hiring plan, not a round number.

SAFEs and Convertible Notes: Cap Table Complexity Before You Even Raise

If you have raised money via a SAFE (Simple Agreement for Future Equity) or a convertible note, those instruments do not appear as shares today. They will convert into shares at your next priced round, typically at a discount to the round price or at a valuation cap, whichever is more favorable to the investor. This means your current cap table understates your fully diluted ownership picture. Before you model your fundraising, you need to model the conversion of all outstanding SAFEs and notes and understand what your fully diluted cap table looks like on the day the next round closes. Two things to check:

1. If you have multiple SAFEs with different valuation caps, the ones with the lowest caps will convert at the most shares, meaning the most dilution to existing shareholders. 2. If your priced round valuation is below the SAFE cap, the SAFE converts at the round price, which means no discount to the new investor and the most dilution to the founders. Always model your cap table post-conversion before entering a priced round negotiation.

What Investors Are Looking For in Your Cap Table

When an institutional investor reviews your cap table, they are checking four things: **Clean structure.** All shareholders are on one document. No side letters or informal agreements that are not captured. No missing option grant paperwork. No shares issued to people who left without a buyback or cliff. **Founder motivation.** If the founding team collectively owns less than 30 to 35% of the fully diluted cap table at seed, this is a yellow flag. Investors want founders to have meaningful ownership to maintain long-term motivation. Low founder ownership can also signal earlier decisions that were made under pressure or without proper advice. **No messy early investors.** Unsophisticated angel investors who were given large stakes on low valuations with no pro-rata rights or drag-along provisions can complicate future rounds. Sophisticated investors will want a clean cap table with standard institutional terms. **Option pool headroom.** Is there enough room in the option pool to hire the team required to justify the round being raised? If the answer is no, the investor will require a refresh before or alongside the investment.

The One Tool You Need

Use either Carta or Pulley to manage your cap table from day one. Both platforms track shares, options, vesting schedules, SAFEs, and generate waterfall analyses automatically. The cost at early stage is low. The cost of managing your cap table in a spreadsheet and then having to reconstruct it for a Series A data room is very high. If you are pre-revenue and have not yet raised any external capital, you can start

with a simple spreadsheet. The moment you issue shares to anyone outside the founding team, move to a proper platform.

Chapter Exercise

Build your current fully diluted cap table using the following structure: founding team shares, any angel or pre-seed investors, outstanding SAFEs or convertible notes (modeled as if they convert at your target seed valuation), and your current or planned option pool. Calculate the founder ownership percentage on both a basic and fully diluted basis. Then model what that percentage becomes if you raise your next round at three different valuations: your floor, your target, and your stretch. Save this as a tab in your financial model.

PART

Part Two: Financial Modeling

A financial model is not a prediction. It is a structured argument about how your business works. This part teaches you how to build one that holds up in diligence: driver-based revenue, a clean assumptions tab, realistic scenarios, and the specific formats that experienced investors use to test whether the founder understands their own business.

020B Stock Options and Equity Compensation: ISOs, NQSOs, RSUs, and the 409A Valuation

Equity compensation is how early-stage companies compete for talent with established companies that can pay higher salaries. Used well, it creates genuine alignment between employees and the company. Used poorly, it creates unexpected tax bills, frustrated employees who do not understand what they hold, and cap table complexity that complicates future financings.

This chapter covers the mechanics of the four primary equity compensation instruments: incentive stock options (ISOs), non-qualified stock options (NQSOs or NSOs), restricted stock units (RSUs), and phantom equity. It covers the 409A independent valuation that determines the exercise price for options, the vesting mechanics that determine when equity is earned, and the critical tax events that every employee should understand before they exercise or sell.

The context is primarily US-incorporated companies, because most early-stage companies used by European founders as legal entities are Delaware C-corporations. Where meaningful differences apply in UK and European entities, they are noted.

Why the Structure of Equity Compensation Matters

The type of equity compensation instrument determines the tax treatment for both the employee and the company. The choice between ISOs and NQSOs is not merely administrative: it determines whether employees pay tax at ordinary income rates (up to 37% in the US federal system) or at long-term capital gains rates (20% for most employees). For an employee with 100,000 options exercised at a 5 dollar strike into shares worth 25 dollars each, the difference between ordinary income and capital gains treatment is a tax bill that differs by 80,000 dollars or more.

The company's choice of instrument also has accounting implications under ASC 718 (US GAAP stock-based compensation rules), which requires recognising the fair value of options as a compensation expense over the vesting period. This is non-cash but it affects EBITDA and gross margin reported to investors. Understanding stock-based compensation expense is essential for anyone building or reading a startup financial model.

Incentive Stock Options (ISOs)

ISOs are the standard equity instrument for US employees of venture-backed startups. Their primary advantage is tax treatment: if an employee holds ISOs for at least two years from the grant date and one year from the exercise date, the gain upon sale is taxed as long-term capital gains rather than ordinary income. This is the preferential tax treatment that makes ISOs the instrument of choice for employees who believe the company will be worth significantly more than the exercise price.

The exercise price of an ISO must be set at or above the fair market value of the common stock at the time of grant, as determined by the 409A valuation. This is not optional: ISOs granted below fair market value lose their preferential tax treatment and create immediate tax liability under Section 409A.

The ISO 100,000 dollar rule limits the annual value of ISOs that can vest in any single year to 100,000 dollars, measured at the exercise price. Any ISOs above this limit automatically convert to NQSOs. For employees receiving large grants at high exercise prices, this limit can be binding.

The Alternative Minimum Tax (AMT) is the most dangerous aspect of ISOs for employees who do not fully understand their tax situation. When an employee exercises ISOs, the spread between the exercise price and the fair market value at exercise is an AMT preference item, potentially triggering AMT liability even if the employee cannot yet sell the shares to pay the tax. Employees who exercised ISOs in companies that subsequently declined in value during the dot-com bust were left

with tax bills on paper gains that no longer existed. This risk is mitigated by early exercises (83(b) elections, described below) but it is a real risk that employees should be advised about by a qualified tax professional.

Non-Qualified Stock Options (NQSOs)

NQSOs can be granted to anyone: employees, contractors, advisors, and board members. They are not subject to the AMT preference treatment that creates risk for ISO holders, and they are not subject to the 100,000 dollar annual vesting cap.

The tax treatment is less favourable: when an NQSO is exercised, the spread between the exercise price and the fair market value at exercise is taxed as ordinary income in the year of exercise, regardless of whether the shares have been sold. The company receives a corresponding tax deduction equal to the ordinary income recognised by the employee.

NQSOs are the standard instrument for contractors and international employees of US companies who cannot receive ISOs (ISOs are only available to employees, not contractors). They are also used for grants above the ISO 100,000 dollar annual vesting limit and for board member compensation.

For European employees of a Delaware C-corporation, NQSOs are the typical structure. The income recognised on exercise may be taxable in both the US (at the federal and potentially state level) and in the employee's country of residence, subject to the relevant tax treaty. French employees face particularly complex treatment: social charges (employer and employee contributions) apply to the gain at exercise under French law, with effective rates that can make NQSO exercise very expensive. UK employees have more favourable treatment through the EMI (Enterprise Management Incentive) scheme, which allows options to be granted with capital gains tax treatment on exercise if certain conditions are met.

Restricted Stock Units (RSUs)

RSUs represent a promise to deliver shares (or cash equivalent) at a future date, subject to a vesting schedule. Unlike options, RSUs have no exercise price: the employee receives shares of the company on vesting, and the tax event is the vesting date, not a future exercise date.

RSUs are the dominant equity compensation instrument at large technology companies (Amazon, Google, Meta) and at late-stage growth companies approaching an IPO. They are less common at early-stage startups for a specific reason: the tax event at vesting creates immediate ordinary income tax liability on the value of the shares received, whether or not the shares can be sold to pay the tax. In a liquid stock (a public company), this is manageable because the employee can sell shares to cover the tax. In an illiquid private company, the employee receives shares they cannot sell but owes tax on their value.

The practical implication is that RSUs are appropriate for growth-stage companies that are likely to have a liquidity event within 12-18 months of a grant, or for companies that have secondary market programs allowing employees to sell shares to cover RSU tax. For companies that are 4-6 years from a potential exit, RSUs create an untenable tax situation for employees and are generally not used.

RSUs with a double trigger are the standard structure for pre-IPO companies that do use RSUs: vesting requires both the standard time-based schedule and a liquidity event (IPO, acquisition, or secondary tender offer). The double trigger prevents the RSU from creating a tax event until there is a market for the shares.

Phantom Equity and Virtual Equity Plans

Phantom equity (also called virtual equity or a virtual option plan, or VSOP in some European structures) is an agreement to pay cash equal to the economic value of a specified number of shares at a future liquidity event, without issuing actual shares. The employee has no ownership stake in the company but has a contractual right to participate in exit proceeds.

Phantom equity is widely used for international employees and contractors of US companies who cannot receive ISOs and for whom

NQSO treatment would create complex multi-jurisdictional tax obligations. It is also used in European entities (GmbHs, SARLs, SRLs) where the legal structure makes issuing actual equity to employees mechanically complex.

The financial model impact of phantom equity is important to capture correctly. Unlike options, phantom equity creates a cash liability on exit, not a share dilution. In an acquisition scenario, phantom equity holders reduce the net proceeds available to shareholders in the same way that transaction bonuses or employee retention payments do. If the company has committed 5% of exit proceeds to phantom equity holders, the effective exit proceeds to shareholders are 5% lower. This must be reflected in the waterfall model used to calculate returns for each investor class.

Romania and France are markets where phantom equity (or local equivalents) has become standard for startup employees precisely because the local legal structures make issuing actual equity in a Delaware parent company to local employees mechanically difficult and tax-inefficient. The VSOP (Virtual Stock Option Plan) structure common in German startups is the best-developed European equivalent.

The 409A Valuation

Section 409A of the US Internal Revenue Code requires that options granted to employees be priced at or above the fair market value of the underlying shares at the time of grant. Fair market value is determined by an independent appraisal: the 409A valuation.

A 409A valuation must be performed by a qualified independent appraiser and must apply a reasonable valuation method. The most common methods are the Black-Scholes option pricing model (for early-stage companies with a single class of stock), the option pricing model (OPM, which allocates value across the cap table based on the relative payoff to each security in different exit scenarios), and the probability-weighted expected return method (PWERM, which applies different weights to alternative future scenarios).

The gap between the preferred stock price (the price paid by investors in a financing round) and the common stock 409A valuation is one of the most structurally important elements of option compensation. In a startup's early stages, common stock is worth significantly less than preferred stock because preferred stock has liquidation preferences and other protective provisions that common stock lacks. A company that raised its Series A at a 10 dollar preferred share price might have a 409A common stock valuation of 2 to 3 dollars, meaning options can be granted with an exercise price of 2 to 3 dollars while the preferred investors paid 10 dollars.

409A valuations are typically required: at the time of initial option grants; following any material event that affects the company's value (a new financing round, a significant revenue milestone, a strategic partnership); and approximately every 12 months at minimum. A 409A valuation has a safe harbour of 12 months, meaning options granted within 12 months of a 409A appraisal are presumed to comply with Section 409A unless a material change in circumstances has occurred.

> *The 409A exercise price is not the price employees pay to own shares. It is the price they pay to exercise the option into shares. The real economic question for an employee is whether the exit value of the shares will exceed the exercise price by enough to justify the tax liability and the risk of holding illiquid equity. This analysis should be done explicitly by every employee receiving a grant, ideally with guidance from a tax professional.*

Vesting Schedules and Cliff Mechanics

The standard vesting schedule for startup employees is four years with a one-year cliff. The cliff means that no equity vests for the first 12 months; at the 12-month anniversary, 25% of the grant vests immediately. The remaining 75% vests monthly over the following 36 months.

The economic rationale for the cliff is employee retention: if the employee leaves before 12 months, they receive nothing. This protects the company from granting equity to people who are not a long-term fit. The risk is that employees who are not performing but are willing to stay until their one-year cliff anniversary become a management challenge.

Founder vesting is the same structure applied to founding equity. Most investors require founders to accept a vesting schedule as a condition of funding, because the alternative (fully vested founder equity from day one) creates a situation where a co-founder can leave immediately after the round closes and own a large unrestricted stake in the company. The standard is four years with a one-year cliff from the company's founding date, which often means that at the time of the first institutional round, one year of the cliff may already have passed.

Acceleration provisions modify the vesting schedule in specific circumstances. Single-trigger acceleration vests some portion of unvested equity upon a change of control (an acquisition) regardless of what happens to the employee post-acquisition. Double-trigger acceleration requires both a change of control and a qualifying termination (being fired without cause or resigning for good reason within a specified period after the acquisition). Double-trigger is the standard in most term sheets because single-trigger acceleration reduces the acquirer's incentive to retain the employee.

The early exercise right (or Section 83(b) election) allows employees to purchase unvested shares at the time of grant rather than waiting for them to vest. The tax advantage is significant: by purchasing unvested shares at the 409A price at the time of grant, the employee resets their capital gains holding period from the grant date and potentially converts what would be ordinary income (on exercise of vested options later) into long-term capital gains. The 83(b) election must be filed with the IRS within 30 days of exercise and cannot be revoked. For companies in their early stages where the 409A value is very low, early exercise is often highly advantageous.

020C Cap Table Modeling Across Multiple Rounds: A Worked Example

A cap table is more than a list of who owns what. It is a dynamic model of how ownership evolves across every financing event, option grant, conversion, and exit. Founders who understand cap table mechanics can anticipate the dilution consequences of every decision before they make it. Those who do not are frequently surprised: by how much they own at Series B, by the economics of a modest exit, and by why an acquisition at a seemingly attractive headline price leaves them with less than they expected.

This chapter builds a cap table from scratch through five stages: incorporation, pre-seed, seed, Series A, and exit. Every number is shown. The goal is to make the arithmetic concrete enough that any founder can adapt the model to their own situation.

Stage 0: Incorporation

Two founders incorporate a Delaware C-corporation. They issue 8,000,000 shares of common stock, split as follows: Founder A receives 4,500,000 shares (56.25%) and Founder B receives 3,500,000 shares (43.75%). These shares are issued at the par value of 0.0001 dollars per share. Total shares outstanding: 8,000,000.

Both founders subject their shares to a four-year vesting schedule with a one-year cliff, beginning from the incorporation date. They each file an 83(b) election within 30 days of share issuance, electing to be taxed on the value of the shares at issuance (0.0001 dollars per share) rather than on the value at vesting. At current par value, this creates a negligible tax event. It starts their capital gains clock from day one.

At this stage the cap table is: Founder A 56.25%, Founder B 43.75%. Total: 8,000,000 shares.

Stage 1: Pre-Seed Round

Twelve months after incorporation, the founders raise a pre-seed round of 600,000 euros from three angel investors. The round is structured as a SAFE with a 3,000,000 euro valuation cap and a 20% discount to the next priced round. The SAFE does not convert immediately; it converts at the next priced round.

Before the pre-seed, the founders create an option pool of 1,000,000 shares, increasing the total share count to 9,000,000. The option pool is created before the SAFE but is included in the pre-money fully diluted count for the purpose of the SAFE valuation cap calculation.

The SAFE investors have invested 600,000 euros. At the 3,000,000 euro cap, their maximum ownership is 600,000 / 3,000,000 = 20% of the post-money equity, assuming conversion at the cap. At this stage, no shares have been issued to the SAFE investors; they hold a future right to shares.

Option pool: 1,000,000 shares reserved. Total authorised shares: 9,000,000. SAFE outstanding: 600,000 euros notional, unconverted.

Stage 2: Seed Round and SAFE Conversion

At month 20, the founders raise a seed round of 2,000,000 euros at a 9,000,000 euro pre-money valuation. This is a priced equity round, issuing new Series Seed Preferred shares. The SAFE converts into Series Seed Preferred at this point.

SAFE conversion: The cap is 3,000,000 euros. The seed round pre-money is 9,000,000 euros. Since the pre-money exceeds the cap, the SAFE converts at the cap price rather than the round price. The fully diluted share count at pre-money seed (excluding SAFE conversion) is 9,000,000 shares. The implied price per share at the 3,000,000 euro cap is 3,000,000 / 9,000,000 = 0.333 euros per share. The SAFE investors receive 600,000 / 0.333 = 1,800,000 new shares of Series Seed Preferred.

Seed round shares: The pre-money valuation is 9,000,000 euros. Including the 1,800,000 shares issued to SAFE investors, the pre-money fully diluted count is 9,000,000 + 1,800,000 = 10,800,000 shares. The implied price per share for the seed round is 9,000,000 /

10,800,000 = 0.833 euros per share. At 2,000,000 euros raised, the seed investors receive 2,000,000 / 0.833 = 2,400,000 new shares of Series Seed Preferred.

Post-seed option pool top-up: The seed lead requires a refresh of the option pool to 15% of the post-seed fully diluted share count. Post-seed shares are 10,800,000 + 2,400,000 = 13,200,000. 15% of 13,200,000 is 1,980,000. The existing pool has 1,000,000 shares, so 980,000 new option shares are created, increasing the total share count to 14,180,000.

Post-seed cap table (fully diluted):

Founder A: 4,500,000 shares, 31.7%

Founder B: 3,500,000 shares, 24.7%

Pre-seed SAFE investors: 1,800,000 shares, 12.7%

Seed lead investor: 1,800,000 shares, 12.7% (assuming full 2M raised by lead)

Seed co-investor: 600,000 shares, 4.2% (assuming 600K co-invest)

Option pool (unissued): 1,980,000 shares, 14.0%

Total: 14,180,000 shares, 100%

Stage 3: Series A

At month 36, the company raises a Series A of 8,000,000 euros at a 32,000,000 euro pre-money valuation. The lead investor requires an option pool of 12% of the post-Series A fully diluted share count.

Option pool top-up before the round: 12% post-round pool. The pre-A share count is 14,180,000. After the A round and pool refresh, total shares will be 14,180,000 + new A shares + new pool shares. Let S be the new A shares and P be the new pool shares. We need P/(14,180,000 + S + P) = 0.12. This circular calculation is solved iteratively or by working with percentages.

Simplified approach: The new option pool is created to bring the total pool to 12% of the post-round fully diluted count. Assuming the existing unissued pool of 1,980,000 already represents approximately

10% of the pre-A count, and the A round will issue approximately 10-12% to the new investor, the top-up required is approximately 500,000-600,000 additional option shares. These are created before the round (diluting founders and existing investors) for a pre-A fully diluted count of approximately 14,680,000 to 14,780,000.

Series A price per share: Pre-money 32,000,000 euros / 14,730,000 (mid-range) = 2.173 euros per share. New A shares: 8,000,000 / 2.173 = 3,680,000 new Series A Preferred shares.

Post-Series A cap table (fully diluted, approximate):

Founder A: 4,500,000 shares, 24.4%

Founder B: 3,500,000 shares, 19.0%

Pre-seed SAFE investors: 1,800,000 shares, 9.8%

Seed investors: 2,400,000 shares, 13.0%

Series A lead: 3,680,000 shares, 20.0%

Option pool (issued and unissued): 2,550,000 shares, 13.8%

Total: approximately 18,430,000 shares, 100%

Combined founder ownership has declined from 100% at incorporation to approximately 43.4% post-Series A, which is within the healthy range for this stage and financing history.

Stage 4: Exit Waterfall

At month 60, the company is acquired for 80,000,000 euros in cash. The waterfall calculation determines how the proceeds are distributed across the cap table.

First, all outstanding preferred shares exercise their liquidation preferences. Assume all preferred shares have 1x non-participating liquidation preferences: they receive their original investment back first, and then participate pari passu with common in any remaining proceeds.

Liquidation preference pool: Pre-seed SAFE investors originally invested 600,000 euros. Seed investors invested 2,000,000 euros.

Series A investors invested 8,000,000 euros. Total liquidation preferences: 10,600,000 euros.

Remaining proceeds after liquidation preferences: 80,000,000 - 10,600,000 = 69,400,000 euros.

All preferred shareholders have the option to convert to common and participate in the 69,400,000 on a pro-rata basis rather than taking their liquidation preference. The conversion is preferable if their pro-rata share of 69,400,000 exceeds their liquidation preference. For Series A investors (20% ownership, 3.68M out of 18.43M fully diluted shares): 20% x 69,400,000 = 13,880,000, which significantly exceeds the 8,000,000 liquidation preference. All preferred holders will convert to common and waive their liquidation preferences.

Distribution of 80,000,000 euros on an as-converted basis (all preferred treated as common):

Founder A (24.4%): 19,520,000 euros

Founder B (19.0%): 15,200,000 euros

Pre-seed investors (9.8%): 7,840,000 euros on 600,000 euros invested = 13.1x return

Seed investors (13.0%): 10,400,000 euros on 2,000,000 euros invested = 5.2x return

Series A lead (20.0%): 16,000,000 euros on 8,000,000 euros invested = 2.0x return

Option pool (13.8%): 11,040,000 euros distributed to option holders who exercise

Note that the Series A return of 2.0x is at the lower end of VC expectations on a single position. This reflects a 60-month hold at 4.4x revenue on a company that raised at a 32M euro pre-money Series A. Whether this is a success or a disappointing outcome depends on the fund's portfolio return profile and the broader vintage context.

What the Waterfall Reveals

Working through this waterfall explicitly shows several things that are easy to miss in abstract terms. First, the pre-seed SAFE investors achieved the highest return multiple (13.1x) because they took the earliest risk at the lowest effective valuation. This is structurally expected and is the economic rationale for the SAFE structure. Second, the Series A investors achieved the lowest return multiple (2.0x) on the highest absolute dollars invested. For a fund that deployed 8 million euros expecting a 3x return, this is a below-target outcome. Third, founder returns are large in absolute terms (35 million euros combined) but the 43.4% ownership represents significant dilution from 100% at inception: the dilution was the cost of access to the capital that funded the growth.

The 80 million euro exit is a successful outcome for founders by any reasonable measure. It is a marginal outcome for the Series A investor by fund return standards. This illustrates why Series A investors are looking for companies that can plausibly return 5-10x on their investment, not 2x. For a 10x return to the Series A investor, the company would need to exit at 400 million euros.

> *Model your exit waterfall before you sign every term sheet. The numbers at exit are determined by the terms you accept at each round, not by the headline valuation. A higher valuation with aggressive liquidation preferences, anti-dilution provisions, and a large forced option pool refresh may leave you with less at exit than a lower valuation with founder-friendly terms. The only way to know is to build the waterfall.*

Cap Table Red Flags

Overhang: An option pool that is too large relative to the company's stage and team is dead equity weight. If 20% of the fully diluted cap table is unissued options that will never be granted, the option pool is diluting everyone for no benefit. Review and right-size the pool at each round.

Too many small shareholders: A cap table with 30 angels, each holding 0.1-0.3%, creates administrative burden, complicates information rights management, and requires unanimous or supermajority consent for some decisions that a cleaner cap table would not. Consolidate through nominee structures or SPVs where possible.

Stacked liquidation preferences: Multiple rounds of participating preferred creates a structure where preferred investors receive their money back and then participate again in the remaining proceeds alongside common. In a modest exit scenario, this can leave founders and employees with very little. Model this explicitly: run the waterfall at 1x, 2x, 3x, and 5x the last post-money valuation and understand who gets what at each level.

Missing 83(b) elections: Founders and early employees who did not file 83(b) elections face tax liability as their equity vests. This is particularly painful in a down round scenario, where equity that has been taxed as income at a high value is now worth significantly less.

Unconverted SAFEs or notes at Series A: Multiple SAFEs with different caps and discounts, or convertible notes with accumulated interest and different discount rates, can create a conversion calculation at Series A that is genuinely confusing and sometimes unfair to one party. Clean up unconverted instruments before a priced equity round wherever possible.

021A The 25 Questions Every VC Will Ask (And the Answers That Work)

The Foundation Questions

1. Why now?

This is not a question about your timeline. It is a question about the enabling conditions that make this problem solvable today when it was not three years ago. Good answers reference technology shifts, regulatory changes, behaviour changes, or infrastructure that recently became available. Bad answers describe how hard you have been working.

2. Why you?

Domain expertise, prior experience with the problem, relevant technical depth, or access to distribution that others lack. This question is asking for a specific, differentiated answer. The worst version is a general competence argument. VCs are not looking for capable generalists. They are looking for people who are unusually well-positioned to solve this exact problem.

3. Why this market?

Size is necessary but not sufficient. VCs want to understand why this market is structurally attractive, not just large. Are incumbents slow-moving? Is there a wedge that creates a defensible entry? Is the market fragmented in a way that rewards a platform?

4. Why will you win?

Sustainable competitive advantage. Not first-mover advantage (which rarely holds) but network effects, proprietary data, switching costs, brand loyalty, or distribution advantages that compound over time. If the answer is “we will execute better”, it is not yet a competitive moat.

5. What breaks first?

This is a stress test question. VCs want to see if you have thought about failure modes. Good founders can name the top two or three things that could kill the business and explain what they are doing to mitigate each one. Founders who deflect this question signal either blind spots or intellectual dishonesty.

The Market Questions

6. How did you size this market?

Show your methodology. Bottom-up is always more credible than top-down. A $1T TAM slide citing a Gartner report is not a market size argument. The number of customers who have the problem, multiplied by what a realistic solution costs them, is.

7. Who are you displacing?

Every product displaces something, even if that something is a spreadsheet or doing nothing. Name it directly. Then explain specifically why you win in a head-to-head comparison. Avoiding this question suggests you have not done the competitive analysis.

8. What does the competitive landscape look like in three years?

VCs are investing across a time horizon of seven to ten years. They need to believe the space will not be dominated by a large incumbent, or that you have a specific reason to survive alongside one. Well-funded competitors entering your space is not automatically a negative. Sometimes it is validation.

The Traction Questions

9. What is your current traction?

Be specific and honest. Revenue, MRR, GMV, active users, pilot customers, LOIs. Whatever you have, name it. Whatever you do not have, do not euphemise. ”Strong pipeline interest” is not traction. A signed pilot with one Fortune 500 is.

10. What is your growth rate?

Month-over-month is more useful than annualised for early-stage companies. Know your last three months. Know what drove acceleration or deceleration. Know what is organic versus what came from a single event.

11. What is your churn rate?

Know both logo churn and revenue churn. Know why customers are churning. Founders who cannot explain their churn have not dug into it. Founders who have dug into it and still have high churn need to explain what they are doing about it.

12. What does your best customer look like?

VCs want to understand your ICP (ideal customer profile) from evidence, not hypothesis. Describe your best customer in specifics: size, role, problem they had before you, outcome they now have, and why they would not go back.

The Model Questions

13. Walk me through your unit economics.

Have CAC, LTV, payback period, and gross margin ready and calculated correctly. Know how each varies by channel and cohort. Know where the numbers are improving and where they are not.

14. How does the business scale?

Where are the operating leverage points? Which costs grow sublinearly with revenue? If every new dollar of revenue requires a proportional increase in headcount, the business does not scale in a VC-fundable way.

15. When do you reach profitability?

Even if profitability is years away, know the path and the key assumptions that determine it. This question is often probing whether you have thought about the business as a going concern, not just a funding-dependent entity.

16. What is your gross margin today, and where does it go?

Gross margin is a business model question, not just an accounting question. The answer signals whether this is a software business, a services business, or something in between. Know your current margin, know the structural drivers, and know what it looks like at scale.

The Team Questions

17. Why is this the right team for this problem?

Not a general competence argument. A specific unfair advantage argument. What does this team know, have access to, or have done before that makes them better positioned than anyone else to solve this?

18. What is missing from the team today?

A founder who knows what they are missing is far more credible than one who believes they have everything covered. Name the gap, name the hire you are planning to make, and explain how you will make sure that person is better at that function than you could be.

19. How do you make decisions when you disagree?

For co-founder teams. VCs have seen co-founder breakups destroy companies. They are looking for evidence of a functioning process, not just mutual admiration. Talk about how you have navigated disagreements before.

The Round Questions

20. How much are you raising and why?

The amount should be tied to specific milestones and a specific runway. Not “18-24 months of runway” as a generic answer, but ”$3M takes us to $1M ARR with 20 customers, which positions us for a Series A at a $20M pre-money.”

21. What milestones does this round fund?

Be specific. Define the milestones in terms that VCs will recognise as Series A proof points: ARR, customer count, NRR, product completion,

or market entry. Vague milestones signal you have not thought through the deployment of the capital.

22. What is your current burn rate and runway?

Know this cold. Investors will check it against your ask. If your monthly burn is $80K and you are raising $3M, that is approximately 37 months of runway. If you say you need 18 months to hit milestones, why are you raising for 37?

23. Who else is in the round?

Social proof matters in fundraising. If you have strong angels or a co-investor already committed, name them. If the round is empty, be honest about where you are in the process. Sophisticated investors will find out; trying to hide the state of the round destroys trust faster than an honest answer.

The Big Questions

24. What does the exit look like?

IPO is rarely the honest answer for pre-Series B companies. M&A is more likely. Name the category of acquirers: strategic buyers who would benefit from your distribution or technology, or financial buyers if the revenue profile supports it. Have a view on the likely valuation range and what drives it.

25. What would make you change your mind about this business?

This is the most revealing question of the 25. Founders who have a clear answer to this have a falsifiable thesis. Founders who do not have an answer, or give an unconvincing one, may be running on conviction without evidence. Investors trust founders who can articulate their own uncertainty.

These 25 questions are not a script. They are a framework. Every investor asks them in a different order with different emphasis. The founders who raise are not the ones who have rehearsed perfect answers; they are the ones who have thought hard enough about their business that an honest answer to each of these questions is already

available. Preparation is what makes a pitch meeting a conversation rather than an interrogation.

021B Meridian: A Seed-to-Series A Case Study

Company Overview

Meridian automates supplier invoice reconciliation for mid-market retailers. The problem: procurement teams at 50–500 employee companies spend 3–5 hours per week manually matching supplier invoices to purchase orders in their ERP. The solution: an integration layer that sits between the retailer's ERP and their supplier network and handles reconciliation automatically, flagging exceptions only.

Founding team: two co-founders. One with eight years of operations experience at a retail chain (domain expert). One with a background in fintech engineering (technical founder). They met at a prior employer, worked together for three years, and share a co-founder agreement that vests over four years with a one-year cliff.

Phase 1: Pre-Seed Round

Meridian raises a pre-seed round of 600,000 euros from two angel investors. One is an operator angel, a former CFO who led a finance transformation project at a large UK retailer. The other is a syndicate of three angels who invest together in B2B SaaS.

Pre-money valuation at pre-seed: 2.4M euros. Post-money: 3M euros. Implied dilution: 20%. The round is structured as an uncapped convertible note with a 20% discount, converting at the seed round.

Why an uncapped note: Meridian does not yet have product-market fit evidence. Valuation is speculative. The operator angel suggests a cap structure would create a ceiling that could complicate the seed round negotiations. Both parties agree to a note with a discount and a maximum conversion share.

Capital deployment plan: 600K euros funds 18 months. The plan is to spend 380K on two hires (one engineer, one part-time customer

success contractor), 120K on cloud infrastructure and tooling, and 100K on customer acquisition experiments. The target at the end of the pre-seed runway: three paying customers, 8K euros MRR, and enough data to build a credible cohort model.

Actual outcome at month 18: four customers, 11.2K euros MRR, average contract value of 2,800 euros per month. Two customers expanded from a trial tier. Average gross margin: 74%. Churn in the period: one customer who was acquired by a larger company that had its own solution.

Phase 2: Building the Seed Pitch

Meridian begins building the seed pitch at month 15 of the pre-seed runway, not month 18. This is a deliberate decision: the first meetings are learning meetings. By month 15, the team wants to be in conversations before they are fundraising out of desperation.

The seed ask: 2.5M euros at a 9M euros pre-money valuation (post-money 11.5M euros). The 20% discount on the convertible note means the angels convert at 7.2M euros effective pre-money. The new money goes in at the same 9M euros pre-money as new investors.

The financial model at seed includes: a 36-month forecast with monthly granularity, a driver-based revenue model tied to new customer acquisition by channel, a headcount plan naming seven hires over 24 months, a gross margin build with COGS broken out, and three scenarios (base, conservative, and stretched).

Key assumptions in the model, all sourced: New logo acquisition rate of 1.5 per month in months 1–12 (based on current conversion rate from pilot to paid). Average contract value growth from 2,800 euros to 3,400 euros by month 24 (based on pricing research and competitor positioning). Gross margin steady at 73–75% (benchmarked against vertical SaaS comparables in fintech infrastructure). Monthly churn rate of 1.2% logo churn (based on observed 14-month average customer life to that point).

The Seed Pitch Deck Structure

Slide 1: Problem. Retailers spend 40+ hours per month reconciling invoices manually. Finance teams at 100-person companies average 3.2 FTE-days per week on this task alone.

Slide 2: Solution. One paragraph, one product screenshot. The headline: reconciliation fully automated, exceptions-only workflow.

Slide 3: Market. Bottom-up: 120,000 mid-market retailers in the UK and DACH region, 40% of which use an ERP compatible with Meridian's integration. Total addressable market of 1.35B euros at current pricing.

Slide 4: Traction. Four customers, 11.2K euros MRR, 74% gross margin, zero logo churn in the last six months.

Slide 5: Business model. Per-seat SaaS with a minimum contract value. Implementation included for the first year.

Slide 6: Unit economics. CAC: 4,200 euros blended. LTV: 38,000 euros at current contract value and observed churn. LTV:CAC ratio: 9.1x. Payback period: 15 months.

Slide 7: Go-to-market. Direct sales to finance directors and heads of procurement. Channel partnership in development with one mid-market ERP provider.

Slide 8: Team. Two co-founders, one senior hire committed to joining.

Slide 9: Financials. 36-month revenue projection and burn profile. The ask: 2.5M euros. Milestones: 50K euros MRR, 15 paying customers, ERP partner agreement signed.

Slide 10: Ask. 2.5M euros seed at 9M euros pre-money.

Phase 3: The Seed Round Process

Meridian runs a structured process over 10 weeks. Week 1–2: warm introductions to 15 seed funds via the operator angel and a few LinkedIn connections from the technical founder's network. Week 3–4: first meetings with 11 of those 15. Week 5–6: second meetings with five. Week 7–8: three partner meetings. Week 9: term sheets from two investors. Week 10: negotiations and close.

Common objections encountered: The market is fragmented, so winning scale requires a land-and-expand strategy that has not been proven yet. The team lacks a sales hire. The 9M euros pre-money is high for 11.2K euros MRR (an approximately 800x revenue multiple on MRR, or 66x ARR, which is at the upper end for a pre-scale SaaS business in the current market).

How each objection was handled: On the land-and-expand question, the team pointed to two customers who had already expanded from pilot to multi-seat within six months. On the sales hire, they offered to make this the first deployment of seed capital. On the valuation, the founders did not move significantly but did include a milestone-based warrant structure that gave investors additional upside if ARR targets were not met by month 18.

Outcome: The seed round closes with a lead investor (a 15M euros B2B SaaS-focused seed fund) taking 1.8M euros at 9M euros pre-money and a co-investor taking 700K euros on the same terms. Total round: 2.5M euros. Lead investor takes one board seat. Co-investor takes observer rights.

Cap Table: Seed Stage

Founders (combined): 72.1% post-seed (pre-seed dilution: 5% option pool top-up, 15% from seed round, conversion of pre-seed note). Pre-seed angels: 11.2% (converted at 7.2M euros effective). Seed lead: 13.6%. Seed co-investor: 5.3%. Employee option pool: 7.8% (refreshed at seed). Total: 110% before buyback of pre-pool reserved shares, normalised to 100% on fully diluted basis.

Phase 4: Seed-to-Series A Execution

The 18 months between seed close and Series A kick-off are the most operationally intense period of the company's life. The financial model is updated monthly. Actuals are tracked against the model every four weeks and presented to the board.

At month 12 post-seed, Meridian is tracking ahead on customers (22 vs. 18 in the base case) but behind on ACV (3,100 euros vs. 3,400 euros

in the model). The team updates the model assumptions, creates a variance explanation, and shares it with investors before they ask.

At month 18 post-seed, the metrics are: MRR 52,400 euros. ARR 628,800 euros. Customer count: 28 paying customers. Gross margin: 76%. Net revenue retention: 112% (expansion from existing accounts exceeding churn). MoM MRR growth rate: 8.2% average over the last six months. Burn rate: 78,000 euros per month. Runway at current burn: 7 months.

The runway number is the trigger. Seven months of runway is tight for a Series A process, which typically takes four to six months from first meeting to close. Meridian begins the Series A process immediately.

Phase 5: Series A

The Series A ask: 8M euros at a 32M euros pre-money valuation. The 32M euros is derived from a revenue multiple of approximately 50x ARR, which is consistent with the lower end of comparable B2B SaaS multiples in the European market for a company with 112% NRR and 28 customers.

Key differences in the Series A financial model versus seed: The revenue forecast now has 24 months of actuals to anchor assumptions. The model includes a sales capacity model (how many AEs, how many deals per quarter, at what ACV). Gross margin at scale is modelled separately from current-state gross margin. The three scenarios now include a path to breakeven (month 38 in the base case). The model includes a detailed hiring plan for 18 new people over 24 months.

The Series A closes in month 24 post-seed. Lead investor: a 120M euros European growth fund with several portfolio companies in the B2B SaaS vertical. Round size: 8M euros. Valuation: 32M euros pre-money. Post-money: 40M euros. Lead investor: 20% ownership at entry.

What This Case Study Illustrates

The model must survive contact with reality. Meridian's ACV came in below target. The team adjusted the model, communicated the variance, and explained the implications. This is what investors mean when they say they want founders who “know their numbers”. They do not mean founders whose numbers are perfect. They mean founders who know why the numbers are what they are.

Round timing is a strategic decision. Starting the seed process at month 15 of an 18-month runway, and starting the Series A at 7 months of runway, are both deliberate choices. Tight runway creates timeline pressure that rarely helps negotiation. Build in time.

The model evolves between rounds. The seed model and the Series A model are structurally different documents. The seed model proves you understand unit economics. The Series A model proves you understand how to scale.

Cap table decisions compound. The decision to use a convertible note at pre-seed, the option pool refresh at seed, and the 20% ownership concession to the Series A lead all shape who owns what at exit. These decisions are made one at a time, but they create a structure that cannot be easily unwound later.

> *Key insight: The founders who raise efficiently are not the luckiest or the most well-connected. They are the ones who understand what they are asking for, who they are asking it from, and what the money needs to produce for everyone who provided it. Every concept in this book lives inside the Meridian story somewhere.*

021C Ten Fundraising Mistakes Founders Make (And How to Not Make Them)

Mistake 1: Raising Too Late

The single most common fundraising mistake is starting too late. Starting too late means beginning investor conversations when runway is under six months. At that point, investors know you are raising from desperation, not optionality. They negotiate harder, move slower, and often pass because the risk of a broken process is higher.

The fix: begin investor relationship-building when you have 12-plus months of runway. Not the formal process, but the first conversations, the warm intros, the LinkedIn visibility. By the time you formally kick off a round, investors should already know who you are.

Mistake 2: Targeting the Wrong Investors

Sending the same deck to a $500M growth fund and a $20M seed fund is not a strategy; it is spam. Growth funds do not do seed deals. Seed funds do not lead Series B rounds. Investors whose fund size, stage, and sector focus do not match your current state will not invest, and the meetings are a waste of time for everyone.

The fix: research fund size, typical check size, stage, sector focus, and fund cycle before reaching out. The 30 minutes of research per investor saves four to six hours of meetings that go nowhere.

Mistake 3: A Weak Traction Narrative

Founders sometimes raise without traction and succeed; but they do so by being honest about what they have. The mistake is not having early-stage traction. The mistake is presenting ambiguous or inflated traction signals: ”strong pipeline”, “significant interest from enterprise

customers", "active conversations with major retailers". These phrases mean nothing and experienced investors know it.

The fix: state what you have precisely. Three paying customers at 2,000 euros per month is a clearer and more trustworthy traction statement than "growing rapidly with strong market validation". Specificity signals credibility.

Mistake 4: Unclear Market Sizing

A TAM slide that says "the global XYZ market is $2.3T according to Grand View Research" is not a market sizing argument. It is a copy-paste. VCs have seen it thousands of times and it signals that the founder has not done original analysis.

The fix: build a bottom-up market size from real data. Number of potential customers, multiplied by annual contract value, gives you a serviceable addressable market. Do that. Show your work. The number will almost always be smaller than the top-down figure, and that is fine. Credibility is worth more than size.

Mistake 5: Poor Cap Table Hygiene

A cap table with 15 small angels from the early days, no option pool, a 60% pre-money discount on a convertible note, or a founding team that no longer has any economic incentive because they diluted too heavily at pre-seed are all structural problems that can block a round. Investors will review the cap table and model the scenario for every stage of future dilution.

The fix: clean up the cap table before you fundraise. Consolidate angels where possible. Refresh the option pool before the round, not after (the economic effect on founders is the same, but the signalling is cleaner). Make sure founding equity economics are still motivating for the people doing the work.

Mistake 6: Over-Engineered Financial Models

A 14-tab model with 600 formulas, Monte Carlo simulations, and a DCF valuation output does not impress investors. It confuses them. A model that is too complex to navigate quickly signals that the founder does not understand what the model is for.

The fix: keep the investor-facing model to five or six tabs: assumptions, revenue, headcount, P&L, cash flow, and summary. Make every number traceable to the assumptions tab. Complexity for its own sake is a signal of insecurity, not sophistication.

Mistake 7: Pitching Without a Process

Sending decks to investors one at a time, waiting for a response, then sending to the next person is not a fundraising process. It is a slow-motion rejection machine. Investors want to see other investors interested. A competitive dynamic, even a manufactured one, creates urgency.

The fix: run a batched process. Identify 20 to 30 target investors. Reach out to all of them in the same two-week window. Sequence conversations so that first meetings cluster, partner meetings cluster, and term sheets are expected in the same period. Momentum is a fundraising tool.

Mistake 8: Negotiating Against Yourself

Founders sometimes volunteer concessions that investors have not asked for. “We could probably do this at a lower valuation”, ”we are flexible on board composition”, “we can add more protective provisions”. These concessions are not appreciated; they are exploited. They signal anxiety, and anxiety signals weakness.

The fix: do not volunteer information that weakens your position. Answer what is asked. Let investors push on the terms they care about. Be willing to move on things that matter to them and do not damage you. But start from your position, not from compromise.

Mistake 9: Ignoring the Second Meeting

First meetings are for qualification. The second meeting is where the investment decision is made. Most founders prepare heavily for the first meeting and think the second meeting will take care of itself. It will not. The second meeting is typically where the VC probes the areas they were uncertain about after meeting one.

The fix: after every first meeting, ask the investor what would make them confident enough to bring this to the partnership. Then build the second meeting around answering that question directly. Do not wait to be told what to prepare.

Mistake 10: Not Managing the Process After a Term Sheet

Receiving a term sheet feels like the finish line. It is not. It is the start of due diligence, negotiation, and legal documentation, any of which can collapse the deal. Founders who celebrate the term sheet and disengage from the process make this phase take twice as long and twice as painfully.

The fix: once you have a term sheet, compress the timeline. Get a lawyer who has closed venture deals before. Have your data room complete before due diligence starts. Be responsive within hours, not days. Every day between term sheet and close is an opportunity for the investor to change their mind.

Ten mistakes: starting too late, targeting the wrong investors, vague traction claims, top-down market sizing, a messy cap table, an over-complicated model, pitching without a process, negotiating against yourself, underpreparing the second meeting, and disengaging after the term sheet. None of these require genius to avoid. They require preparation and discipline. The founders who raise efficiently are, almost without exception, the ones who have avoided most of these ten.

021 How to Build a Financial Model for Your Startup (2026 Guide)

What is a Startup Financial Model?

A startup financial model is a quantitative representation of your business: how it generates revenue, what it costs to operate, and how much cash it needs to reach key milestones. A good model is not a static spreadsheet — it is a living decision-making tool that connects your assumptions directly to your outcomes.

Typical planning horizon	3–5 years
Core outputs	P&L, cash flow, balance sheet
Most common investor ask 3+	Monthly for Year 1–2, quarterly for Year
#1 reason models fail drivers	Inputs not connected to real business
Standard scenario count aggressive)	Optionally 3 (base, conservative,

Why Does a Financial Model Matter for Startups?

Investors do not fund ideas — they fund evidence that a team understands the economics of their business. A financial model is the clearest signal of that understanding.

According to DocSend's 2024 Startup Fundraising Report, investors spend an average of 23 minutes reviewing a startup's financial documents during due diligence — more time than any other section of the pitch. McKinsey research on high-growth companies found that startups with documented financial models had a 35% higher Series A conversion rate from seed than those without.

I have seen this dynamic firsthand across multiple fundraising processes

rounds backed by top-tier European and US venture funds. The deals that moved fastest were the ones where the founder could answer any investor question by pointing directly to a cell in their model.

What Are the Core Components of a Startup Financial Model?

A complete startup financial model has five interconnected components. Missing any one of them creates blind spots that experienced investors will find immediately.

1. Revenue Model — How does money come in? This must be

driver-based. Eg if it is B2B sales based: How many opportunities can a salesperson review, how many leads can they qualify, how many wins can they bring in each month? That determines the number of customers whcih can then be multiplied by average contract value × churn rate, an assumptions that says "revenue grows 20% per year" is not acceptable. **2. Cost Structure** — Split between COGS (what it costs to deliver your product) and OpEx (what it costs to run the company). Gross margin is one of the first numbers a Series A investor checks.

3. Headcount Plan — Personnel is typically 60–80% of a startup's burn. Model headcount by role, start date, salary, and associated costs (benefits, payroll overhead). Vague headcount equals untrustworthy burn. If hiring globally beware of varied overhead, a Romanian employee has less than 10% overhead whereas a French one can exceed 45%. If using Employer of Record, this also has to be included

4. Cash Flow Statement — Profit does not equal cash. A startup can show accounting profit and run out of money. The cash flow statement shows when you actually receive and spend money, which determines your real runway. The ratio of average days receiveable can give you a good sense on how long it takes for you to get paid each invoice on average. **5. Balance Sheet** — Often skipped at early stages. Investors use it to verify the model is internally consistent. If your balance sheet doesn't balance, everything else is

suspect. It is of secondary priority in my opinion as the cash flow and P&L have to be healthy first before you should worry about the balance sheet.

> *Key insight: The most common failure in startup financial models is building a revenue forecast disconnected from operational drivers. Revenue should be the output of modeling customers, conversion rates, pricing, and retention — not an input. I have rebuilt models for funded startups where the entire revenue tab was a growth rate assumption pasted over five years. That is not a model — it is a spreadsheet with hope in it.*

How to Build a Driver-Based Revenue Model

Driver-based modeling means your revenue is derived from the underlying mechanics of your business, not assumed. This applies whether you are a marketplace, a SaaS platform, or a services business.

Identify your primary growth driver. For SaaS: new logo acquisition rate. For marketplace: GMV and take rate. For services: billable hours or active client count. Every other number flows from this single driver.

Build your funnel top-down. Start with market reach (impressions, leads, qualified pipeline), apply conversion rates at each stage, and arrive at new customers per period. This structure lets you test the sensitivity of revenue to each lever independently. Use industry averages as inputs to avoid getting questioned about these and not having a solid answer.

Layer in retention. New revenue minus churned revenue equals net revenue. A model that shows revenue growing without accounting for churn is wrong by construction. Model monthly or annual churn explicitly and calculate net revenue retention (NRR) — investors will ask for it. Working with Indeed Flex (previously Syft), one of Europe's leading flexible staffing marketplaces, reinforced a principle I apply to every marketplace model: GMV growth and revenue growth are not the same number. Take rate compression, geographic mix shift, and

worker-to-client ratio all affect revenue independently of GMV. Model each driver separately or your forecast will be wrong in three different directions at once.

1. **Start with the unit** — one customer, one transaction, one contract. What do the economics of a single unit look like?
2. **Build the cohort** — group customers by acquisition month. Track each cohort's revenue and cost over time. This reveals LTV/CAC naturally rather than as a calculation bolted on at the end.
3. **Aggregate across cohorts** — sum all active cohorts per period to get total revenue. This is the most robust revenue forecasting method for subscription and recurring revenue businesses.
4. **Sense-check against top-down** — does your bottom-up model imply a market share that is realistic? If you are projecting $50M in Year 3 and your TAM is $200M, that is a 25% market share — possible but it requires a compelling narrative.

How to Model Unit Economics

Unit economics are the financial building blocks of your business. They answer one question: is acquiring one customer profitable? Without this, you cannot answer whether you should grow faster or slower. **Customer Acquisition Cost (CAC):** Total sales and marketing spend in a period ÷ new customers acquired in that period. Model CAC by channel, not as a blended average — blended CAC hides which channels are producing and which are destroying value.

Customer Lifetime Value (LTV): Average revenue per customer per month × gross margin % ÷ monthly churn rate. This gives the net present value of a customer relationship. If you have growing cohorts this formula fails as the value nears infinity the closer you get to 0% churn rate. If that's your case feel free to reach out and I can advise how to deal with it, and I am quite likely going to be willing to invest in your business assuming CAC is not bigger than the 1st year worth of LTV. **LTV:CAC ratio:** The standard benchmark is 3:1 minimum for a healthy SaaS or marketplace business. Below 3:1, you are likely

acquiring customers at a loss when accounting for margin. Above 10:1 may indicate you are underinvesting in growth.

CAC Payback Period: CAC ÷ (monthly revenue per customer × gross margin %). This is how many months until a new customer has paid back their acquisition cost. Best-in-class SaaS companies target under 12 months. Most early-stage startups run 18–36 months — which is acceptable, as long as churn is low and the trajectory is improving. Modeling the unit economics for Simplifii and Haus across different market conditions — including multi-currency operations across the UK, US, and UAE — made clear that unit economics look very different depending on the market. CAC in Dubai for an early-stage marketplace can be 2–3× the equivalent in London, while LTV potential is often higher. Multi-market models must run unit economics independently per geography before aggregating.

> *Key insight: Investors funding a pre-profitable startup are making a bet that LTV:CAC improves with scale. Your model must show how and when — through pricing power, channel efficiency, or retention improvements — or the bet is unsubstantiated. Every VC I have worked with, from Creandum to B2Ventures, runs this analysis themselves within the first hours of receiving a model.*

What Scenarios Should a Startup Model Include?

Every investor-ready financial model needs at least three scenarios. Single-scenario models signal that the founder has not stress-tested their own assumptions — a significant red flag in a competitive fundraising process.

Conservative	Base Case	Aggressive
Revenue growth	Lowest realistic	Most likely
CAC trend	Flat or increasing Modest	Significant improvement
Churn rate	Higher than	At current rate

Hiring plan	Lean	Planned
Best used for	Stress testing	Investor

> If your business cannot survive on conservative assumptions without another raise, investors will see it and your negotiating position deteriorates.
> Across the raises I have worked on, the scenario that generates the most investor questions is always the conservative case — specifically:
> ”At what point does this scenario require additional capital, and what would trigger you to shift to this plan?“ Founders who have not built this scenario cannot answer that question.

How to Structure a Financial Model for Investor Due Diligence

A model built for internal use and a model built for investors are different documents. Here is what sophisticated investors from funds like Profounders and Boost Capital actually examine during diligence:
Assumption transparency. Every key input — growth rate, churn, gross margin, CAC — should be in a clearly labelled inputs tab with a source or rationale. "Industry benchmark" is not a source. "Median gross margin for B2B SaaS with ACV >$10k per OpenView's 2024 SaaS Benchmarks report" is.

Monthly granularity for Year 1–2. Annual totals hide cash gaps. A company can look profitable annually while going cash-negative for three consecutive months. Every experienced investor checks monthly, because they have been caught by annual averages before.

A clear bridge to the raise. The model should make obvious: how much you are raising, what milestones that capital funds, and what the business looks like at the end of the runway. This is not narrative — it is structural. The cash flow statement should show the raise as an inflow

and the spending plan as outflows, with the resulting runway visible in a single row.

Sensitivity analysis. Show what happens if your top two assumptions are wrong. If revenue is 20% below base case, when do you run out of cash? Investors will run this themselves — having it pre-built signals confidence, not weakness.

Common Financial Modeling Mistakes Founders Make

Straight-line revenue growth: Writing "revenue grows 15% per month" without connecting it to customers, pricing, or retention. It is not a model — it is a wish with formatting.

Ignoring COGS: Many first-time founders model gross revenue without cost of goods. Gross margin is one of the most scrutinised metrics in any raise. Skipping it destroys credibility in the first five minutes of a diligence call.

Circular references and broken formulas: A model with formula errors is worse than no model. It tells investors the founder has not pressure-tested their own numbers. Always audit before sharing.

Single scenario only: Shows you have not stress-tested your business. Every experienced investor immediately asks for the downside case.

Headcount as a lump sum: "Salaries: $800k" is not a headcount plan. Model each role, start date, and loaded cost (salary + benefits + payroll taxes, typically 1.2–1.3× base salary).

Confusing cash and revenue: Deferred revenue, payment timing, and accounts receivable all create gaps between when you recognise revenue and when cash arrives. Especially critical for enterprise deals with net-60 or net-90 payment terms.

Single-currency models for multi-market businesses: If you operate across the UK, US, and UAE simultaneously, a single-currency model masks FX exposure, local CAC differences, and margin variation by market. This is an error I see repeatedly in cross-border startups.

A strong startup financial model is driver-based, scenario-tested, and built with investor due diligence in mind from day one. Revenue must flow from real business drivers — customers, pricing, retention — not from assumed growth rates. Unit economics (CAC, LTV, payback period) are not optional metrics to add later; they are the foundation of every credible growth narrative. Model monthly for your near-term horizon, include at least three scenarios, and ensure your assumptions are sourced and auditable. Having worked across $20M+ in raised capital and $100M+ in exits across UK, US, and Dubai-based marketplaces and SaaS businesses, the consistent truth is this: the founders who raise faster, at better terms, are the ones who understood their model well enough to defend every number in it.

My Takeaways from Raising 5 Rounds Across 7 VCs And One 9 figure Exit

What Does It Take to Build a Model That Closes Funding Rounds?

A financial model that closes funding rounds is not the most sophisticated spreadsheet in the room, it is the most honest one. It translates a founder's vision into a set of numbers that are both aspirational and defensible, and it forces you to find the risks before the investors do.

The rounds behind this article:

Seed Round, Series A ,(lead by CreandumVC) Venture Debt Debt facility, Convertible Note Bridge instrument, Exit to Indeed (corporate acquirer) **Total outcome$xxxM+ exit, 55x revenue multiple**

How I Got Here: The Context

I am not going to pretend I knew what I was doing the first time I built a fundraising model. I made assumptions I couldn't defend, built revenue lines I believed in but couldn't source, and presented scenarios that looked compelling until an investor asked the third question. What

changed was not the tool, it was the discipline. Over years of working with Indeed Flex (formerly Syft), a two-sided labour marketplace that scaled from startup to acquisition, I built and rebuilt the financial model through every stage of the company's life: first capital, growth rounds, venture debt, a nearly-done Series B, and ultimately the diligence process for the exit.

The investors who backed us, including Creandum, Profounders, were not easy rooms. Each of them had seen hundreds of models. The three principles below are what I believe made ours different.

Principle 1: Let the Data Show You the Way

Most founders build a financial model to convince investors. I build financial models to *find out what's true* about a business, and then use that truth to make the case.

This distinction matters enormously. When you treat a model as a persuasion tool, you start with the number you want and work backwards to justify it. When you treat it as a discovery tool, you start with the real drivers of your business, customers, conversion rates, churn, unit economics, and let the numbers tell you what is possible.

The practical result: the model becomes a map of your risks and your opportunities simultaneously. When I translate a founder's vision into a full model, I can identify which assumptions are fragile (and therefore need to be stress-tested or mitigated before the investor sees them) and which are conservative (and therefore represent upside in the story).

> **Key insight: The best fundraising models are not optimistic, they are accurate, with an honest upside case layered on top. Investors who have seen hundreds of models can tell the difference in seconds.**

This principle also means your model becomes a management tool, not just a fundraising document. At Indeed Flex, the model we used to raise was the same model we used to track performance against plan. When you build this way, variance analysis is not a threat, it is a conversation.

Principle 2: Always Have a Story Behind the Numbers

A number without a story is just a number. A number with a story is a conviction.

In every model I have built, every material assumption, every growth rate, every churn figure, every CAC estimate, must have a source or a rationale attached to it. Not buried in a footnote. Stated clearly in an assumptions tab, so that when an investor asks "why is your CAC $320 and not $500?", the answer is not "we think that's achievable" , it is "here is the channel mix, here are the comparable benchmarks from OpenView's SaaS report, and here is our current pilot data." This does not mean your numbers need to be conservative. Aspirational numbers are fine; required, even, for early-stage fundraising. What they cannot be is undefended.

The rule of thumb I have worked with across multiple rounds: hitting 80% of your base case forecast is considered good performance in most investor conversations. The model does not need to be right, it needs to demonstrate that you understand your business well enough that when you miss, you can explain exactly why, and what you are doing about it.

> **Key insight: Investors are not funding your forecast, they are funding your understanding of your own business. The story behind the numbers is the proof of that understanding.**

Principle 3: I Know Nothing

This is the most important principle and the hardest to maintain when you have been working on a business for two years and you believe, deeply, that you are right.

Questioning everything means: every assumption in the model gets challenged before an investor gets to challenge it. Every growth rate gets tested against comparable companies. Every cost estimate gets sense-checked against actual quotes. Every unit economics figure gets

triangulated against your own data, industry benchmarks, and the logical implications of your business model.

In practice, this looks like running a pre-mortem on your own model. Ask: what are the three assumptions in this model that, if wrong, would make the business unviable? Then either fix those assumptions to be more defensible, or build a clear narrative for why you believe them despite the risk.

I have rebuilt large sections of financial models during live fundraising processes because this questioning process found something that didn't hold up. It is not comfortable. It is significantly better than having an investor find it for you.

The questions I ask every model before it leaves my hands:

1. Does revenue come from real business drivers or from an assumed growth rate?
2. Can I defend every material assumption with data, a benchmark, or high conviction logic?
3. Does the conservative scenario show a path to survival without another raise?
4. Is there a circular reference or formula error anywhere in this file?
5. If the CEO reads this model for the first time, does it reflect their vision accurately?

What This Looks Like Across a Full Funding Journey

The application of these three principles looks different at each stage: Stage: What the model must prove Hardest question

Pre-seed: You understand the unit economics hypothesis "Why will CAC come down?"

Seed: You can acquire customers and retain them "What does Month 12 churn look like?"

Series A: The unit economics work at current scale "Show me LTV:CAC by channel"

Venture debt: You have predictable cash flows "What's your worst-case runway?"

Growth / Series B: You have a credible path to profitability "When do you break even?"

Exit / M&A DD: Your actuals validate your historical forecasts "Explain this variance"

Each stage requires a different model emphasis, but the same three principles apply throughout. The data leads. The story is sourced. You question everything.

The Honest Truth About Needing To Be A "Genius"

After the news from a recent client, I found myself asking whether I was just lucky or whether I had actually done something right across these rounds.

The honest answer is: the principles are not complicated. What is rare is the discipline to apply all three of them, every time, under the pressure of a fundraising process when the temptation to shade a number or leave an assumption undefended is very real.

The founders who raise faster, at better terms, are not smarter. They are more prepared. Preparation, in startup finance, looks like a model built on real drivers, defended by a sourced story, and stress-tested by someone willing to question their own assumptions before anyone else gets the chance.

That is what this series is about.

Frequently Asked Questions

How many rounds is typical for a startup before exit?

There is no standard. Indeed Flex (formerly Syft) went through two seed rounds, a Series A, venture debt, and a convertible note before the exit to Indeed—six financing events in total. The number depends on capital efficiency, growth rate, and market conditions. What matters more than the count is that each round has a clear purpose: specific milestones the capital will fund.

What made the Indeed Flex exit possible at a 55x revenue multiple?

Several factors—but the financial model's contribution was specific: it gave the acquirer confidence in the quality of the revenue base, and it enabled a robust defence of client cohort value when the acquirer attempted to cut the valuation. The model was not just a forecast. It was a structured representation of the business that held up under six months of diligence.

Is financial modeling experience necessary for a founder to raise?

You do not need to build the model yourself—but you must understand it deeply enough to defend every number in a partner meeting. Founders who outsource the model entirely and cannot answer questions about their own assumptions lose credibility fast. The model is a test of how well you understand your own business.

What is the most common reason a model fails in investor diligence?

Revenue disconnected from real drivers. Investors who have seen hundreds of models immediately identify the difference between a revenue line built from business mechanics (customers × conversion × ACV × retention) and a revenue line that is a growth rate applied to last year's number. The former invites questions. The latter ends conversations.

Summary

Three principles have held across every funding round and every model I have built: let the data lead you to the truth about your business, build a sourced story behind every material assumption, and question everything before anyone else can. These are not sophisticated ideas—they are disciplines that most models violate because the pressure to look good overrides the pressure to be accurate. A model built on these principles does not just raise money. It becomes the

operating foundation of the company—the document that ties the CEO's vision to the CFO's reality, and gives investors confidence that the team in front of them knows exactly what they are building.

022 Driver-Based Modeling: Build Revenue From Reality, Not From Hope

What Is Driver-Based Modeling?

A driver-based model connects every financial output to a measurable business input. Revenue is not a number you choose. It is a number that falls out when you define how your business actually works. Think of it this way. A traditional financial forecast says: "Revenue will be $2M next year." A driver-based model says: "We will generate 10,000 leads per month at a 3% conversion rate, producing 300 new customers monthly. At an average contract value of $550, that produces $165,000 in monthly new revenue. Layer in 5% monthly churn on the existing base, and the model calculates what total revenue looks like." The difference is not cosmetic. It is structural. The second version tells you exactly which levers to pull, which assumptions to challenge, and which numbers need to hold for the business to work.

Traditional Forecast	Driver-Based Model
Revenue = $2M (assumed)	Revenue = f(leads, conversion, ACV, retention)
Costs = % of revenue	Costs = headcount + vendor contracts + unit costs
Growth = 15% MoM	Growth = output of improving individual drivers
Untestable	Every assumption is independently testable
Breaks under investor questioning	Invites investor questioning productively

Why Investors Require Driver-Based Models

An investor evaluating your startup needs to answer one question before writing a check: "Do these founders understand how their business creates value?" A revenue line that says "grows 20% per month" answers nothing. It is a wish formatted as a number.

When Creandum evaluated the Indeed Flex model during our Series A process, the first thing they did was trace every revenue line back to its operational driver. They wanted to see the funnel: how many employers were we reaching, what was the conversion rate from first shift to recurring client, what was the average gross margin per placement. They did not ask "what is your revenue projection." They asked "walk me through how a customer becomes revenue."

This is how every sophisticated fund operates. Profounders, B2Ventures, Boost Capital, they all run the same mental test: can the founder explain the chain from market activity to cash in the bank? If the answer lives in the model as a connected set of drivers, the conversation goes well. If the answer is "we assumed 15% growth," the meeting is effectively over.

> **Key insight: Investors are not checking whether your projections are right. They are checking whether you understand the mechanics well enough that when the projections are wrong, you can diagnose exactly why and adjust.**

The Five Categories of Business Drivers

Every startup, regardless of business model, has drivers that fall into five categories. Identifying yours is the first step to building a proper model.

1. Acquisition Drivers

These are the inputs that determine how many new customers you bring in. For a B2B SaaS company, this might be monthly leads, demo booking rate, and close rate. For a marketplace like Indeed Flex, this

was inbound employer inquiries, outbound sales calls, and conversion from first shift posted to active recurring client. For a direct-to-consumer brand, it could be monthly site visitors, add-to-cart rate, and checkout completion rate.

The key discipline: break the funnel into at least three stages. Top (awareness or reach), middle (engagement or qualification), and bottom (conversion). Single-stage acquisition assumptions like "we will get 100 new customers per month" are not driver-based. They are just smaller wishes.

2. Monetization Drivers

Once you have a customer, how much do they pay? This is not just "price." It includes average contract value, number of seats or units per customer, billing frequency, expansion revenue potential, and discount rates. For Indeed Flex, the monetization driver was hours worked per client per week multiplied by the gross margin per hour. For SaaS, it is typically seats multiplied by price per seat, with an expansion multiplier for upsell.

Model these separately from acquisition. A common mistake is to combine "new customers multiplied by average revenue" into a single line. That hides whether growth is coming from more customers or more revenue per customer, and those two trajectories require completely different strategies.

3. Retention Drivers

Churn destroys forecasts faster than any other variable. Model it explicitly: monthly logo churn (percentage of customers who cancel), revenue churn (percentage of MRR lost), and net revenue retention (which includes expansion). These three numbers tell very different stories, and sophisticated investors will ask for all three.

At Indeed Flex, we tracked client retention both by logo (did the employer come back next month?) and by revenue (how much did they spend compared to last month?). Logo retention was strong, but

revenue retention was even stronger because repeat clients increased their usage over time. That distinction was material during the exit diligence.

4. Cost Drivers

Costs are not percentages of revenue. They are driven by specific operational decisions: hiring plan, vendor contracts, infrastructure usage, customer support volume, marketing spend by channel. Each cost category should connect to an operational metric.

Headcount is the biggest one for most startups: 60-80% of burn is salaries. Model it role by role, not as a lump sum. Infrastructure costs scale with usage (users, API calls, data storage). Customer support costs scale with customer count and ticket volume. Marketing spend is a choice variable that drives your acquisition metrics.

5. Timing Drivers

When things happen matters almost as much as whether they happen. Payment terms (net-30, net-60, net-90) create gaps between revenue recognition and cash arrival. Hiring ramp time means a new salesperson does not produce at full capacity for 3-6 months. Seasonal patterns affect acquisition differently by quarter. Marketing spend precedes the revenue it generates by weeks or months.

Ignoring timing is one of the most dangerous mistakes in startup modeling. A company can be profitable on an annual basis while running out of cash in Q2 because all the big enterprise contracts pay net-90 and all the hiring happened in Q1.

How to Build a Driver-Based Model: Step by Step

Step 1: Map your value chain

Draw out how one unit of revenue is created. Start from the earliest input (a website visitor, a sales call, a marketplace listing) and trace every step until cash hits your bank account. Write down every

conversion rate, time lag, and cost at each step. This map becomes the skeleton of your model.

Step 2: Identify your primary driver

Every business has one driver that matters most. For SaaS, it is usually new logo acquisition rate. For marketplaces, it is typically liquidity (the match rate between supply and demand). For e-commerce, it is traffic multiplied by conversion. Find yours. This becomes the single cell in your model that, when changed, ripples through everything else.

Step 3: Build the assumptions tab

Create a dedicated tab where every driver lives as a named, sourced assumption. Growth rate in leads per month: 8% (source: last 6 months average). Close rate: 12% (source: CRM data, January to June). Average contract value: $6,200 (source: last 50 closed deals). Monthly logo churn: 4.5% (source: cohort analysis). Every number needs a source. If the source is "our best guess," say so. Investors respect honesty. They do not respect unsourced confidence.

Step 4: Wire the formulas

Connect your assumptions to your output tabs. Revenue should be a formula that references the assumptions tab, not a hardcoded number. New MRR = new customers (from acquisition model) multiplied by ACV (from assumptions tab). Churned MRR = existing MRR multiplied by churn rate (from assumptions tab). Net new MRR = new MRR minus churned MRR. Total MRR = prior month total plus net new MRR.

Do this for every financial line. COGS should reference cost-per-unit assumptions. Headcount costs should reference the hiring plan. Marketing spend should reference the budget allocation by channel.

Step 5: Validate with a sense check

Once the model is wired, look at the outputs and ask: does this imply something impossible? If your model shows $50M in revenue in Year 3 and your total addressable market is $200M, that is a 25% market share, possible for some businesses, absurd for others. If your model shows 500 customers by Month 18 and you have two salespeople, check the implied quota per rep. If your model shows $0 in customer support costs while serving 1,000 clients, something is wrong.

Top-down sense checks catch the errors that bottom-up construction misses.

A Worked Example: SaaS Startup

Here is a simplified driver-based revenue model for a B2B SaaS company at seed stage.

Driver	Value and Source
Monthly website visitors	15,000 (Google Analytics, trailing 3-month avg)
Visitor to demo rate	2.5% (Hubspot, trailing 3-month avg)
Demo to close rate	18% (CRM, last 2 quarters) Average contract value (annual)
Months to full ACV (ramp)	2 months (observed onboarding pattern)

2.5% = 375 demos. 375 demos multiplied by 18% = 67.5 new customers per

month (round to 68). 68 new customers multiplied by $600 MRR = $40,800

in new monthly revenue. After accounting for churn on the existing base

and the two-month ramp period, the model produces a net MRR growth

trajectory that is fully traceable to real business metrics.

If an investor asks ”what happens if your close rate drops to 12%?“

you change one cell and the entire model updates. That is the power of driver-based construction.

A Worked Example: Two-Sided Marketplace

Marketplace models require drivers on both sides. Working on the Indeed Flex model taught me this the hard way: you cannot model a marketplace like a SaaS business. Supply and demand interact, and the model has to reflect that interaction.

Driver (Supply Side)	Driver (Demand Side)
Workers registered per month	Employers posting shifts per month Worker activation rate
Average shifts per active worker	Average shifts posted per active employer
Fill rate (shifts matched/shifts	Gross margin per filled shift posted)

is the output of fills, and fills depend on the balance between supply and demand. A marketplace with 10,000 workers and 50 employers will have very different economics than one with 2,000 workers and 500 employers, even if total GMV is similar. The driver-based model has to capture that balance, or your projections will be structurally wrong.

> *Key insight: The test of a good driver-based model is not whether the output looks right. It is whether changing any single input produces a cascade of realistic, traceable changes through every tab. If you change churn and nothing happens to your headcount plan for customer success, the model is not truly driver-based.*

Common Mistakes in Driver-Based Modeling

After building or reviewing hundreds of models, these are the mistakes I see most frequently:

Using a single blended conversion rate. Channel-level conversion varies enormously. Organic search converts at 4-8% for most SaaS companies. Paid social converts at 0.5-2%. Blending these into "3% average" hides the fact that one channel is working and another is burning cash.

Not separating new vs. expansion revenue. A model that shows $100K MRR growth without distinguishing between new customer acquisition and existing customer expansion is incomplete. The growth rate of each is driven by completely different factors and funded by different budgets. **Assuming constant churn.** Churn typically improves as your product matures and your customer base shifts toward better-fit accounts. But it can also spike during price increases, market downturns, or product transitions. Model it as a variable, not a constant.

Disconnecting marketing spend from acquisition. If your marketing budget doubles and your lead count stays the same, the model is lying. Marketing spend should directly drive your top-of-funnel metrics. If the relationship is not linear, model the diminishing returns explicitly.

Investor Perspective: What VCs Check First

When Profounders received our model, the first tab they opened was the assumptions tab. The second thing they did was trace the revenue line back to see if it was formula-driven or hardcoded. The entire evaluation took about 15 minutes before they decided to take a deeper meeting. Here is what the best investors check, in roughly this order:

Check	What They Are Looking For
Revenue construction	Is it built from drivers or just a growth rate?

Assumption sourcing	Is every key input sourced or at least labelled?
Churn treatment	Is retention modeled explicitly, with cohort data?
Unit economics	Does CAC, LTV, and payback derive from the model?
Cash flow timing	Does the model capture payment lags and burn timing?
Scenario sensitivity	Can they change one driver and see what breaks?

How Cohort Forecasting Works

The mechanics are straightforward once you understand the structure.

Step 1: Define your cohorts

Group customers by acquisition month. January 2025 cohort = all customers who made their first purchase or signed up in January 2025. Each cohort gets its own row in the model.

Step 2: Track revenue per cohort per month

For each cohort, record the revenue generated in Month 0 (their first month), Month 1, Month 2, and so on. This creates a triangle-shaped data structure. The January cohort has the most months of data. The most recent cohort has only one month.

Step 3: Calculate retention curves

For each cohort, calculate the percentage of Month 0 revenue that remains in each subsequent month. If the January cohort generated $10,000 in Month 0 and $8,500 in Month 1, the Month 1 retention rate is 85%. Do this for every cohort.

Here is where the magic happens: you will see patterns. Maybe your average Month 1 retention is 85%, Month 3 is 72%, Month 6 is 58%, and Month 12 is 45%. Or maybe your cohorts show improving retention

over time because your product is getting better. Both patterns are extremely valuable information that aggregate analysis would never reveal.

Step 4: Build the forecast

For future months, apply the observed retention curves to existing cohorts and the assumed acquisition rate to new cohorts. Total revenue for any future month = sum of all active cohorts' expected revenue for that month.

This means your forecast accounts for the fact that older cohorts will continue to shrink (or grow, if your NRR is above 100%) while new cohorts come in. The interaction between these two forces is what determines your actual revenue trajectory.

A Practical Example

Let us build a simple cohort model for a SaaS startup with $500 average MRR per customer.

Month	New Customers Acquired	Retained from Prior Cohorts	Total Active

Month 2: 80%, Month 3: 75%, Month 4: 72%, Month 5: 70%, Month 6+: 69%

(stabilized).

January cohort revenue in June: 20 customers multiplied by $500 multiplied by 70% retention = $7,000. February cohort revenue in June: 25 multiplied by $500 multiplied by 72% = $9,000. March cohort: 30 multiplied by $500 multiplied by 75% = $11,250. April cohort: 35 multiplied by $500 multiplied by 80% = $14,000. May cohort: 40 multiplied by $500 multiplied by 88% = $17,600. June cohort (new): 45 multiplied by $500 multiplied by 100% = $22,500.

Total June revenue: $81,350. An aggregate model would not tell you that

$22,500 of that is brand new (and fragile), while $7,000 is from a highly retained, stable base. The cohort model gives you both numbers.

> *Key insight: The single most powerful thing a cohort model reveals is the difference between growing because you are acquiring fast and growing because your customers are staying. One is sustainable. The other is a treadmill. Investors know which one they are betting on.*

Why This Saved an Exit

During the Indeed Flex acquisition process, the acquirer challenged our revenue projections. Their concern was that headline revenue growth could be masking customer loss. They were right to be cautious. It is a common pattern.

What saved the conversation was cohort-level data. We could show that not only were older cohorts retained, they were actually spending more per month. The January 2019 cohort was generating 30% more revenue per surviving customer in January 2021 than they were at acquisition. This expansion pattern, visible only at the cohort level, proved the revenue base was strengthening, not just growing. It was one of the key factors that supported the valuation multiple.

How to Handle Common Complications

What if my cohorts are too small to be statistically meaningful? This is common at very early stages. If you are acquiring 5-10 customers per month, individual cohorts will be noisy. Group them into quarterly cohorts instead, or combine your earliest months into a single "founding cohort." The key is to have enough customers per group that the retention rate is not being driven by one or two outlier accounts. **What about expansion revenue?**

Cohort models handle expansion naturally. If the June revenue from the January cohort is higher than the May revenue, your retention rate for that period is above 100%. This is net revenue retention at the cohort level. SaaS businesses with strong expansion (upsell, cross-sell,

seat growth) will show cohort curves that flatten or even grow over time. This is the single strongest signal an investor can see in a model.

What about different pricing tiers or segments?

Build separate cohort models for each material segment. An enterprise customer acquired in January behaves very differently from an SMB customer acquired in January. If you blend them into one cohort, you lose the ability to see that your enterprise retention is 95% while your SMB retention is 60%. Those are two completely different businesses disguised as one. Indeed Flex modeled cohorts separately by client size and market (UK vs. other) because the retention patterns were materially different.

Investor Perspective

When a fund like Creandum or B2Ventures evaluates a SaaS or marketplace model, the cohort analysis is where they spend the most time. Here is what they look for:

Signal	What It Tells the Investor
Stable retention after Month 6	Predictable revenue base NRR above 100% per cohort
Declining retention in recent	Product, market, or ICP problem cohorts

immediate credibility. A founder who only has aggregate revenue numbers gets asked for cohort data and loses time and confidence.

Cohort-based revenue forecasting is the gold standard for any startup with recurring or repeat revenue. It separates the health of your existing customer base from the performance of your acquisition engine, making your forecasts more accurate and your investor conversations more credible. Build the structure early, even if your cohorts are small. Track retention at the cohort level, account for expansion, and segment by customer type if the behavior differs

meaningfully. This is not just a better forecasting method. It is a better way to understand your business.

024 The Assumptions Tab: The Most Important Sheet in Your Entire Financial Model

What Is an Assumptions Tab?

An assumptions tab is a dedicated sheet in your financial model where every variable input is stored in one place. Instead of burying your growth rate in a formula on the revenue tab, your churn rate in a cell on the retention sheet, and your hiring timeline scattered across the headcount plan, you centralize them all.

Every formula in every other tab of the model should reference back to this sheet. When you change an assumption, the entire model updates. When an investor wants to understand what you believe about your business, they look at one tab instead of hunting through 15 sheets. Here is the simplest way to think about it: the assumptions tab is the control panel. Everything else is the machine. If you built the machine but forgot the control panel, nobody can operate it, including you.

Why Investors Open This Tab First

I have sat in rooms with analysts from Creandum, Profounders, and B2Ventures, and the pattern is always the same. They open the model. They look for a tab called "Assumptions" or "Inputs" or "Drivers." If it exists and is well-structured, they nod and start reading. If it does not exist, they start clicking through random tabs trying to find where the inputs live, and their confidence in the founder drops with every click.

The reason is simple: the assumptions tab is the fastest way to evaluate two things at once.

First, the quality of the founder's thinking. Are the assumptions realistic? Are they sourced? Do they demonstrate an understanding of the business drivers? A founder who lists "monthly churn: 3%, source: industry average (OpenView SaaS Benchmarks

2024)" has done the work. A founder whose churn rate is buried in a formula with no documentation has not.

Second, the model's testability. Can the investor change an input and see what happens? A centralized assumptions tab with clearly labeled cells makes scenario analysis trivial. A model where assumptions are hardcoded into formulas across six different tabs makes scenario analysis impossible without rebuilding the model.

> **Key insight: An investor can evaluate the quality of your entire model in under five minutes by reading your assumptions tab. If the assumptions are sourced, logical, and comprehensive, they trust the model. If the tab is missing or thin, the rest of the model becomes suspect.**

What Belongs in the Assumptions Tab

Every material input that drives a financial output. Here is the comprehensive list, organized by category:

Revenue Assumptions

Assumption	Example	Source
Monthly lead volume	8,000	Google Analytics, Q4 average
Lead-to-demo conversion rate	3.2%	Hubspot, trailing 6 months
Demo-to-close rate	15%	CRM, last 2 quarters
Average contract value (ACV)	$8,400/year	Last 40 closed deals
Monthly logo churn	4.2%	Cohort analysis, 12-month
Net revenue retention	108%	Existing customer data, annual

Expansion rate	2.1% MoM	Upsell records, 6 months
Sales cycle length	42 days average	CRM pipeline data

Cost Assumptions

Cost Item	Example Value	Notes
Starting headcount	8 FTEs	Current team size
Average salary by function	Engineering: $115K, Sales: $85K + OTE	Fully loaded with benefits
Payroll overhead multiplier	1.25x	Benefits, tax, insurance
Hiring timeline	2 engineers Q2, 1 AE Q3	Per hiring plan
Infrastructure cost per user	$0.42/user/month	AWS billing, last 3 months
Marketing spend (monthly)	$25K current, scaling to $40K by Q4	Aligned with growth targets
CAC by channel	Paid: $380, Organic: $85, Referral: $40	Blended across channels

Cash Flow Assumptions

Assumption	Example	Notes
Payment terms (receivables)	Net-30 for SMB, Net-60 for enterprise	Affects cash collection timing
Payment terms (payables)	Net-30 for most vendors	Standard vendor terms
Billing frequency	80% monthly, 20% annual upfront	Annual contracts aid cash flow

Fundraise timing	Month 8 of model (close by October)	Aligned with board timeline
Monthly burn rate (current)	$92K	Trailing 3-month average

Market Sizing Assumptions

Assumption	Example	Reasoning
TAM	$4.2B	Source: Gartner 2024 market report
SAM	$800M	Source: geographic and segment filter
Growth rate in leads	8% MoM	Source: trailing 6-month trend
Seasonal adjustment	Q1 −15%, Q4 +20%	Source: 2 years of data

How to Structure the Tab

The layout matters. A well-structured assumptions tab reads like a reference document, not a data dump. Here is the structure I use for every model:

Column A: Category. Revenue, Costs, Cash Flow, Market. Group assumptions logically so they are easy to scan.

Column B: Assumption name. Use plain language. "Monthly logo churn rate" not "churn_m."

Column C: Value. This is the cell that every other tab references. Format consistently (percentages as percentages, currencies with the right symbol, counts as integers).

Column D: Source. Where did this number come from? Internal data, industry benchmark, analyst report, or best estimate. Be specific. "Industry average" is not a source. "OpenView SaaS Benchmarks 2024, median for $1-5M ARR companies" is.

Column E: Date. When was this assumption last validated? This matters because a churn rate from 18 months ago may no longer reflect your current product. Investors look at assumption freshness. **Column F: Notes.** Any additional context. "Likely to improve after Q3 product release." "Conservative relative to comparable companies." "This is our weakest assumption; see sensitivity analysis." *Key insight: The best assumptions tabs I have seen, including the one we built for the Indeed Flex exit, had a color-coding system. Green for assumptions backed by 6+ months of internal data. Yellow for assumptions backed by benchmarks or limited data. Red for assumptions that are educated guesses. This gives anyone reading the model an instant visual map of where the certainty lives and where the risk is.*

Connecting Assumptions to Scenarios

The assumptions tab is what makes scenario planning possible. Without it, building three scenarios means copying your entire model three times. With it, you add three columns to the assumptions tab: conservative, base, and aggressive.

Each scenario has its own value for each assumption. Conservative churn: 6%. Base churn: 4.2%. Aggressive churn: 3%. A single toggle cell at the top of the tab (set to 1, 2, or 3) tells the model which column to read. The entire model shifts with one click.

This structure is exactly what investors want to see. They can sit in a meeting, say "what if your churn is 6% instead of 4%?", and you can switch the scenario in real time. That builds confidence faster than any slide deck.

The Seven Mistakes That Break Assumptions Tabs

1. **Hardcoded numbers in formulas.** If your revenue tab says =B12*0.15, that 0.15 should be a reference to the assumptions tab. Every time you type a number directly into a formula, you are hiding an assumption.

2. **Unsourced assumptions.** "We believe" is not a source. Every material assumption needs data, a benchmark, or at minimum an honest label saying "founder estimate."
3. **Stale dates.** An assumption validated 18 months ago might as well be a guess. Update the date column when you revalidate, and flag assumptions older than 6 months for review.
4. **Missing assumptions.** If your headcount plan has salary numbers that do not trace back to the assumptions tab, you have a gap. Audit by trying to change every number in the model from the assumptions tab alone. If anything does not update, it is hardcoded somewhere. **5. Too many assumptions.** A seed-stage model should have 15-30 key assumptions. More than 50 and you are over-engineering. Not every minor input needs to be an assumption. Office supply costs do not need to be on this tab.
6. **No categorization.** A flat list of 30 assumptions with no grouping is hard to navigate. Group by category and add section headers. **7. No scenario columns.** If your assumptions tab only has one value per assumption, you cannot run scenarios without rebuilding. Add the three-scenario structure from the start, even if you only use the base case initially.

A Real-World Test: The 5-Minute Investor Audit

Here is what happens when an experienced investor receives your model: Minute 1: They look for the assumptions tab. If it does not exist, they close the file and reply "could you send a version with documented assumptions?" This has happened to three founders I have worked with. It is recoverable, but it wastes time and signals unpreparedness. Minute 2: They scan the assumption names. Are the revenue drivers granular (leads, conversion, ACV, churn) or vague ("growth rate")? Granular earns trust. Vague triggers skepticism.

Minute 3: They check sources. Are sources specific and recent? Or are they generic or missing? This is where most assumptions tabs fail. Minute 4: They test the model. They pick one assumption, change it, and check whether the outputs update correctly. If the model breaks or

does not respond, it tells the investor the assumptions are decorative, not functional.

Minute 5: They look at the scenario toggle. Can they flip between conservative, base, and aggressive with one click? If yes, they start engaging with the scenarios. If no, they start marking their own version of scenarios in a side spreadsheet, which means you have lost control of the narrative.

The assumptions tab is where your entire financial model lives or dies. It centralizes every key input, sources it, dates it, and makes it editable. It is the first thing investors open, the foundation of scenario analysis, and the fastest way to demonstrate that you understand your own business. Build it with categories, sources, dates, notes, and scenario columns. Audit it by changing each assumption and confirming the model responds. Update it monthly as real data replaces estimates. This single tab determines whether your model is a credible decision-making tool or a beautifully formatted guess.

025 How to Build a Headcount Plan That Investors Trust and Founders Can Actually Use

Why Headcount Planning Matters More Than Founders Think

Personnel cost is the largest single expense category for almost every startup before profitability. At Indeed Flex, headcount represented over 70% of our operating burn at every stage from seed through Series A. Every other cost category, infrastructure, marketing, office space, was secondary.

This means that if your headcount plan is wrong, your entire burn rate is wrong. If your burn rate is wrong, your runway calculation is wrong. And if your runway is wrong, you will either run out of money sooner than expected or raise at the wrong time. The downstream consequences of a bad headcount plan are existential.

Investors understand this, which is why headcount is one of the first areas they scrutinize. When Profounders reviewed our model, they did not just look at the total salary line. They opened the headcount tab, looked at each planned hire, checked the timing against revenue milestones, and asked: "Why do you need this role in Q2 instead of Q3?" *Key insight: A headcount plan is not a wish list. It is a sequenced investment thesis that says: "If we hire these people in this order, we believe the business will reach these milestones by these dates." Every role needs a justification tied to a business outcome.*

The Structure of a Proper Headcount Plan

Every headcount plan should be built as a table with these columns:

Column	What It Contains

Role title	Specific title (e.g., ”Senior Backend Engineer“, not
Department	Engineering, Sales, Marketing, Operations, G&A
Start month	When this person's first paycheck hits
Base salary (annual)	Gross annual salary in local currency
Payroll overhead multiplier	1.2x to 1.5x depending on geography Fully loaded monthly cost
Ramp time	Months to full productivity (especially sales roles)
Revenue linkage	What revenue or capacity this role enables

Understanding Payroll Overhead

This is the number first-time founders almost always underestimate. Base salary is not what an employee costs you. The actual cost includes benefits (health insurance, pension/401k), employer-side payroll taxes, equipment, software licenses, and if you are using an Employer of Record for international hires, the EOR fee.

The multiplier varies significantly by geography:

Geography	Typical Overhead Multiplier
United States	1.25-1.35x (health insurance, FICA, state taxes)
United Kingdom	1.15-1.25x (NI contributions, pension auto-enrolment)
France	1.40-1.50x (social charges are among the highest in Europe)

Romania / Eastern Europe	1.08-1.15x (lower social contributions)
UAE / Dubai	1.05-1.12x (minimal payroll taxes, visa costs instead)
Via Employer of Record	Add 10-20% on top of local overhead for EOR fee

UK, US, and Europe, the difference in overhead between a French employee and a Romanian employee was nearly 40 percentage points. In a 50-person team spread across three markets, that variation means the difference between 18 months of runway and 14 months. It is not a rounding error.

> *Key insight: If your startup operates across multiple countries or uses an Employer of Record, model each employee's loaded cost individually. A blended "1.25x overhead" will understate your costs in high-overhead markets and overstate them in low-overhead ones. The net effect is usually an unpleasant surprise when actuals come in above forecast.*

How to Sequence Hiring

The order in which you hire matters as much as who you hire. A common mistake is to plan all hires for "Q1" because the money just arrived. In practice, hiring takes 2-4 months from decision to first day, and a new hire takes another 2-6 months to reach full productivity. Your model needs to reflect both lags.

Hiring for revenue generation

Sales roles should be timed based on your acquisition model. If each account executive can close 8 deals per month once ramped, and your model projects needing 16 new deals per month by Q3, you need to have hired the second AE by Q1 at the latest (assuming a 3-month ramp). Work backwards from the revenue target to the hire date, not forwards from the funding date to the role.

Hiring for capacity

Engineering and product hires should be timed based on product milestones. If your revenue model assumes a new product line launching in Q3, the engineers building it need to start in Q1. Customer support hires should scale with customer count, typically 1 support person per 200-400 customers for SaaS, though this varies by product complexity. **Hiring for operations**

Finance, HR, and operations hires are often deferred too long. Most startups should hire their first finance or operations lead around the 15-20 employee mark or $1.5M ARR, whichever comes first. Waiting longer creates the kind of operational debt that manifests as inaccurate financial reporting, compliance gaps, and messy cap tables.

A Worked Example: Seed Stage SaaS (12-Month Headcount Plan)

Here is a realistic headcount plan for a seed-stage SaaS company that just raised $2.5M.

Role	Start	Base	Loaded	Dept	Revenue Link
CTO (existing)	M1	$140K	$175K	Eng	Core product
CEO (existing)	M1	$120K	$150K	G&A	Strategy
Sr. Backend Eng	M1	$130K	$163K	Eng	Core product
Frontend Engineer	M2	$115K	$144K	Eng	Core product
Account Exec #1	M2	$80K+ OTE	$130K	Sales	8 deals/mo

Product Designer	M3	$105K	$131K	Eng	User growth
Account Exec #2	M5	$80K+ OTE	$130K	Sales	8 deals/mo
Customer Success	M5	$75K	$94K	Ops	Retention
Marketing Lead	M6	$100K	$125K	Mktg	Lead gen
Backend Eng #2	M7	$125K	$156K	Eng	Scale
SDR	M8	$55K+ OTE	$88K	Sales	Pipeline
Finance/ Ops Lead	M9	$95K	$119K	G&A	Ops debt

How the Headcount Plan Connects to the Rest of the Model

A well-built headcount plan is not a standalone document. It feeds directly into three other parts of your financial model:

P&L: Salaries and benefits flow into operating expenses. Split by department (Engineering, Sales, Marketing, G&A) so investors can see your cost allocation.

Cash flow: Salary payments typically hit monthly, but signing bonuses, equipment purchases, and recruiter fees are lumpy. Model the timing, not just the annual total.

Revenue capacity: If your revenue model assumes 24 new deals per month by Q4, your headcount plan must show enough ramped salespeople to deliver that. If it does not, the investor sees the gap instantly.

What Investors Look For in a Headcount Plan

Signal	What It Means to an Investor
Role-by-role detail	Founder understands their burn composition
Start dates aligned to milestones	Hiring is strategic, not reactive Loaded costs, not just base salary
Department allocation visible	Cost structure is transparent

Headcount is most startups' largest cost, and it deserves a dedicated, detailed plan. Model every role individually with base salary, payroll overhead, start month, ramp time, and a clear business justification. Sequence hires based on business milestones, not funding dates. Account for geographic variation in overhead costs. Connect the plan directly to your revenue model and cash flow statement. This is not busy work. It is the most operationally useful part of your financial model, and it is one of the first things every investor will scrutinize.

026 COGS vs. OpEx: Why Getting This Wrong Destroys Your Gross Margin Story

What Is COGS?

Cost of Goods Sold represents the direct costs of delivering your product or service. If you did not have the customer, you would not have the cost. The technical definition is: costs that vary directly with revenue.

For a physical product company, this is intuitive: raw materials, manufacturing, shipping. For software and marketplace businesses, it is less obvious, which is where the mistakes happen.

COGS for a SaaS Business

In SaaS, COGS typically includes: hosting and infrastructure costs (AWS, GCP, Azure) that scale with user count, third-party software costs baked into the product (payment processing, data APIs, messaging services), customer support and onboarding staff (salaries of people directly serving customers), and DevOps or site reliability engineering costs. It does not include: engineering salaries for product development (that is R&D, an operating expense), sales and marketing costs (operating expense), office rent (operating expense), or executive salaries (operating expense).

COGS for a Marketplace

At Indeed Flex, COGS included the direct costs of fulfilling a shift: worker payments, employer liability insurance, payment processing fees, and the cost of the operations team that handled shift matching and dispute resolution. These costs existed only because shifts were being filled. No shifts, no COGS.

The sales team that acquired employers? That was OpEx. The engineering team that built the platform? OpEx. The office in Holborn?

OpEx. The distinction is clear once you apply the test: would this cost exist if we had zero transactions this month?

COGS for a Services Business

For consulting, professional services, or agency models, COGS is the cost of the people delivering the service and any direct costs associated with delivery (travel, subcontractors, specialized tools used exclusively for client work). The salaries of consultants billing hours are COGS. The salary of the office manager is OpEx.

Cost Item	COGS or OpEx?	Reasoning
AWS	hosting (scales with users)	COGS
Stripe	payment processing fees	COGS
Customer	support team salaries	COGS
Engineering	team salaries (product	OpEx (R&D)
Sales team salaries	OpEx (Sales & Marketing) Google Ads spend	OpEx (Sales & Marketing) Office rent
CEO	salary	OpEx (G&A)
Legal	and accounting fees	OpEx (G&A)
Worker	payments in a marketplace	COGS

Why Gross Margin Is the Number Investors Check First

Gross margin = (Revenue minus COGS) / Revenue. It tells investors how much money is left over from each dollar of revenue after paying for the direct costs of delivery. Everything else (salaries, marketing, office, growth) gets funded from what is left.

A SaaS company with 75% gross margin has $0.75 of every revenue dollar available to fund growth. A marketplace with 25% net revenue

margin has $0.25. A services company with 40% gross margin has $0.40. These numbers dictate the entire growth equation: how fast you can scale, how much you can spend on acquisition, and how quickly you reach profitability.

Business Model	Typical Gross Margin Range
SaaS (software delivery)	70-85%
Marketplace (net revenue basis)	60-75%
Marketplace (GMV basis)	15-30%
E-commerce (own inventory)	30-50%
Professional services	35-55%
Hardware + software	40-60%

question will be: ”What is in your COGS?“ Because either you have a structural delivery cost problem (bad sign), or you have misclassified operating expenses as COGS (which means the model is unreliable). Both answers lead to difficult conversations.

> *Key insight: Gross margin misclassification is one of the most common model credibility killers I have seen across multiple fundraising processes. Moving one engineer's salary from OpEx to COGS can swing your gross margin by 5-10 percentage points at early stages. Investors will check your classification. If it does not match industry standards, they will reclassify it themselves, and now your numbers and their numbers disagree before the conversation even starts.*

How to Model COGS Correctly

Step 1: List every cost in your business

Write down every expense, from AWS bills to the CEO's salary. Do not categorize yet. Just list.

Step 2: Apply the incremental customer test

For each cost, ask: "If we added 100 more customers tomorrow, would this cost increase?" If yes, it is COGS (or at least partially). If no, it is OpEx. AWS hosting scales with users? COGS. Office rent stays the same regardless? OpEx. Customer support team needs more people as customer count grows? COGS.

Step 3: Handle the gray areas

Some costs do not fit neatly. DevOps engineers maintain the infrastructure (COGS-related) but also build tooling (R&D, which is OpEx). The standard approach: if more than 50% of their time is on infrastructure maintenance, classify the role as COGS. Otherwise, OpEx. Be consistent and document your allocation logic in the assumptions tab. Similarly, customer success roles that are partly onboarding (COGS) and partly upselling (Sales, OpEx) can be split. A 70/30 COGS/OpEx allocation is defensible if documented. A 100% COGS allocation for a role that spends half its time on expansion revenue is not. **Step 4: Benchmark your gross margin**

Once you have classified everything, calculate your gross margin and compare it to comparable companies at your stage and business model. If your SaaS gross margin is 55% and the benchmark is 75%, something is likely misclassified, or you have a genuine delivery cost problem that needs to be addressed in the investor narrative.

How COGS Changes as You Scale

One of the most valuable things a model can show is how gross margin improves with scale. Most software businesses have economies of scale in COGS: the cost of hosting 1,000 users is not 10x the cost of hosting 100 users. Payment processing fees might decrease as volume increases (negotiated rates). Customer support can become more efficient with better tooling and knowledge bases.

Model this trajectory explicitly. Show gross margin at current scale, at 2x revenue, and at 5x revenue. If gross margin improves from 68% to 78% over that range, it tells a compelling story about the leverage in

your business model. If it stays flat or declines, you need to understand and explain why.

At Indeed Flex, gross margin improved meaningfully as we scaled because the per-shift operational cost decreased with volume (better matching algorithms, lower support ticket rates per shift, negotiated insurance rates). Modeling that improvement was critical to the Series A narrative because it showed the business would become significantly more profitable at scale.

The Investor Conversation About Margins

When Creandum and B2Ventures looked at our model, the gross margin discussion went like this:

Investor: "Walk me through your COGS."

Us: "Per filled shift, we have worker payment, employer insurance, payment processing, and a per-shift allocation of the operations team. Here is the breakdown."

Investor: "What happens to gross margin at 3x volume?"

Us: "Insurance rate negotiation at higher volume gives us 2 points. Operations efficiency gives us 3 points. Payment processing scale discount gives us 1 point. We model gross margin improving from 22% to 28% net."

Investor: "That is clear. Show me the sensitivity if insurance costs rise 15%."

That conversation only works if your COGS classification is clean and your model has the structure to show margin evolution. If COGS is a single number with no breakdown, none of those questions can be answered, and the investor loses confidence in the model.

COGS is the cost of delivering your product. OpEx is the cost of running your company. Getting the classification right determines your gross margin, which is one of the most scrutinized numbers in any investor evaluation. Use the incremental customer test: if the cost scales with customers, it is COGS. Model each COGS component individually so you can show margin evolution at scale. Benchmark against

comparable companies to ensure your classification matches industry standards. Document your allocation logic in the assumptions tab. A misclassified model does not just produce wrong numbers. It produces wrong conversations, and those are much harder to fix.

027 The Cash Flow Statement Founders Always Get Wrong

Why Profit and Cash Are Not the Same Thing

This is the most dangerous misunderstanding in startup finance. Revenue recognition and cash receipt are two different events. When you close a $60,000 annual contract, your P&L might recognize $5,000 per month in revenue starting immediately. But if the customer pays net-60, you will not see that cash for two months. If they pay annually upfront, you receive $60,000 on day one but can only recognize $5,000 per month in revenue.

The same asymmetry exists on the cost side. You hire an engineer in January, their first paycheck hits in February, but the recruiter fee lands in March. You sign a 12-month software license and pay the full amount upfront, but the expense is recognized monthly.

These timing differences are small individually. Collectively, they can create a gap of 2-4 months between what your P&L says and what your bank account shows. For a startup burning $100K per month, a 2-month gap is $200K, and that can be the difference between having runway and not having runway.

> **Key insight: I have seen funded startups with positive unit economics and growing revenue come within weeks of running out of cash because they managed from the P&L instead of the cash flow statement. The P&L tells you whether the business model works. The cash flow statement tells you whether the business survives long enough to prove it.**

The Three Sections of a Cash Flow Statement

A startup cash flow statement has three sections. Each captures a different type of cash movement, and together they produce your net cash position and runway.

1. Cash Flow from Operations

This is the cash generated or consumed by your core business activities. Start with net income (from the P&L), then adjust for non-cash items and timing differences.

Key adjustments include:

Accounts receivable changes. If you invoiced $50K this month but only collected $30K, the difference ($20K) is cash you earned but do not have yet. It reduces your operating cash flow.

Accounts payable changes. If you received $25K in vendor invoices but only paid $15K, you are holding $10K longer. That temporarily boosts cash flow. It does not mean you are richer; it means you have not paid the bill yet.

Deferred revenue. If a customer pays annually upfront, you receive $60K in cash but only recognize $5K in revenue this month. The remaining $55K is deferred revenue, a liability on your balance sheet but cash in your bank account. This is why annual prepayment is so valuable for startups: it improves cash flow even though it does not accelerate revenue recognition.

Depreciation and amortization. These are P&L expenses that do not consume cash (the cash was spent when the asset was purchased). Add them back to get from accounting profit to cash profit.

2. Cash Flow from Investing

This captures cash spent on long-term assets: equipment, capitalized software development costs, office buildout. For most pre-Series A startups, this section is small. A laptop purchase, maybe some office furniture. If you are capitalizing development costs (which is an accounting choice that has tax implications), those go here rather than in operations.

The key rule: if the spend creates an asset on the balance sheet rather than an expense on the P&L, it belongs in investing cash flow. **3. Cash Flow from Financing**

This is where fundraising appears. Equity investment, convertible notes, venture debt drawdowns, and loan repayments all live here.

When you raise $2.5M, it shows up as a positive cash flow from financing. When you repay venture debt principal, it shows as negative.

This section is often the simplest but the most important for runway calculation. The timing of your raise directly affects when your cash position peaks and how long the runway extends.

How to Build the Cash Flow Statement: Step by Step

Step 1: Start with your P&L

Pull net income for each month from your P&L. This is your starting point. Every adjustment moves from accounting reality to cash reality. **Step 2: Model your receivables cycle**

Determine your average collection period. If your customers pay net-30, model a one-month lag between revenue recognition and cash receipt. If you have a mix of payment terms (SMB pays net-15, enterprise pays net-60), model each segment separately.

A practical approach: calculate your Days Sales Outstanding (DSO). DSO = (Accounts Receivable / Revenue) x 30. If your DSO is 45 days, you are collecting cash about 6 weeks after invoicing. Your cash flow model needs to reflect that 6-week delay.

Step 3: Model your payables cycle

When do you actually pay your bills? Most startups pay vendors on net-30 terms. Payroll hits biweekly or monthly. AWS bills on the 3rd of the following month. Map each major cost category to its payment timing. **Step 4: Add your deferred revenue**

If customers pay upfront (monthly, quarterly, or annually), the cash arrives before the revenue is recognized. Model the timing of each payment pattern. A SaaS company with 30% annual prepay customers receives a significant cash inflow at contract signing that smooths out the monthly burn.

Step 5: Calculate net cash flow and cumulative cash position Net cash flow = cash in minus cash out for each month. Cumulative cash position = prior month's cash balance plus this

month's net cash flow. The cumulative cash row is your runway visualized as a number. When it hits zero, the company stops.

The Runway Calculation

Runway is not "total cash divided by monthly burn." That formula only works if burn is constant, and it never is. Runway should be read directly from the cash flow model: the month where cumulative cash goes negative is the month you run out of money.

For fundraising purposes, you need three versions of this number:

Scenario	What It Tells You
Base case runway	How long you last if things go roughly as planned
Conservative runway	How long you last if revenue underperforms by 20-30%
Zero-revenue runway	How long you last if revenue stops entirely tomorrow

base case gives you 18 months of runway but your conservative case gives you 9, that gap is a red flag. It means a moderate revenue miss puts you in emergency fundraising mode within three quarters.

At Indeed Flex, we tracked all three runway numbers weekly during fundraising periods. The zero-revenue runway was the number the board watched most closely, because it answered the question: ”If everything goes wrong, how much time do we have to fix it?“

> *Key insight: The most dangerous moment for a startup's cash is not when things go badly. It is when things go well but cash lags behind. Rapid growth often means higher upfront costs (hiring, marketing, inventory) before the revenue from that growth converts to cash. Growing fast while*

> *cash-flow-negative is a controlled burn. Growing fast while ignoring the cash flow statement is an uncontrolled one.*

A Worked Example: SaaS Startup Cash Flow

Here is a simplified three-month cash flow for a SaaS company doing $80K MRR with a $2.5M raise closing in Month 1.

Line Item	Month 1	Month 2	Month 3
Net income (from P&L)	−$45,000	−$38,000	−$32,000
Add back: depreciation	+$2,000	+$2,000	+$2,000
Change in accounts receivable	−$12,000	−$8,000	−$5,000
Change in deferred revenue	+$18,000	+$5,000	+$3,000
Change in accounts payable	+$3,000	−$1,000	+$2,000
Cash from operations	−$34,000	−$40,000	−$30,000
Equipment purchases	−$8,000	$0	$0

cash outflow was only $34K because of deferred revenue (customers paying upfront) and payable timing. The $2.5M raise in financing creates the runway. At a $40K monthly cash burn from operations, the company has roughly 62 months of runway, but that assumes burn stays constant. In reality, burn will increase as the team grows, which is why the model needs to project cash flow 18-24 months out, not just

calculate a static ratio.

The Five Cash Flow Mistakes That Kill Startups

1. **Modeling revenue as cash.** Booking $100K in revenue this month does not mean $100K arrived in your bank account. If your average DSO is 45 days, roughly half of that revenue is still in accounts receivable. Enterprise SaaS companies with net-60 or net-90 terms are especially vulnerable here.
2. **Ignoring payment timing on the cost side.** Annual software licenses, quarterly insurance premiums, recruiter fees on signing, these create lumpy cash outflows that monthly burn rate averages miss entirely. A $36K annual license paid upfront in January means January's cash burn is $33K higher than February's.
3. **Forgetting the fundraise itself has costs.** Legal fees for a seed round run $15-30K. Due diligence, data room setup, travel for partner meetings, these add up. The cash from the raise arrives net of these costs, not gross.
4. **Not modeling the gap between hiring and revenue.** You hire two salespeople in Month 3 at a combined loaded cost of $25K per month. They ramp for 3 months. They start closing deals in Month 6. Cash from those deals arrives in Month 7 or 8. That is a 5-month gap where costs are real but revenue contribution is zero. If your cash flow model does not capture this, your runway is overstated.
5. **Using annual averages instead of monthly actuals.** A company can be cash-positive for the year while going cash-negative for three consecutive months. Monthly granularity in the cash flow statement is not optional. It is the only way to see the dips that could kill you.

How Investors Use Your Cash Flow Statement

During the Series A process with Creandum, the cash flow statement was where the conversation shifted from "is this a good business?" to "is

this a good investment?" The P&L and unit economics told them the model worked. The cash flow statement told them:

Question	Where the Answer Lives
How much runway does the current	Cumulative cash row, base case raise buy?
What if revenue is 25% below plan?	Conservative scenario cash flow How capital-efficient is this team? Operating cash flow vs. headcount growth
When will this company need to	Month where conservative cash goes raise again?

month by month, explaining why each quarter looks the way it does,

demonstrates operational command. A founder who can only discuss revenue

and hopes the cash side works out signals inexperience.

The cash flow statement is the financial statement that keeps your company alive. It captures the timing differences between revenue and cash, expenses and payments, and financing events and burn. Build it using the indirect method: start with net income, adjust for receivables, payables, deferred revenue, and non-cash items. Calculate runway from the cumulative cash row, not from a static ratio. Model monthly for at least 18-24 months. Run it under conservative assumptions and know exactly when your cash goes to zero. This is the statement that answers the question every founder and every investor needs answered: how long do we have, and what do we need to do before the money runs out?

028 How to Build Three Scenarios That Prove You Have Actually Thought About What Could Go Wrong

Why Single-Scenario Models Fail

A financial model with one set of numbers makes an implicit claim: "This is what will happen." Every investor in the room knows that is false. Startups operate in conditions of extreme uncertainty. No one knows the exact churn rate, conversion rate, or hiring timeline. A model that pretends to know these things signals either naivety or dishonesty, and neither is a good look in a fundraising meeting.

Across every fundraising process I have worked on, from seed through the Indeed Flex exit, the scenario that generated the most investor questions was the conservative case. Not because investors are pessimists, but because the conservative case answers the question they care about most: "What happens when things do not go as planned, and do these founders have a plan for that?"

A founder who presents three scenarios is saying: "I know this could go multiple ways, and I have thought about what I will do in each case." That is what operational readiness looks like, and investors fund it.

The Three Scenarios Defined

Conservative: The Survival Case

This is not a disaster scenario. It is the scenario where things are harder than expected. Growth is slower. Churn is higher. Hiring takes longer. Sales cycles stretch. The question the conservative case answers: can this business survive and hit meaningful milestones without needing an emergency raise?

Specifically: if your conservative case shows you running out of cash before reaching a fundable milestone, your model has a structural problem. Either you need to raise more, spend less, or redefine the milestones.

Base: The Plan Case

This is what you believe will happen given your current trajectory, market conditions, and execution capacity. It should be grounded in real data: historical growth rates, current pipeline, observed churn, known hiring plans. The base case is not optimistic. It is realistic. The distinction matters because investors expect you to hit 70-80% of your base case. If your base case is already aggressive, hitting 70% of it might mean the business is in trouble.

Aggressive: The Upside Case

This is the scenario where things go better than planned. A key partnership accelerates growth. Churn drops faster than expected. A new channel outperforms. The aggressive case should be plausible, not fantastical. If your aggressive case requires capturing 40% of your TAM in 18 months, it is not a scenario; it is a fairy tale.

The aggressive case serves a specific purpose in fundraising: it shows the return potential that justifies the investment. Investors need to see a path to a 10x return in their fund model. The aggressive case is where that path lives.

Attribute	Conservative	Base	Aggressive
Revenue growth	50-60% of base	Current trajectory Churn rate	1.3-1.5x base
Conversion rate	0.7-0.8x base	Current observed Purpose	Prove survival

base, CAC at 0.8-0.9x base, hiring pulled forward 1-2 months, conversion at 1.1-1.2x base. Purpose: show return potential.

How to Build Scenarios That Actually Work

Step 1: Identify your top 5 assumptions by impact

Not all assumptions deserve scenario variation. Focus on the ones that move the output most. For most startups, these are: new customer acquisition rate, churn or retention, average revenue per customer, hiring pace and cost, and one business-specific variable (take rate for marketplaces, expansion revenue for SaaS, utilization rate for services).

Run a quick sensitivity: change each assumption by 20% and see which ones move total revenue or cash most. The top 5 by impact become your scenario variables.

Step 2: Define scenario-specific values for each variable

This is where most founders get lazy. They apply a flat multiplier to revenue instead of varying the underlying drivers. That produces scenarios that look different but are not structurally different. The correct approach: for each of your top 5 assumptions, define a conservative, base, and aggressive value. Monthly lead growth: conservative 4%, base 8%, aggressive 12%. Monthly churn: conservative 6%, base 4%, aggressive 2.5%. ACV: conservative $5,500, base $7,200, aggressive $9,000.

When you vary the inputs, the outputs change in interconnected ways that a flat revenue multiplier would never capture. Lower lead growth means fewer customers, which means less support hiring needed (lower cost), which partially offsets the revenue miss. That interplay is what makes scenario planning useful.

Step 3: Wire the scenario toggle

In your assumptions tab, add three columns next to each assumption: conservative, base, aggressive. Add a single toggle cell at the top (a dropdown or a value of 1, 2, or 3). Use an INDEX or CHOOSE formula so each assumption reads from the correct column based on the toggle. The result: changing one cell switches the entire model between scenarios. Every tab, every calculation, every output updates simultaneously. This is the structure investors want to see because it lets them test assumptions in real time during a meeting.

Step 4: Check each scenario for internal consistency

A common mistake: the conservative revenue scenario still assumes the same hiring plan as the base case. If revenue is 40% lower, you would not hire the same number of salespeople. Adjust headcount and variable costs in each scenario to reflect the operational decisions you would actually make.

Conservative scenarios should include cost cuts or deferrals that a founder would realistically implement if growth stalled. Aggressive scenarios should include the additional investment (marketing, hiring) needed to capture the upside.

Step 5: Identify the decision points

For each scenario, mark the month where a key decision must be made. In the conservative case, which month do you need to start cutting costs? In the aggressive case, which month do you need to accelerate hiring? These decision triggers are more valuable than the financial projections themselves because they turn scenarios into an operating playbook.

> **Key insight: The real value of scenario planning is not predicting the future. It is pre-deciding your responses. When an investor asks "what will you do if churn spikes to 7%?", you do not need to think on the spot. You pull up the conservative scenario and walk them through the cost adjustments, hiring deferrals, and revised milestones you have already mapped out.**

What Each Scenario Must Show

Scenario	Must Demonstrate
Conservative	Runway to a meaningful milestone without emergency raise
Conservative	Which costs get cut first and when Conservative

Base	Capital efficiency (revenue per dollar raised)
Base	Unit economics improvement trajectory
Aggressive	Return potential for investors (10x path)
Aggressive	What incremental investment captures the upside
Aggressive	Market share implied (is it realistic?)

The Conversation Scenarios Enable

During the Indeed Flex fundraising processes, scenarios transformed investor meetings from interrogations into collaborations. Instead of the investor poking holes and the founder defending, the conversation became: "Let us look at what happens if your UK take rate drops 2 points. The model shows that in the conservative case, and here is how it flows through to EBITDA and runway."

That kind of conversation only happens when the model has working scenarios. Without them, the founder says "I will have to get back to you on that," which is the fundraising equivalent of "I have not done the work."

Profounders specifically asked us to run a custom scenario during one meeting: "What if you win the enterprise contract pipeline but it takes 6 months longer than planned?" Because the model was built with scenario infrastructure, we could adjust the sales cycle assumption on the assumptions tab, toggle a modified case, and show them the answer in under two minutes. That moment built more trust than any slide in the deck.

Common Scenario-Building Mistakes

Flat revenue scaling. Conservative = base revenue x 0.7. This tells the investor nothing about what drives the miss or how you would respond. Change the underlying drivers, not the output.

Same cost structure across scenarios. If revenue drops 30%, you would not maintain the same hiring plan. Adjust costs to reflect realistic operational responses.

Aggressive case that implies impossible market share. If your aggressive case requires 20% market share in Year 2 and you are a seed-stage company, it undermines rather than supports the story. **No decision triggers.** Scenarios without action points are just three sets of numbers. Add decision markers: "In the conservative case, we defer the third engineering hire and reduce marketing spend by 30% starting Month 6."

Too many scenarios. Three is standard. Five is acceptable for specific investor requests. Seven or more suggests you do not know which ones matter.

Three-scenario planning is not a checkbox exercise. It is the structure that turns a financial model from a static forecast into a decision-making tool. Build scenarios by varying the underlying drivers, not just the revenue line. Ensure cost structures adjust realistically in each case. Mark decision triggers so each scenario becomes an operating playbook. Wire a toggle in the assumptions tab so the entire model switches with one click. The conservative case proves you can survive adversity. The base case shows your plan. The aggressive case shows the return. Together, they tell investors something more valuable than any single projection: that you have thought about what could go wrong and already know what you will do about it.

029 How to Model a Two-Sided Marketplace: Lessons From Building the Indeed Flex Financial Model

Why Marketplace Models Are Different

A SaaS company acquires one type of customer and charges them a subscription. A marketplace operates two customer types whose behavior is interconnected. A staffing marketplace like Indeed Flex needed workers (supply) and employers (demand). Without enough workers, employers would not post shifts. Without enough employers, workers would not sign up. The financial model had to capture that interdependence or it was structurally wrong.

The most common mistake founders make is modeling a marketplace like a SaaS business: project total customers, multiply by average revenue, done. This ignores the dynamics that define marketplace economics. Fill rate (the percentage of demand that gets matched to supply) is a variable that changes based on the supply-demand ratio. Take rate can compress as volume grows. Unit economics differ by side, by geography, and by segment.

> **Key insight: A marketplace model is a system model, not a customer model. You are modeling the interaction between two groups and the transactions that result. If your model does not capture the supply-demand balance, it will produce projections that look reasonable in a spreadsheet and are wrong in reality.**

The Architecture of a Marketplace Model

At Indeed Flex, the financial model was built around seven interconnected modules. Each one fed the next, and changing any single input rippled through the entire system.

Module 1: Supply Acquisition and Retention

Workers registered per month, activation rate (percentage who completed their first shift), retention at 30, 60, and 90 days, and reactivation rate for dormant workers. The supply pool at any given time was the cumulative registered workers minus churned ones, plus reactivated ones. The key insight from building this: worker retention was not a simple churn curve. Workers would go dormant for weeks and then reactivate when a shift matched their availability. The model needed to distinguish between true churn (permanently left) and dormancy (temporarily inactive), because the two required completely different strategies and had different financial implications.

Module 2: Demand Acquisition and Retention

Employers onboarded per month, conversion rate from sign-up to first shift posted, retention (percentage still posting shifts after 3, 6, 12 months), and frequency (shifts posted per active employer per month). Demand-side metrics were more stable than supply-side because employer contracts were stickier, but frequency varied significantly by segment. A large hospitality chain posted 50+ shifts per week. A small office posted 2-3 per month.

Module 3: Matching and Fill Rate

This is the module that makes marketplace modeling hard. Fill rate = shifts filled / shifts posted. It depends on the supply-demand balance in each market. A market with 500 available workers and 100 shifts per day has a very different fill rate than one with 500 workers and 1,000 shifts per day.

We modeled fill rate as a function of supply density (available workers per shift). Below a certain density threshold, fill rate dropped sharply. Above it, fill rate plateaued near 90-95%. This S-curve relationship was critical for the model because it meant that adding supply in a market with low density had a disproportionate impact on revenue.

Module 4: GMV Calculation

GMV (Gross Merchandise Value) = shifts filled x average shift value. Average shift value = hours per shift x hourly rate. This sounds simple, but hourly rates varied by role type, geography, and time of day. Night shifts and weekend shifts commanded premiums. London rates were higher than Manchester rates. The model needed to capture the mix, not just the average.

Module 5: Revenue and Take Rate

Revenue = GMV x take rate. The take rate was the percentage of each transaction that the marketplace kept as revenue. But take rate was not a constant. It varied by client segment (enterprise clients negotiated lower rates), by worker type, and by market. Over time, as the marketplace scaled and competition increased, there was pressure on take rate.

Modeling take rate as a single number would have been wrong by 3-5 percentage points depending on the mix. We segmented it by client size (SMB, mid-market, enterprise) and geography, which produced much more accurate revenue projections.

Module 6: Cost of Fulfillment (COGS)

For each filled shift: worker payment, employer liability insurance, payment processing fee, and a per-shift allocation of the operations team cost. Gross margin per shift was the take rate minus these costs. At early stages, gross margin per shift was thin, 15-20% of net revenue. At scale, it improved to 25-30% as insurance rates came down with volume and operations became more efficient.

Module 7: Unit Economics by Cohort

Each employer cohort's LTV was calculated as: average monthly gross profit per employer x average retention duration. Each cohort had its own acquisition cost (tracked by channel and month). The LTV:CAC ratio by cohort was the number that told us whether the marketplace was working, and it was the number the acquirer scrutinized most during exit diligence.

The Supply-Demand Balance: Modeling the Flywheel

The marketplace flywheel is the dynamic that makes or breaks a two-sided platform: more supply attracts more demand, which attracts more supply. In the model, this translates to:

Higher worker density in a market leads to higher fill rates. Higher fill rates lead to better employer experience. Better employer experience leads to higher employer retention and referrals. More employers lead to more shifts posted. More shifts lead to better worker experience (more work available). Better worker experience leads to higher worker retention.

Modeling this loop requires careful thought about which direction to start from. We chose to model supply first (worker acquisition and retention), then demand (employer acquisition and retention), then the matching layer (fill rate as a function of the supply-demand ratio), and finally economics (revenue, COGS, unit economics).

The flywheel creates a non-linear relationship between investment and output. Doubling marketing spend in a market that is already past the liquidity threshold might improve fill rate by 2 points. Doubling spend in a market below the threshold might improve fill rate by 15 points. The model needed to capture this non-linearity, or capital allocation decisions would be wrong.

> **Key insight: The financial model for a marketplace is also the strategic model. It tells you which markets to invest in (those closest to liquidity threshold), which to slow-play (those too far from threshold for current budgets), and which to monitor (those already at high fill rates where additional supply has diminishing returns). At Indeed Flex, the model directly informed market expansion sequencing.**

How This Model Supported a 9-Figure Exit

When Indeed acquired Syft, the financial model was the centerpiece of the valuation discussion. The acquirer's diligence team ran their own

analysis on our data, but the structure of our model gave them a framework to work within.

Three specific elements of the marketplace model made a material difference during exit:

Cohort-level employer economics. The model showed that employer cohorts were not just retained, they were spending more over time. This expansion pattern, visible only because we tracked cohort economics, was one of the key factors supporting the revenue multiple.

Fill rate dynamics by market. The model demonstrated that newer markets were still climbing the fill rate curve, meaning significant untapped value in the existing footprint. The acquirer could see where growth would come from without entering new markets.

Take rate sensitivity. When the acquirer asked what would happen if take rate compressed by 2 points, we could show the impact across every segment and geography within minutes. That responsiveness built confidence in the model and, by extension, in the management team.

Common Marketplace Modeling Mistakes

Modeling revenue without modeling both sides. If your model shows revenue growing without an explicit supply and demand layer, it is a SaaS model wearing a marketplace costume. Model each side independently and connect them through the matching function.

Using a constant take rate. Take rates evolve with scale,

competition, and client mix. Model the components that affect take rate rather than hardcoding a single percentage.

Ignoring geographic variation. A marketplace in London and a marketplace in Manchester have different supply costs, demand densities, and competitive dynamics. Aggregate models hide these differences and produce misleading unit economics.

Confusing GMV with revenue. A $10M GMV marketplace with a 15% take rate generates $1.5M in revenue. Investors have seen founders pitch GMV as if it were revenue. It destroys credibility instantly.

Not modeling the liquidity threshold. Every marketplace has a point where supply-demand density is sufficient for reliable matching. Below it, the marketplace does not work. Above it, network effects kick in. If your model does not capture this threshold, it cannot tell you how much you need to invest to get a new market to viability.

A two-sided marketplace model is a system model that captures the interaction between supply and demand. Build it in modules: supply acquisition and retention, demand acquisition and retention, matching and fill rate, GMV, revenue and take rate, COGS, and unit economics by cohort. Model the supply-demand balance explicitly because it drives fill rate, which drives revenue, which drives everything. Segment by geography and customer size. Track take rate evolution. And remember that the marketplace flywheel is a real dynamic that the model must capture, not just a slide in the pitch deck.

030 SaaS vs. Marketplace Financial Models: The Key Differences That Change Everything

The Fundamental Structural Difference

A SaaS business has one customer type who pays a subscription. The relationship is direct: the company builds a product, the customer pays to use it, recurring revenue accumulates. The financial model is built around that single customer relationship.

A marketplace has two (or more) customer types whose behavior is interdependent. Revenue comes from transactions between those groups, and the marketplace takes a cut. The financial model must capture both sides and the interaction between them.

This is not a minor distinction. It changes how you model revenue, how you calculate unit economics, how you think about growth, and what investors look for when they evaluate your numbers.

Revenue: Subscription vs. Transaction

SaaS Revenue Model

Revenue is driven by: number of customers x price per customer x retention rate. You acquire a customer, they pay monthly or annually, and they stay until they churn. New MRR comes from new logos and expansion (upsell, additional seats). Lost MRR comes from churn and contraction. Net new MRR = new + expansion - churn - contraction. The beauty of SaaS revenue from a modeling perspective: it is highly predictable once you have a few months of data. Churn rates stabilize, expansion patterns emerge, and you can forecast with reasonable confidence 12-18 months out.

Marketplace Revenue Model

Revenue = GMV x take rate, where GMV = number of transactions x average transaction value. The number of transactions depends on the activity of both sides: how many units the supply side offers and how many the demand side wants. The matching between these two sides (fill rate, conversion rate) adds a layer that does not exist in SaaS. Marketplace revenue is inherently more volatile than SaaS revenue because it depends on transaction volume, which fluctuates with seasonality, supply availability, and demand cycles. A staffing marketplace has higher volume in Q4 (holiday season). A real estate marketplace has higher volume in spring. These patterns must be modeled explicitly.

Dimension	SaaS	Marketplace	Consumer
Revenue driver	Subscriptions (predictable)	Transactions (variable)	Purchases or ad impressions
Revenue volatility	Low (contracted)	Moderate (seasonal)	High (trend-dependent)
Gross margin	70–85%	40–70%	30–60%
Key metric	ARR, NRR	GMV, take rate	DAU, ARPU
Growth lever	New logos + expansion	Supply/demand density	Virality + retention
Churn impact	Direct revenue loss	Supply/demand imbalance	Usage decline

Revenue volatility is moderate to high (volume-dependent). Gross margin on net revenue is 50-75%. Key metric is GMV, take rate, fill rate. Growth lever is liquidity on both sides. Churn impact is indirect: losing supply reduces fill rate, which reduces demand retention.

Unit Economics: Direct vs. Two-Sided

SaaS Unit Economics

Straightforward to calculate: CAC is total sales and marketing spend divided by new customers. LTV is average revenue per customer x gross margin / churn rate. LTV:CAC ratio is the benchmark that every SaaS investor knows. Payback period is CAC / (monthly revenue x gross margin). These formulas apply cleanly because there is one customer type.

Marketplace Unit Economics

Immediately more complex. You have acquisition costs for both sides. Worker acquisition cost and employer acquisition cost are separate budgets with separate channels. LTV must be calculated for the monetized side (usually demand) but supply-side retention affects demand-side LTV. At Indeed Flex, we calculated employer LTV as: average monthly gross profit per employer / monthly employer churn rate. But that LTV was only achievable if worker supply was sufficient to maintain high fill rates. If worker supply dropped, fill rates dropped, employers got a worse experience, and churn increased. The supply-side investment was a prerequisite for demand-side LTV, not separate from it.

This interconnection means marketplace unit economics cannot be calculated in isolation the way SaaS metrics can. You have to model the system, not just the customer.

> **Key insight: The biggest mistake I see in marketplace models is applying SaaS unit economics frameworks without adaptation. LTV:CAC for a marketplace must account for the cost of maintaining the other side of the market. An employer LTV:CAC of 5:1 is meaningless if you need to spend $200K per market to build sufficient worker supply. The real question is: total investment to reach liquidity in a market vs. total value extracted from that market over time.**

Cost Structure: Fixed Platform vs. Variable Fulfillment

SaaS Cost Structure

COGS is primarily hosting, infrastructure, and customer support. These scale slowly relative to revenue because software delivery has near-zero marginal cost. A SaaS company serving 1,000 customers might spend $20K/month on AWS. Serving 10,000 customers might cost $60K. That is 3x cost for 10x customers, which is why SaaS gross margins are structurally high.

OpEx is dominated by engineering (building the product) and sales & marketing (acquiring customers). The cost structure is front-loaded: you invest heavily in product and GTM, then revenue scales faster than costs. This is the fundamental SaaS value proposition.

Marketplace Cost Structure

COGS includes fulfillment costs that scale nearly linearly with transactions: payment processing, insurance, supply-side payments, logistics, quality assurance. Every transaction has a direct cost. Gross margins are structurally lower because delivery is not

zero-marginal-cost.

However, marketplaces benefit from operational leverage at scale: the per-transaction cost of matching, support, and operations decreases as the platform gets more efficient. At Indeed Flex, the cost of filling the 10,000th shift was meaningfully lower than the cost of filling the 100th shift because the matching algorithm was better, the operations team was more experienced, and insurance rates were negotiated at higher volume.

Cost Category	SaaS Behaviour	Marketplace Behaviour
COGS	Scales slowly (near-zero marginal cost)	Scales with transaction volume
Customer support	Scales with customer count	Scales with both sides of marketplace
Engineering	Front-loaded, then maintenance mode	Ongoing for matching optimisation
Sales & marketing	Scales linearly with growth targets	Must acquire both supply and demand

marginal cost). Support scales with both sides of the market. Engineering is ongoing for matching optimization. Sales & marketing must fund both sides. Gross margin starts low and improves with operational leverage.

Growth Dynamics: Linear vs. Network Effects

SaaS growth is relatively linear: invest X in sales and marketing, acquire Y customers. Double X, roughly double Y (with some diminishing returns). The growth rate is a function of investment and execution. Marketplace growth has a network effect component: each new participant on one side makes the platform more valuable for the other side. This means growth can be super-linear once you cross the liquidity threshold, but it also means growth is painfully slow before you reach that threshold.

The financial model must reflect this non-linearity. A marketplace model that shows constant percentage growth in GMV is ignoring the most distinctive feature of the marketplace business model. Instead, model the supply-demand flywheel explicitly: show how supply density drives fill rate, which drives demand retention, which drives demand-side referrals, which attracts more supply.

What Investors Evaluate Differently

SaaS Investor Focus	Marketplace Investor Focus
ARR growth rate	GMV growth rate and take rate stability
Net revenue retention	Fill rate and supply-demand balance CAC payback period
Path to profitability via margins	Path to profitability via operational leverage

different from a SaaS evaluation. They wanted to see supply density by market, fill rate curves over time, take rate by segment, and the relationship between market maturity and unit economics. A SaaS-style model showing ARR growth and NRR would not have answered their core question: does the flywheel work, and does it get stronger with scale?

Hybrid Models: When Your Business Is Both

Some businesses have elements of both. A SaaS platform with a marketplace layer (like Shopify, which charges subscriptions and takes a cut of transactions). A marketplace that adds premium subscription features for one side.

Simplifii and Haus both had hybrid elements: platform subscription revenue plus transaction-based or service-based revenue. In these cases, model each revenue stream separately with its own driver structure. Do not blend them into one revenue line because the economics, growth dynamics, and investor questions are different for each stream. The hybrid model typically requires a SaaS module (subscribers x price x retention) and a marketplace or transaction module (transactions x value x take rate), running in parallel and sharing a common cost structure. The investor presentation should clearly distinguish between the two revenue types and show the gross margin for each.

SaaS and marketplace financial models differ in every dimension: revenue mechanics, cost structure, unit economics, and growth dynamics. SaaS models are built around a single customer relationship with subscription revenue. Marketplace models are system models built around the interaction between two customer groups and the transactions that result. Applying a SaaS framework to a marketplace (or vice versa) produces a model that looks structurally correct but generates projections that are wrong in ways that investors will immediately identify. Choose the right framework for your business,

model the drivers that actually matter, and present metrics that reflect the true economics of how your company creates value.

031 Multi-Currency Financial Models: How to Run Finance Across UK, US, and Europe

What Is the Multi-Currency Problem in Startup Finance?

Scenario	FX Risk	Impact on Model
UK startup with US	GBP/USD movement	Revenue recognised in USD, customers
USD-functional company	GBP/USD movement	UK team costs fluctuate in with UK ops
EUR entity paying USD	EUR/USD movement	Vendor costs fluctuate in contracts
Multi-entity group	All cross-entity	Intercompany eliminations consolidation

The Functional vs. Presentation Currency Distinction

The first structural decision in a multi-currency model is distinguishing functional currency from presentation currency. **Functional currency** is the currency of the primary economic environment in which the entity operates. For a UK startup that earns most revenue in GBP and pays most costs in GBP, the functional currency is GBP. For a US Delaware C-corp that invoices in USD and pays its key staff in USD, the functional currency is USD, even if it has a UK subsidiary.

Presentation currency is the currency in which you report financial statements. Many international startups present in USD for investor reporting, even when the functional currency of the operating

entity is GBP or EUR. This requires translating the financial statements from functional to presentation currency at period-end rates.

> *Key insight: The distinction matters because translation from functional to presentation currency creates exchange differences that sit in equity (not the P&L). Many founders confuse these translation differences with actual operating losses or gains, which distorts the reading of the model.*

How to Build a Multi-Currency Model That Works

Step 1: Define the functional currency for each entity.

If you have a UK entity and a US entity, each has its own functional currency. Transactions within each entity are recorded in functional currency. Intercompany transactions require explicit FX translation. **Step 2: Apply a single functional currency for the model consolidation.**

For investor reporting, pick one currency — typically USD — and translate all entities into it. Apply average rates for P&L items, closing rates for balance sheet items. Flag this methodology in the assumptions tab.

Step 3: Build an FX assumption table.

Do not hard-code exchange rates into revenue or cost lines. Build an assumptions table with the FX rates you are using, their source, and the date of the rate. Use forward rates where available for forecasting rather than spot rates, since forward rates reflect market expectations. **Step 4: Create an FX sensitivity.**

A 10% movement in your main cross-currency pair is a realistic scenario over a multi-year model horizon. Show what a 10% depreciation of your revenue currency (or appreciation of your cost currency) does to gross margin and operating cash flow. This is a standard investor diligence question for any international business.

Step 5: Model intercompany charges explicitly.

If the UK entity provides services to the US entity (or vice versa), these intercompany charges need to be at arm's length for tax purposes

and eliminated at consolidation for financial reporting. Build these flows explicitly rather than netting them out informally.

The Specific Challenges by Entity Structure

UK Ltd operating with USD revenue:

Revenue recorded in USD must be translated to GBP at the transaction rate or a monthly average rate. If GBP strengthens against USD, the translated GBP revenue is lower even if USD revenue grew. This is not a business problem — it is a translation effect — but it looks like one in the P&L if it is not clearly labelled.

US Delaware C-corp with UK subsidiary:

The UK subsidiary's costs (largely GBP) are translated into USD for consolidation. A weak GBP environment reduces the USD cost of the UK team. A strong GBP environment increases it. Model the UK subsidiary's P&L in GBP, then translate at the assumption rate for the US entity consolidation.

Multi-entity with EUR operations (e.g. France, Italy, Romania): EUR-denominated costs and any EUR revenue need their own FX assumptions. The EUR/USD rate and EUR/GBP rate both matter. Build each entity's P&L in its functional currency, then translate to the presentation currency for the consolidated view.

The FX Assumptions Most Models Get Wrong

Using spot rates for multi-year forecasts. Spot rates are

appropriate for today. For years two through five of a forecast, use forward rates or an explicitly stated assumption with a rationale. "We assume stable rates throughout the forecast period" is an acceptable assumption if stated explicitly. Implicitly assuming stable rates without flagging it is a model credibility issue.

Not separating transaction exposure from translation exposure. Transaction exposure arises when you have a contractual obligation in a foreign currency (e.g. a USD invoice outstanding in a

GBP entity). Translation exposure arises when you consolidate foreign subsidiaries. These require different treatments.

Ignoring hedging costs. If the business uses FX forwards or other hedging instruments to manage currency exposure, the cost of those hedges belongs in the model. Hedging reduces volatility but is not free.

Multi-currency operations require multi-currency models. The key steps are: define functional currency for each entity, use a single presentation currency for investor reporting, build an explicit FX assumption table with sourced rates, run a 10% sensitivity on major cross-currency pairs, and model intercompany charges explicitly. The most common errors are using spot rates for multi-year forecasts, conflating translation effects with operating performance, and ignoring the cost of hedging. Get the structure right and the model will reconcile to actual bank balances across currencies — which is the final test of whether it has been built correctly.

032 How to Model Gross Margin for Different Business Models

Why Gross Margin Modelling Varies by Business Model

Gross margin measures the revenue remaining after the direct costs of delivering your product. Because "direct costs of delivery" means something different for a SaaS company than for a marketplace or a professional services firm, the structure of the COGS model and the resulting gross margin benchmarks differ significantly across business types.

SaaS Gross Margin: What to Include

For a pure software-as-a-service business, COGS typically includes: Cloud hosting and infrastructure (AWS, GCP, Azure) scaled with customer count or usage

Third-party API costs consumed per customer (payment processors, mapping services, communication APIs)

Payment processing fees on each transaction

Customer success salaries attributable to direct customer onboarding and support (not account growth)

Amortisation of capitalised software development costs if applicable Product engineering salaries (building features) belong in R&D (OpEx), not COGS, unless they are maintaining the delivery infrastructure. Getting this wrong artificially raises gross margin.

The key modelling question for SaaS is whether hosting and infrastructure costs are variable (per customer or per usage) or largely fixed (flat regardless of customer count). Variable infrastructure costs compress gross margin as volume grows until economies of scale kick in. Fixed infrastructure costs mean gross margin improves significantly as revenue scales.

Marketplace Gross Margin: What to Include

Marketplace COGS depends on whether the business takes transaction risk or simply facilitates transactions.

Pure take-rate model (no fulfilment risk): COGS includes payment processing, fraud and trust and safety operations that scale with transaction volume, and any insurance costs tied to individual transactions. Gross margins in this model are typically high (60-80%) because the primary cost is just friction on the transaction, not delivery of the underlying service.

Marketplace with operations (e.g. staffing, logistics, services): COGS includes the cost of the service being fulfilled, direct operations staff managing the supply side at volume, compliance and verification costs per transaction, and any insurance or indemnity tied to individual placements. Gross margins here are lower (35-60%) because the business is taking on more of the economic risk of the transaction.

> *For a staffing marketplace like Indeed Flex, the key COGS lines were worker pay (the largest single line), employer obligations (NI, holiday pay), compliance and verification operations, and payment processing. The resulting gross margin reflects the spread between the rate charged to employers and the rate paid to workers plus the direct costs of managing that relationship.*

Professional Services Gross Margin: The Utilisation Model

Professional services gross margin is primarily driven by two factors: the billing rate per hour and the utilisation rate (percentage of billable hours actually billed).

Gross margin formula for professional services:

Gross Margin = (Revenue - Direct Delivery Costs) / Revenue Direct Delivery Costs = Staff Time Costs + Direct Project Costs Staff Time Costs = (Senior FTE cost × hours) + (Junior FTE cost × hours) The key

lever is the ratio of junior to senior staff on deliverables. Higher junior leverage (more junior staff per senior) typically improves gross margin at the cost of quality risk. Higher senior intensity improves quality but compresses margin.

In the model, build utilisation assumptions explicitly. A team of 10 that is 70% utilised on billable work produces very different economics than one that is 90% utilised. And the trajectory — whether utilisation improves as the team matures and processes stabilise — drives the gross margin trend in the forecast.

How to Model Gross Margin Trajectory at Scale

For every business model, gross margin should improve as scale increases. The question is by how much and through which mechanism. **SaaS:** Gross margin improves as fixed infrastructure costs are spread over more customers. Model this explicitly by separating fixed and variable COGS components and showing the fixed cost dilution effect per customer.

Marketplace: Gross margin improves as payment processing rates improve with volume (negotiated rates) and as operations costs per transaction decline through process automation and scale. Show these as explicit assumptions with a rationale.

Professional services: Gross margin improves as utilisation increases and as the junior-to-senior ratio in delivery is optimised. Build utilisation as an explicit driver.

The model should show the gross margin trajectory over the forecast period, with clear attribution of what drives each change. An investor who sees gross margin improving from 45% to 65% over five years without an explanation of the mechanism will ask. Build the explanation into the model before they do.

Common Gross Margin Modelling Mistakes

Flat gross margin across the entire forecast period with no explanation

Engineering costs in COGS for product development work

Customer success costs entirely in Sales and Marketing with nothing in COGS

No separation of fixed vs. variable COGS components

Gross margin that significantly outperforms the business model benchmark with no explanation

Gross margin varies by business model and should be modelled with that context in mind. Build COGS using the delivery test, benchmark the output against comparable business models, separate fixed and variable components, show the trajectory and explain the mechanism for improvement, and know the answers to the questions investors will ask about any line that diverges from benchmark. Gross margin is not just a number — it is the clearest signal of how efficiently your business delivers value relative to what it costs to deliver it.

033 Sensitivity Analysis: Pre-Building the Questions Investors Will Ask

What Is Sensitivity Analysis in a Startup Financial Model?

Sensitivity analysis shows how a change in one or more input assumptions flows through to a change in the model's key outputs — typically runway, gross margin, and revenue. It is distinct from scenario planning (which changes multiple assumptions together to reflect different business environments) and is instead focused on isolating the impact of individual variables.

Key distinctions:

Sensitivity Analysis	Scenario Analysis
What	One variable at a time
Purpose	Identify which inputs drive the Show coherent future states most variance
Output	Tornado chart or sensitivity
Best for	Identifying risk concentration

Why Investors Use Sensitivity Analysis in Diligence

When a VC analyst goes through a financial model in diligence, one of the standard checks is to change the key assumptions by 10-20% and see what happens to the bottom line. This is not adversarial. It is how experienced investors understand which assumptions the model is most dependent on.

If the model has not been built with this flexibility in mind — if assumptions are hard-coded rather than centralised — the analyst

cannot run the test and flags this as a model quality issue. If the model has sensitivity analysis pre-built, the analyst can see the founder has already done the work and understands their own key risks. The question "what are the two or three assumptions your model is most sensitive to?" is one of the most common investor diligence questions. Pre-building the sensitivity is the equivalent of pre-building the answer.

The Four Most Important Sensitivities for Early-Stage Models

1. CAC (Customer Acquisition Cost)

If CAC is 20% higher than modelled, what happens to marketing spend, to the CAC payback period, to LTV:CAC, and to monthly burn? For most early-stage businesses, CAC is one of the most uncertain and highest-impact variables. Build a sensitivity that shows the impact across a range of CAC scenarios.

2. Churn rate

If monthly churn is 1 percentage point higher than modelled, what happens to net revenue retention, to total customer count at year-end, and to the revenue line? For SaaS businesses especially, the compounding effect of churn means a small change in churn rate produces large changes in the revenue line over 18-24 months.

3. Revenue growth rate

If the core growth assumption is 20% below plan, what is the revised runway? What is the new point at which the company needs its next round? This is often the most powerful single sensitivity because it directly addresses the question investors most want answered: what is the downside scenario?

4. Gross margin

If COGS are 10% higher than modelled (due to higher infrastructure costs, pricing pressure, or operational inefficiency), what happens to gross margin and operating leverage? For businesses targeting high gross margins, the sensitivity shows how robust the unit economics are to cost pressure.

How to Build a Sensitivity Table in Excel

A sensitivity table (or data table in Excel terminology) shows the output of a model formula across a range of input values. The structure is straightforward:

One-variable sensitivity (1D data table):

Column header: range of values for the input (e.g. CAC from $200 to $400 in $25 increments)

Row output: the metric being measured (e.g. 18-month runway, gross margin)

Built using Excel's Data > What-If Analysis > Data Table

Two-variable sensitivity (2D data table):

Row header: range of values for input A (e.g. CAC)

Column header: range of values for input B (e.g. churn rate) Cell values: output metric (e.g. monthly burn at month 18) Built the same way but using both row input cell and column input cell

The key requirement for both: the input must flow from a single cell reference in the model (not be hard-coded in multiple places). Centralised assumptions tabs are what make sensitivity analysis possible without rebuilding.

> *Key insight: Sensitivity analysis is only possible if the model is built with centralised assumptions. A model where growth rates, CAC, and churn are hard-coded in multiple tabs cannot be stress-tested efficiently. The assumptions tab is what makes sensitivity analysis fast.*

The Tornado Chart: Visualising Which Inputs Matter Most

A tornado chart ranks input variables by the size of their impact on a key output (typically revenue, runway, or gross margin). The widest bar is the variable the business is most sensitive to. The narrowest is the least critical.

Building a tornado chart requires running a high and low test for each variable (e.g. ±20%) and recording the change in output. The result is a ranked list that immediately communicates the model's key risk concentrations.

For most early-stage companies, the ranking looks roughly like this:

1. Revenue growth rate (widest bar, highest impact)
2. Churn rate (for subscription businesses)
3. CAC (for sales-led or performance-marketing-driven businesses) 4. Gross margin
5. Headcount growth (typically lower impact in early models) A tornado chart in the investor deck communicates two things simultaneously: the founder has done the work of understanding their key uncertainties, and the business is not dependent on a single fragile assumption.

Common Sensitivity Analysis Mistakes

Running sensitivities on outputs that are already robust, ignoring the fragile ones

Hard-coding assumptions that prevent efficient testing

Showing only upside sensitivities (what happens if CAC is lower than expected)

Not including sensitivity results in investor materials, doing the analysis only internally

Sensitivity ranges that are too narrow to reflect realistic uncertainty

Sensitivity analysis turns a financial model from a static forecast into a tool for understanding risk. The four variables most worth testing in early-stage models are CAC, churn, revenue growth rate, and gross margin. Build sensitivity tables using centralised assumption cells, visualise the results in a tornado chart for investor materials, and use the analysis to pre-answer the questions investors will ask. A founder who has run their own sensitivities and can speak to the results is in a structurally stronger position in any diligence conversation than one who has not.

034 How to Audit Your Financial Model Before Sharing It With Investors

Why Auditing Your Own Model Matters

Founders who share unaudited models make a specific kind of mistake. The model goes to a diligence analyst who finds a formula error on the revenue tab, a cash balance that does not reconcile to the P&L, or an assumption that is inconsistent across two different tabs. The technical error may be minor. The signal it sends is not.

A model error tells an investor one or more of the following: the founder has not reviewed this themselves, the numbers should not be trusted at face value, there may be other errors not found yet. None of these are the signal you want to send when someone is deciding whether to write a check.

Audit checklist summary:

Category	**Time

Step 1: Structural Integrity Check

Before reviewing any numbers, verify the model's structural health. **Circular references:** A circular reference occurs when a formula refers back to its own cell, directly or through a chain of references. Excel will flag these with a warning. They should be eliminated or, if intentional (e.g. in an interest model), explicitly documented. Circular references in unintended places are a sign the model has not been reviewed.

Broken links: If the model references external files, check whether those links are live. Broken external links return errors that cascade through dependent formulas.

Error values: Scan every tab for #REF!, #VALUE!, #DIV/0!, and #NAME? errors. These indicate broken formulas. Address each one before the model leaves the building.

Tool: Excel's Formulas > Error Checking function surfaces all errors in the active sheet. Run it on every tab.

Step 2: Mathematical Consistency

The most important mathematical check is the three-statement reconciliation: the P&L, the cash flow statement, and the balance sheet must tie together.

The cash reconciliation test: The ending cash balance in the cash flow statement must match the cash line on the balance sheet for every period. If they diverge, there is an error somewhere in the model structure.

The retained earnings test: The cumulative net income from the P&L must match the change in retained earnings on the balance sheet. If they do not reconcile, equity movements have not been captured correctly. **The P&L to operating cash flow bridge:** Starting from net income, the bridge to operating cash flow must account for all non-cash items (depreciation, SBC) and working capital movements. The bridge should reconcile exactly.

> *Key insight: A model that passes the three-statement reconciliation test has almost certainly been built correctly in structure, even if individual assumptions are wrong. A model that fails it has a structural error that will be found in diligence.*

Step 3: Assumption Centralisation

Every key assumption in the model should live in exactly one place: the assumptions tab. When reviewing the model before sharing, the test is straightforward: open the assumptions tab, find the CAC assumption, and trace the formula chain from that assumption cell through to the revenue model.

If the formula chain breaks — if you find the CAC value hard-coded somewhere else in the model, or if the number in the model differs from the assumptions tab — you have found a centralisation failure. These failures mean that changing an assumption in one place does not change it everywhere, which produces inconsistencies that analysts find during diligence.

What to look for:

Any numeric constants in formula cells that should be references to the assumptions tab

Assumptions that appear in multiple places with slightly different values

Key inputs that are not in the assumptions tab at all

Step 4: Internal Consistency

Internal consistency means the model tells one coherent story across all tabs. Inconsistencies — where the headcount plan does not support the revenue assumptions, or the gross margin calculation uses different COGS than the P&L — are often more damaging than simple formula errors because they suggest the model was assembled rather than designed. **Key consistency checks:**

Headcount to revenue: If revenue grows 3x, does headcount grow enough to support that? If the model shows headcount flat while revenue triples, either the model assumes the product scales with no marginal operational support (which needs explicit justification) or the headcount plan has not been connected to the revenue model. *COGS consistency:* Does the COGS build on the gross margin tab match the COGS line on the P&L? These should be the same formula, not separate calculations.

Revenue bridge: Does the revenue from the detailed revenue build match the revenue line on the P&L? This sounds obvious but is frequently wrong in models assembled from multiple sources.

Step 5: Logic Review

The final audit step is a logic review — reading the model as an investor would and asking whether each assumption and output is credible.

Growth rate credibility: Is the growth rate in year one consistent with current traction? Is the growth rate in year five credible given the market size and the competitive environment? Investors who see hockey-stick models with no explanation of what changes in year two will ask.

Margin trajectory: Does gross margin improve over time, and is the mechanism explained? Does operating leverage appear in the model (OpEx growing slower than revenue at scale)?

Sign checks: Revenue should be positive. Costs should be negative (or shown as positive in a cost-only view, consistently). Mixed sign conventions are a common source of confusion and errors.

Runway check: Does the cash runway calculation use the cash flow statement or the P&L? It must use cash flow. A runway number derived from the P&L is wrong for any business with working capital.

The 90-Minute Audit in Practice

The full audit described above takes two to three hours for a complex model. For a standard early-stage model, a focused 90-minute audit covers the most important ground:

1. Run error checking on every tab (20 min)
2. Verify cash reconciliation (15 min)
3. Test assumptions centralisation on five key inputs (20 min) 4. Check headcount-to-revenue and COGS-to-P&L consistency (20 min) 5. Read the model as an investor: does the story make sense? (15 min) This is the minimum before the model goes to any serious investor.

Auditing a financial model before sharing it is not optional. It is the final step between building a model and using it as an instrument for raising capital. The five checks — structural integrity, mathematical consistency, assumption centralisation, internal consistency, and logic

review — take less than three hours and eliminate the most common model credibility killers. A model that passes this audit does not guarantee a successful raise. But a model that fails it in diligence can stop one.

035 Excel vs. Google Sheets for Startup Financial Modeling: The Honest Answer

The Honest Comparison

The Excel vs. Google Sheets debate for startup financial modelling has a pragmatic answer: use whichever tool you will build the best model in, and know the limitations of each before you commit.

Both can produce excellent startup financial models. Both can produce terrible ones. The tool is not the constraint. The thinking behind the model is.

Direct comparison:

Feature	Excel	Google Sheets
Performance with	Handles complexity better	Slows significantly large models
Collaboration	Requires OneDrive/SharePoint	Native real-time or manual sharing
Formula library	Larger, more advanced	Smaller, catching up (XLOOKUP, dynamic arrays)
Investor	Standard for formal diligence Acceptable at seed; expectation	less common at Series A+
Version control	Manual (save copies)	Automatic revision history
Offline access	Full functionality	Limited without sync AI integration

Cost	Microsoft 365 from ~£8/mo	Free with Google account

When Google Sheets Is the Right Choice

Early-stage models (pre-seed through seed): At this stage, the model is relatively simple — revenue build, cost structure, cash runway, basic three-statement. Google Sheets handles this comfortably and the collaboration benefits are real. Multiple stakeholders (cofounder, advisor, fractional CFO) can work simultaneously without version conflicts.

When the model is being actively built with collaborators: Real-time collaboration in Google Sheets is genuinely better than anything Excel offers outside Microsoft 365 co-authoring. If the model is being built iteratively with input from multiple people, Google Sheets reduces friction.

When the investor is not yet formal: Seed investors who are assessing the business at a high level often receive Google Sheets models. The format is not a concern at this stage.

When Excel Is the Right Choice

Series A and beyond: At Series A, diligence is more thorough. External advisors and analysts may work on the model. Excel is the professional standard for this environment. Sending a Google Sheets link to a Series A firm's diligence team is not a disqualifier but it is unusual.

Complex multi-tab models with many formula dependencies: Excel handles computational load better. A model with cohort analysis, three-statement accounting, sensitivity analysis, scenario toggles, and a data room summary — running across 15+ tabs with thousands of formula cells — will perform more reliably in Excel.

When the investor specifies a format: Some investors and their advisors have tools built around Excel. If they ask for an Excel file, send an Excel file.

When data tables and advanced analysis are needed: Excel's Data Table function for sensitivity analysis has no direct equivalent in Google Sheets. Power Query, pivot table complexity, and dynamic array functions are all more capable in Excel.

The Formula Gaps That Matter

Google Sheets has been catching up to Excel rapidly but there are still meaningful formula gaps for financial modelling:

XLOOKUP: Available in Excel 365, not natively in Sheets (though VLOOKUP/INDEX-MATCH work as alternatives). XLOOKUP's error handling is cleaner for model building.

Dynamic arrays (FILTER, SORT, UNIQUE): Excel's dynamic array functions are more powerful for building flexible summary tables. Sheets has equivalents but with different syntax and limitations. **Data Tables:** Excel's What-If Analysis > Data Table has no direct Google Sheets equivalent. For sensitivity analysis built natively in the spreadsheet, Excel is the only option.

Power Query: For connecting to external data sources or automating data cleaning in the model, Excel's Power Query has no Sheets equivalent. This matters more for operational models than for investor-facing models.

The Practical Recommendation

Build early drafts in whichever tool you think in. If you are more comfortable in Google Sheets, start there.

Before any formal investor process, convert to Excel. The conversion takes less than 30 minutes for most models (File > Download > .xlsx in Google Sheets). Review all formulas after conversion — some Sheets functions do not translate directly to Excel and will produce errors. Build the investor-facing model in Excel with a named assumption tab, clean formatting, and a summary tab. Keep a Google Sheets version for ongoing collaboration with your team if that is where they work best.

The version that goes to investors should be in Excel unless you have specific context that Sheets is fine for this particular process.

Excel is the professional standard for investor-facing financial models. Google Sheets is a better collaboration tool for iterative model building with a team. Use Google Sheets for early drafts and active collaboration. Convert to Excel before any formal investor process. Know the formula limitations of each tool before they become problems in a diligence context. The model's quality is determined by the thinking behind it, not the tool it is built in — but the tool you choose should not create unnecessary friction in the process.

036 The Balance Sheet Founders Skip (And Why That's a Mistake)

Why the Balance Sheet Exists

The balance sheet answers a single question: at this exact point in time, what does the company own, what does it owe, and what is left over for shareholders? Every other financial statement is about movement across time. The P&L shows revenue and costs over a period. The cash flow statement shows how cash moved in and out. The balance sheet is different. It is a photograph of financial position at a specific date. The fundamental equation that governs every balance sheet ever produced is:

Assets = Liabilities + Equity

This equation cannot be violated. If it does not hold, there is a mistake somewhere in the model. This is what makes the balance sheet the most powerful validation tool in financial modeling. A model that balances is not automatically correct, but a model that does not balance is definitely wrong.

What Goes on a Startup Balance Sheet

Most startup balance sheets are simpler than founders expect. The complexity comes later, at growth stage, with stock-based compensation accounting, capitalized development costs, and complex debt structures. At pre-seed through Series A, the balance sheet typically has six to ten line items.

Current Assets

Current assets are assets expected to be converted to cash within twelve months. Cash and cash equivalents: your actual bank balance. This number must match the ending balance in the cash flow statement exactly. If it does not, there is an error. Accounts receivable: revenue that has been recognized on the P&L but not yet collected from

customers. If you have net 30 payment terms and customers pay on average in 35 days, you will carry a receivable balance. A SaaS company billing monthly with immediate payment collection will have minimal receivables. Prepaid expenses: costs you have paid in advance that have not yet been expensed. Annual software subscriptions paid upfront are a common example. You pay in January, but the expense is recognized monthly across the year. At the end of January, eleven months of that subscription sit as a prepaid on the balance sheet.

Non-Current Assets

Non-current assets are held for longer than twelve months. Capitalized software development costs: if your engineers are building the core product, some of that cost can be capitalized under accounting standards rather than fully expensed in the period. This is more relevant as you scale and matters for GAAP-compliant reporting. At early stage, many startups expense all development costs. Equipment: laptops, servers, physical assets. Depreciated over their useful life. Right-of-use assets: if you have an office lease, under modern accounting standards (IFRS 16 / ASC 842), the lease creates a right-of-use asset on the balance sheet.

Current Liabilities

Accounts payable: expenses that have been recognized on the P&L but not yet paid to suppliers. If you accrue for vendor invoices with 30-day terms, those balances sit here until paid. Deferred revenue: one of the most important and most frequently missed balance sheet items for startups. If a SaaS company collects an annual subscription upfront in January, only one month of that subscription has been earned by the end of January. The remaining eleven months are a liability. The company owes the customer eleven more months of service. Accrued expenses: costs incurred but not yet invoiced. Payroll accruals for the last few days of the month, for example. Short-term portion of debt: any principal due within twelve months on term loans or venture debt.

Non-Current Liabilities

Long-term debt: venture debt, term loans, or other borrowings with repayment beyond twelve months.

Equity

Paid-in capital: the cumulative amount invested by shareholders. Every equity round increases this number. Accumulated deficit (or retained earnings): the running total of all net income and losses since the company was founded. For venture-backed startups that are burning cash to grow, this will be a large negative number.

Why Founders Skip It and Why That Is Expensive

There are three reasons founders omit the balance sheet from their financial models. The first is that it feels redundant. "The P&L shows the costs and the cash flow statement shows the cash. What does the balance sheet add?" The answer is that it adds structural integrity. A model that includes a balance sheet that balances is a model where the three statements are correctly connected to each other. Most financial model errors manifest as a balance sheet that will not balance. The second reason is that it is harder to build. The P&L is relatively intuitive. The cash flow statement is more complex but follows a clear logic. The balance sheet requires understanding how transactions flow between statements, and that takes some time to learn. But "harder to build" is not a reason to omit a component that investors expect to see at Series A. The third reason is that errors hide in other statements and only reveal themselves in the balance sheet. A founder who skips the balance sheet can feel confident the model is correct while carrying a systematic error that would be instantly visible if the balance sheet were present. Skipping it does not remove the error. It removes the ability to detect it. What do investors actually look for when they open the balance sheet? Three things primarily. First, does it balance? An investor who opens a model and sees a balance sheet where assets do not equal liabilities plus equity will stop trusting every other number in the model. Second, what is the working capital position? Working capital is current assets minus current liabilities. A startup with negative working capital is spending cash faster than it is collecting it, even if the P&L shows a small profit. Growing businesses often have negative working capital because receivables lag payables and payroll. Understanding this dynamic and modelling it correctly matters. Third,

deferred revenue. A large and growing deferred revenue balance is actually a positive signal. It means customers are paying upfront, the company is collecting cash before recognizing revenue, and cash timing is favorable relative to accounting profit. This is one of the reasons SaaS businesses with annual prepaid contracts carry negative working capital despite strong unit economics.

How to Build the Balance Sheet Step by Step

Step 1: Start with equity

The equity section should be built first because it derives directly from the other statements. Paid-in capital increases each time an equity round is raised. Accumulated deficit starts at zero and decreases by the amount of net loss each period, as pulled directly from the P&L.

Step 2: Add liabilities

Accounts payable derives from the payment timing assumptions in your cost model. If suppliers are paid on 30-day terms and monthly costs are GBP 200k, you will typically carry roughly GBP 200k in accounts payable at any given time. Deferred revenue requires careful calculation. Take cumulative annual subscriptions collected in the period, subtract the portion that has been earned based on how many months of service have been delivered, and the remainder is the deferred revenue liability. Many SaaS models get this wrong by either omitting it entirely or confusing cash collected with revenue recognized. Accrued expenses accumulate through the period and are reset when paid. If you run a monthly payroll with a cutoff on the 25th, the last five days of each month generate an accrued payroll liability.

Step 3: Build current assets

Cash must equal the ending balance from the cash flow statement. There is no version of the model where these two numbers should differ. Accounts receivable flows from your revenue model and payment timing. If 20% of revenue is collected 30 days after invoicing, the receivables balance at month end equals approximately 20% of that month's revenue.

Step 4: Add non-current assets

For most early-stage software startups, this section is small. If you are capitalizing any development costs, the balance here grows by capitalized amounts and decreases by amortization charges.

Step 5: Test the balance

Build a check cell: Assets minus (Liabilities plus Equity). This cell must equal zero. If it does not, work backward through each section until you find the mismatch. Common culprits are deferred revenue not linked to the revenue model, equity not updated for new rounds, or accumulated deficit not pulling from the P&L correctly.

The Most Common Balance Sheet Errors

Deferred revenue omitted. Extremely common in SaaS models. A company collecting GBP 1M in annual subscriptions upfront has a GBP 916k deferred revenue liability in month one. Leaving this out makes the balance sheet impossible to balance correctly and overstates how much revenue has been earned. **Equity rounds not in paid-in capital.** The cash from a fundraise appears in the cash flow statement under financing activities. It must also appear as an increase in paid-in capital in the equity section. Founders who update the cash flow statement but forget to update equity will see the balance sheet go out of balance by exactly the amount of the round raised. **Accumulated deficit not formula-driven.** If the accumulated deficit is hardcoded rather than pulling cumulatively from the P&L net income line, any future P&L change will not flow through to equity. The balance sheet becomes decorative rather than functional. **Receivables not connected to revenue timing.** If the model has payment terms assumptions, receivables should change as those assumptions change. A model where cash collected always equals revenue recognized, but payment terms are stated as net 30, is internally inconsistent.

Chapter Exercise

Add a balance sheet tab to your financial model. Build the equity section first by pulling paid-in capital from your funding history and

accumulated deficit from the cumulative net income total in your P&L. Then build the liabilities section, making sure deferred revenue is linked to your revenue recognition timing. Add current assets, ensuring cash ties exactly to your cash flow statement. Build a balance check cell. If the model does not balance, work backward through each section to find where the disconnect is. This exercise will surface any structural errors in your three-statement model before an investor does it for you.

037 How to Model Pre-Revenue Startups

What Investors Are Actually Evaluating

When a VC opens a pre-revenue financial model, they are not trying to assess the accuracy of the revenue forecast. They know the numbers will be wrong. Every early-stage investor knows this. What they are assessing is the quality of the thinking. Specifically, they are asking:

Does this founder know what it costs to build the product they are describing? Does the timeline to first revenue reflect how long product development actually takes, or does it assume everything goes to plan? Have they talked to real customers, or are the pricing and conversion assumptions pulled from a blog post? Do they understand which assumptions drive the most variance in outcomes? Have they thought about what happens if the plan is wrong? A pre-revenue model that answers those questions earns credibility. One that does not answers them by omission.

The Three-Section Structure

Pre-revenue models work best when organized into three distinct sections: the cost model, the revenue model, and the milestone and runway output. These sections serve different purposes and should be built in this order. **Section 1: The Cost Model**

Build costs first. This is the most reliable section of any pre-revenue model because costs are largely within the founder's control. Costs also represent the strongest evidence of operational realism. A founder who has thought carefully about what it costs to build and run the business they are describing is demonstrating operational literacy. The cost model should include every category of spend with sufficient detail to be credible. Headcount is the most important line item. Build it role by role. List each hire, their planned start date, their gross salary, and the associated employer costs (national insurance, benefits, payroll taxes,

and any Employer of Record fees if hiring internationally). A Romanian developer hired through an EOR carries roughly 8 to 12% employer overhead. A French employee carries 40 to 47%. A US employee varies significantly by state. These differences matter at small team sizes. Include the timing lag between the decision to hire and the point at which the person is productive. For a senior engineer, the hiring process alone typically takes eight to twelve weeks. An experienced enterprise sales hire may take three to four months to close and another three to four months to reach full productivity. Your headcount plan needs to reflect this, or your revenue ramp will be optimistic. Infrastructure and tooling costs are often underestimated. List the specific services you depend on: cloud hosting, monitoring, analytics, CRM, security tooling, product analytics. Price each one at the scale you expect to reach by the end of the forecast period, not at the current free tier. Professional services tend to surprise founders who have not raised before. Legal costs for incorporation, shareholder agreements, and option plan documentation typically run GBP 8,000 to GBP 20,000 depending on complexity and jurisdiction. Annual accounting and audit costs start at GBP 5,000 to GBP 8,000 for early-stage companies and increase with transaction volume. Go-to-market costs need a timeline. When do you start spending on sales and marketing? What are the channels? What are the expected conversion rates at each stage of the funnel? If you cannot answer these questions, the go-to-market cost section is not ready. **Section 2: The Revenue Model**

Revenue in a pre-revenue model is a hypothesis, and it should be labelled as one. The job of this section is not to predict revenue accurately. It is to document the specific assumptions you are making about how revenue will develop, why those assumptions are reasonable, and what the key drivers are. For a B2B SaaS company, the revenue model typically flows through a small number of connected assumptions: how many leads enter the pipeline per month, what is the conversion rate from lead to qualified opportunity, what is the conversion rate from qualified opportunity to closed deal, and what is the average contract value. Each of those four numbers should have a note explaining where it comes from. "Conversion rate from qualified

opportunity to close: 25% based on conversations with two enterprise sales leads who have operated in this market, adjusted downward from a typical 30% benchmark to reflect our early-stage positioning" is a model assumption. "Conversion rate: 25%" is a placeholder. Churn and expansion should both be present even at pre-revenue stage, at minimum as flagged assumptions. You will not know your actual churn rate yet. But stating that you are assuming annual churn of 8% because comparable B2B SaaS products at this price point typically see 5 to 12% annual churn shows you have thought about it. Not modelling churn at all is a signal. If you are a marketplace business, you need to model both sides of the market. The supply-side assumptions (number of service providers, activation rate, utilization rate) drive the supply economics. The demand-side assumptions (number of customers, booking frequency, average transaction value) drive the revenue. The balance between supply and demand is the most important assumption in any marketplace model and the one most often oversimplified. **Section 3: Runway and Milestones**

This is the output section that investors spend the most time on. It should answer three questions directly. How long does the capital being raised last under the base case? Under the conservative case? Under the worst case? What specific milestones will be reached with this capital? Not vague descriptors ("build product," "acquire customers"), but specific and measurable outcomes: ship the core product to ten pilot customers by month eight, reach GBP 50k MRR by month fourteen, close two enterprise contracts by month sixteen. What metrics, achieved by when, create the conditions for the next fundraise? The milestone section is where founders most often underdeliver. Milestones should be tied to the cost model (these are what the capital funds) and the revenue model (these are the commercial outcomes those costs are designed to produce). A milestone document that is disconnected from the financial model is not a financial model output. It is a slide deck.

The Most Important Pre-Revenue Assumption

Time to first revenue is the single assumption that drives more variance than any other in a pre-revenue model. It determines when cash starts coming in, how long the company is fully burning without offset, and how much runway remains when the company becomes commercially active. Most pre-revenue models underestimate time to first revenue. The most common reasons are: product development takes longer than planned (software almost always does), the sales cycle is longer than assumed because the product is not yet proven, onboarding and integration for the first customers takes time that the model did not account for, and pricing conversations take longer than expected once a real dollar figure is on the table. A useful rule: take your estimated time to first revenue, apply a 30 to 40% buffer, and check whether the round you are raising still provides enough runway in that scenario. If it does not, you either need to raise more capital or run leaner until first revenue arrives.

Scenario Modeling for Pre-Revenue Companies

Every pre-revenue model needs at least three scenarios, and they need to be built from different assumption sets rather than from multiplying a single revenue number up or down. **Base case:** assumptions as modelled. Time to first revenue at the planned date. Hiring plan as specified. Revenue ramp as per the model. **Conservative case:** first revenue delayed by three to six months. Two key hires delayed or replaced with more junior alternatives. Revenue in year one at 50 to 60% of the base case. **Capital efficiency case:** focused on extending runway. Minimum viable headcount. Aggressive prioritization of only the activities required to reach the milestone that unlocks the next round. This scenario answers the question: if everything takes longer and costs more than planned, how long can the company survive on this capital? The question investors are asking when they look at the conservative and capital efficiency cases is not whether the business fails in those scenarios. It is whether the founder has thought about them. A founder who has only modelled the base case has only thought about the best case.

What Not to Model

Do not model revenue starting in month one unless the product already exists and customers are in active conversations. Investors who have funded pre-revenue companies know what the first six months look like. Immediate revenue signals either that the timeline is wrong or that the founder has not been honest about the current state of the product. Do not present a single revenue scenario as a forecast. At pre-revenue stage, anyone claiming precision in revenue forecasting is either not being honest or does not understand the uncertainty inherent in the exercise. Do not model headcount productivity at 100% from day one. Build in a ramp: a new sales hire who starts in month three typically contributes meaningful revenue by month six or seven, not month four. A new engineer who joins mid-sprint contributes full velocity after three to four weeks of context.

FAQ

How detailed should a pre-revenue model be?

Detailed enough to answer the core questions without false precision. A 50-row model with honest assumptions and clear notes is more fundable than a 200-row model built on inputs that cannot be defended. Build the level of detail you can actually explain and defend in a 30-minute investor conversation. **Should pre-revenue models include a P&L?**

Yes. Even with zero revenue, the P&L shows the cost structure and the trajectory toward breakeven. It also forces you to think about gross margin, which becomes an important assumption as soon as revenue appears in the model. **How should a pre-revenue model handle taxes?**

At early stage, most pre-revenue companies carry accumulated losses that fully offset any future tax liability for years. Include a tax line, assume zero tax liability for the forecast period given the accumulated deficit, and note the assumption.

Chapter Exercise

Build the three-section pre-revenue model for your business. Start with the cost model: list every hire, every tool, every fixed cost by month for the next 24 months. Then build the revenue model with an explicit assumptions tab where every input has a note explaining its basis. Finally, build the runway calculation under three scenarios. For each scenario, identify the single most important assumption that determines whether the company reaches the next fundraising milestone.

038 Financial Model Red Flags: What Breaks Investor Confidence

Why Models Lose Investor Confidence

Before the investor asks a single question, the model has already communicated something. A model that is well-structured, internally consistent, and built from explicit assumptions communicates that the founder has thought carefully about their business. A model that contains red flags communicates the opposite. Importantly, the red flags that matter most are not technical errors. Most technical errors in financial models are detectable and fixable. The red flags that kill deals are the structural ones: projections that imply the founder does not understand their economics, assumptions that reveal they have not spoken to real customers, and a single-scenario model that suggests they have not considered what happens when things go wrong.

Red Flag 1: The Unexplained Hockey Stick

The hockey stick growth curve is not a red flag in itself. Many businesses legitimately show slow early growth followed by a period of acceleration. A SaaS business that achieves product-market fit in year one and shifts from founder-led to a repeatable outbound sales motion in year two can credibly show acceleration in year three. The model should name that inflection point, explain the mechanism, and show it reflected in the headcount plan and cost structure. The red flag is a hockey stick with no explanation. Revenue doubles in month eighteen with no corresponding change in sales headcount, no new channel coming online, no pricing change, no seasonality. Just a curve that bends upward because the spreadsheet formula was written that way. When an investor sees an unexplained inflection, the first move is to change the formula. What does the model look like if the growth rate in year two matches year one? If the answer is that the model raises significantly less capital than the round being pitched, the red flag has become a structural problem. The fix is straightforward: every period

where the growth rate materially changes, there should be an explanation in the assumptions tab. Not a vague statement. A specific mechanism. "Hire first sales leader in month 12, first two SDRs in month 14, expecting 3x increase in outbound pipeline by month 18."

Red Flag 2: Gross Margin That Beats the Benchmark Without Justification

Every business model has a gross margin range that sophisticated investors know. Pure software SaaS: 70 to 85%. Managed services SaaS: 50 to 70%. Marketplace with an operational component: 35 to 60%. Professional services: 20 to 45%. E-commerce: 30 to 55%. When a model shows gross margin significantly above the high end of the relevant benchmark, the first investor move is to check what is in COGS and what has been placed in OpEx. The most common COGS misclassifications are:

Customer success costs placed entirely in headcount rather than partially in COGS. For a SaaS business, customer success is at least partially a cost of delivering the service. Industry practice varies, but putting 100% of customer success in S&M or G&A to inflate gross margin is a classic error. Payment processing fees placed in OpEx. Stripe fees, payment gateway costs, and merchant processing charges are a cost of revenue, not an operating expense. Hosting and infrastructure placed in R&D. If your product runs on AWS, those hosting costs are part of delivering the product to customers. Classifying them in R&D can inflate reported gross margin. A gross margin that looks too good should be treated as a diligence prompt, not a compliment. The question to ask before the investor does is: if I moved every arguable cost out of OpEx and into COGS, what does gross margin look like? If the answer is materially lower than what the model shows, the classification needs to be revisited.

Red Flag 3: Assumptions with No Source

There is a meaningful difference between a model built from research and one built from guesses. Both can produce the same

number. Only one is defensible in a partner meeting. "CAC: GBP 400" is a placeholder.

"CAC: GBP 400 based on our last 90 days of paid acquisition on LinkedIn and Google, where we are seeing a 3.2% click-to-trial conversion rate and a 14% trial-to-paid conversion rate. We expect this to decrease to GBP 480 by month 18 as we exhaust the most efficient targeting segments" is an assumption. When an investor asks "walk me through your CAC assumption" and the founder says "I researched comparable companies and GBP 400 seemed reasonable," the follow-up question is immediate: which companies, what stage, what channel mix, what ACV? If the answer is vague, the CAC is a placeholder dressed up as research. The assumptions that investors probe most frequently are: CAC and the channel model that drives it, time-to-revenue for each product tier, gross margin and COGS components, churn rate and its basis, and headcount productivity assumptions. For each of these, the note in the assumptions tab should answer: what is the source, what is the range of comparable benchmarks, and where on that range are we assuming we will land?

Red Flag 4: No Downside Scenario

A model with only one scenario has made an assumption: that the plan will unfold as planned. Every investor who has funded more than three companies knows that plans rarely unfold as planned. The absence of a downside scenario communicates one of two things: either the founder has not modelled it, or they have and decided not to show it. Neither inspires confidence. A well-constructed conservative scenario is not a pessimistic version of the model. It is a realistic version that challenges the most uncertain assumptions. For most early-stage companies, the most uncertain assumptions are: time to first meaningful revenue, CAC in new channels, enterprise sales cycle length, and key hire timing. A conservative scenario that shows the business still viable, albeit with tighter runway, is actually a positive signal. It shows the founder has stress-tested the plan and the business model holds even when things do not go perfectly. A conservative scenario where the company runs out of cash in month nine creates a different

conversation, but it is better to surface that in the model than for the investor to discover it themselves.

Red Flag 5: Headcount Flat as Revenue Triples

If revenue grows from GBP 1M to GBP 3M ARR over 18 months but headcount grows from 12 to 14 people, the model is implying either that the additional GBP 2M of revenue requires almost no additional work, or that the headcount model has not been connected to the revenue model. The first situation is plausible in specific circumstances: a product with near-zero marginal cost of delivery and no manual customer success component. But it is rare and needs to be explicitly explained and justified. The second situation is a modeling error. A revenue model and a headcount model that are not connected are not a financial model. They are two separate spreadsheets. Every material revenue assumption should flow through to a headcount implication. If outbound sales is the growth driver, how many SDRs and AEs are needed to generate that pipeline? If customer success is needed to retain and expand that revenue, how many CSMs? If the product needs to scale, how many engineers? These connections are not optional. They are what makes the model coherent.

Red Flag 6: Using the P&L to Calculate Burn Rate

Burn rate is a cash concept. It should be calculated from the cash flow statement, not from the P&L. This distinction matters in two important situations. First, if the company has significant deferred revenue (annual subscriptions collected upfront), the P&L will show less revenue than cash collected. Using P&L net loss to calculate burn will overstate how much cash is being consumed. Second, if the company has significant accruals that are not yet paid in cash (a large vendor invoice received but not yet paid, for example), the P&L expense will run ahead of cash outflows. Using P&L net loss to calculate burn will understate actual cash burn. The correct definition: monthly cash burn is the net decrease in cash from operating activities as shown in the cash flow statement. Runway is current cash divided by monthly

burn. Founders who calculate runway from the P&L and end up with a different answer than investors are usually making this mistake. The investor's number is right.

Red Flag 7: Formula Errors

A model sent to investors with REF errors, VALUE errors, or circular references has not been reviewed before it was sent. This is one of the fastest ways to lose credibility in diligence. The errors themselves may be minor. But the signal is that the model was not treated as a document representing the company. The pre-send review takes two hours. Open every tab. Check every formula. Look for cells with values that should be formulae and are not. Run a scenario test: change the revenue growth assumption by 10% and check that the change flows correctly through the P&L, the cash flow statement, and the balance sheet. If any tab does not update correctly, there is a broken link.

Red Flag 8: Founder Cannot Defend the Numbers

This is the hardest red flag to fix in the short term and the most damaging in the moment. If an investor asks "how did you get to this CAC number" and the founder needs to refer to the screen to remember, that is a signal. If the answer is "the model says GBP 400," that is a bigger signal. The model is supposed to represent the founder's thinking. It is built from assumptions the founder owns and can explain. A founder who has outsourced the model to a financial consultant and is presenting it without having built it themselves is in a difficult position when the investor starts asking second-level questions. Build the model yourself, or at minimum, build it with enough involvement that you can answer any question about any assumption without needing to look at the spreadsheet.

The Pre-Send Audit

Before any version of the model goes to investors, run through this checklist. Can every inflection in the revenue growth curve be explained by a specific business event named in the assumptions? Does gross

margin fall within the benchmark range, or is the gap explicitly justified? Does every material assumption have a note with a source? Is there a conservative scenario that shows the company's minimum viable outcome? Does headcount growth connect logically to revenue growth? Does burn rate come from the cash flow statement? Are there any formula errors on any tab? Can you answer 20 minutes of investor questions without looking at the screen? Eight checks. Two to three hours. The cost of not running them is measured in deal delays, valuation adjustments, and investor trust that is difficult to rebuild once lost.

Chapter Exercise

Send your model to a founder or advisor who will ask hard questions and request that they specifically probe the four most important assumptions: revenue ramp, gross margin, CAC, and headcount plan. For each assumption they challenge, write a two-sentence note in the assumptions tab explaining the basis and the range of comparable benchmarks. Repeat this until you can answer every material assumption question without referring to the screen.

039 How to Update Your Model Mid-Fundraise Without Losing Consistency

Why Models Change Mid-Fundraise

A fundraising process for a seed or Series A round typically runs 12 to 20 weeks from first meetings to close. Over that period, three types of changes happen to a financial model. The first is investor-driven. A VC challenges a specific assumption in a meeting and the founder agrees the challenge is valid. The CAC assumption was too optimistic given what the founder has learned from early customer acquisition. The gross margin does not fully account for customer success costs. The revenue ramp in year two assumed a sales hire who now will not be in place until month 14 rather than month 10. The second is data-driven. New monthly actuals come in. A month that was forecast in the model can now be replaced with real numbers. Sometimes the actuals are better than plan. Often they reveal gaps in the model's assumptions. Either way, the model should be updated. The third is strategic. The pricing model changes. A new enterprise tier is added. The go-to-market motion shifts from inbound to outbound. A key partnership materializes that changes the customer acquisition model. All three types of changes are legitimate. The problem is execution: how do you make these updates without inadvertently creating a model where different tabs reflect different versions of the truth?

The Architecture That Makes Consistent Updates Safe

The single most important structural choice in any financial model is centralized assumptions. Every material input lives in exactly one cell in the assumptions tab. Every formula in the model that uses that input references that cell. When the assumption changes, changing the single source cell flows the update through the entire model automatically. This sounds obvious. It is surprisingly rare in practice.

The most common alternative is a model where key assumptions are repeated in multiple tabs. CAC appears in the revenue model tab, the marketing spend tab, and the board summary tab. Each time someone updates the model, they might update one or two instances and miss the third. The model then contains three different values for the same assumption, none of which are labelled as the current version. If your model is built this way, the fix before sending to investors is to consolidate every material assumption into the assumptions tab, build a reference from every formula that uses it, and verify the update flows correctly by changing a value and checking all downstream effects. The categories of assumptions that most commonly need consolidating: Revenue assumptions: average contract value, conversion rates at each funnel stage, churn rate, expansion revenue percentage. These often appear separately in the revenue model, the LTV calculation, and the summary metrics tab. Cost assumptions: salary bands by role, employer overhead percentages, headcount productivity ramp timelines. These often appear in a headcount tab but are disconnected from the scenario model. Timing assumptions: first revenue month, first sales hire start date, product launch date. These are frequently hardcoded in different cells across multiple tabs rather than referenced from a single source.

Version Control: The Practice That Prevents Disaster

Every time a material update is made to the model, save a new version with a consistent naming convention. The format that works: CompanyName_Model_v[number]_[YYYYMMDD].xlsx Example: Simplifii_Model_v4_20250308.xlsx

This naming convention creates a complete audit trail of every version sent to every investor. When an investor references a number in a follow-up email three weeks after your last meeting, you can identify exactly which version they are working from and compare it to the current version. What not to do: overwrite the previous version without saving a copy first. Use naming conventions like "final," "final_v2," "FINAL_REAL," "latest." These are all jokes that become nightmares during diligence. What to do: create a folder structure that makes version management automatic. A folder called "Model Versions" with

every version saved chronologically. A separate folder called "Investor Versions" with the specific version sent to each investor, named by investor name and date. When sending an updated model, specify exactly what changed. Not "updated model attached." Something like: "Updated model attached, version 5 dated 8 March 2025. Changes from v4: CAC assumption updated from GBP 320 to GBP 380 reflecting February actuals. This moves base case 18-month runway from 14 months to 12 months. Conservative scenario runway moves from 10 months to 9 months. No other material changes." This communication practice does two things. First, it respects the investor's time by telling them exactly what to look at. Second, it creates a written record of every model update, which is valuable if there are any questions during diligence about when assumptions changed and why.

Handling Different Versions with Different Investors

One of the more delicate situations in a fundraise is managing investors who received different versions of the model. An investor who received version 2 in January may have built their own analysis on those numbers. If version 5 in March has materially different assumptions, that investor needs to know. The rule: any investor in active conversation who received a previous version should receive the updated version with a clear change summary. Do not assume they will ask. Do not wait until the next meeting. Send it proactively with an explanation. If an assumption change makes the business look materially worse (a revenue forecast cut, a runway compression, a gross margin reduction), the proactive update is particularly important. Investors who discover that numbers have changed without being told directly will interpret the lack of communication as concealment. The damage from that interpretation is far greater than the damage from the underlying assumption change. The one exception: if an investor has already passed on the deal, there is no obligation to send updates. But if they are still engaged, tracking the company, or have expressed interest in re-engaging at the next milestone, keep them on the update list.

The Three Most Destructive Mid-Raise Update Mistakes

Updating the model without updating the assumptions tab.

This happens more often than any other consistency error. The founder changes a number in the revenue model tab because an investor challenged the assumption. But the assumptions tab still shows the old value. An investor who opens the model and compares the assumptions tab to the revenue model will find the discrepancy immediately. The model looks inconsistent, and the assumption cannot be defended because it is not documented. **Sending an update without explaining what changed.**

An investor who received the previous version needs to know what is different. Sending version 5 without noting that it is version 5, without explaining what changed from version 4, forces the investor to compare documents line by line. Most will not do this. They will either use the old version or ask you to summarize the changes, which means an additional email exchange that could have been avoided. **Changing assumptions in one tab but not in dependent tabs.** A pricing change that flows through the revenue model but not through the LTV calculation. A headcount timing change that affects the cost model but not the productivity ramp assumptions in the revenue model. A churn rate change that updates the revenue forecast but not the cohort retention chart used in the investor summary. The fix is systematic: after every material update, open every tab and ask whether this change has any implication for the numbers here. Do not rely on memory. Do not assume the formulas are connected. Check each one.

Knowing When to Hold and When to Update

Not every investor challenge should result in a model update. Part of fundraising is defending your assumptions when you believe they are right. If an investor challenges your CAC assumption and you have 90 days of actual data supporting it, the right response is to defend the assumption, not to change it. The distinction is between a legitimate new input that changes your view, and an investor preference for more conservative assumptions. The former warrants updating the model.

The latter warrants a conversation about why your assumptions are defensible. If you do change an assumption in response to investor feedback, the change note should reflect honest reasoning: "updated based on actual data now available" or "updated based on a revised view of the enterprise sales cycle following conversations with prospects." It should not say "revised based on investor feedback" unless the feedback contained new information you did not previously have.

When an Update Should Trigger a Conversation, Not Just a File

Some model updates are administrative. You updated last month's actuals, nothing material changed, here is the current version. Send it with a brief note and move on. Other updates cross a threshold where the model update alone is insufficient. These are the moments when the change is significant enough that an investor deserves a call before they see the revised file. The situations that warrant a proactive call rather than just an updated model: A key assumption is being changed because the business performed materially differently than planned, not because of new data or new thinking. If revenue in the most recent month was 30% below plan, that is a conversation, not just a spreadsheet update. A scenario that was previously described as the base case is now looking like the optimistic case. If the conservative scenario has become the base case, investors need to hear that framing before they see the numbers. Runway falls below a threshold that changes what is possible. If the model update reveals that runway at current burn is now under nine months rather than fourteen months, that changes the context for every other conversation in the fundraise. The principle is: investors should never receive a materially worse model without also receiving a conversation that provides context. The context does not change the numbers. But it gives the founder the opportunity to explain what they have learned, what they are doing about it, and why the thesis is still intact. Receiving a worse model without context leaves the investor to draw their own conclusions. Those conclusions are rarely as optimistic as the founder's.

FAQ

How often should the model be updated during an active fundraise? Material changes should be reflected as soon as possible. Minor updates (new monthly actuals with no material variance from plan) can be batched into a monthly update. Do not send multiple versions per week unless there is a genuinely significant change each time. Frequent small updates create noise and suggest the model is not stable. **What if different investors are at different stages and have received different versions?** Track it. A simple spreadsheet with investor name, version received, and date sent is all you need. When a material update is issued, every investor in active conversation gets the new version simultaneously with the same change summary. No investor should feel they received information later than another. **Should the model include a "last updated" cell that is visible on the summary tab?** Yes. Put the version number and last updated date on the summary tab. This removes any ambiguity about which version an investor is looking at and signals that the model is actively maintained.

Chapter Exercise

Take your current model and identify every place where a key assumption appears in more than one tab. Consolidate each of those assumptions into the assumptions tab and replace the other instances with cell references. Then save the current version with a proper version naming convention. Create a folder for model versions and move all previous versions into it. Finally, write a one-paragraph change summary describing the most recent material update to the model, as if you were sending it to an investor. This exercise will surface whether your model architecture is set up for safe iteration.

040 From Model to Board Report: Keeping the Spreadsheet Alive Post-Raise

The Shift That Has to Happen

The fundraising model and the operating model serve different masters and different timelines.

This transition should happen within 30 days of round close. The capital is in the bank. The plan that was funded now becomes the baseline against which all future performance is measured. Every month that passes without actuals being tracked against that plan is a month where variance accumulates without visibility.

Building the Actuals Layer

The first operational update to the fundraising model is adding the actuals layer. For every period that has passed since the model was built, create four columns alongside the existing forecast. **Plan:** the original projection from the fundraising model. Do not change this retrospectively. The plan is the plan. Its value comes from being a fixed reference point. **Actual:** what actually happened, pulled directly from the accounting system. For revenue, this is recognised revenue per the accounts, not cash collected. For costs, this is the expense per the profit and loss account. **Variance (absolute):** actual minus plan. Positive variance on revenue is good. Positive variance on costs means overspend. **Variance (percentage):** absolute variance divided by plan. Variances above 10% in either direction should be explained. The actuals come from the accounting system. If the company is using Xero, QuickBooks, or another cloud accounting tool, the P&L and cash flow data exports to Excel in a format that can be mapped to the model categories. Setting up this mapping correctly once makes every subsequent monthly import a 20-minute exercise rather than a

three-hour one. The key prerequisite is a chart of accounts in the accounting system that maps directly to the categories in the financial model. If the model has a "Customer Success" line but the accounting system has all payroll in a single "Salaries" account, the mapping is impossible without a manual allocation. Align the chart of accounts to the model structure before the accounting system has too many historical transactions to reclassify.

The Monthly Update Cadence

The operational model runs on a monthly cycle. The target is to have the prior month fully updated and a revised forward forecast ready within ten business days of month end. **Week 1 of the new month:** Close the prior month's books. The finance lead or CFO exports the P&L, balance sheet, and cash flow statement from the accounting system. Actuals are imported into the model. **Week 1 to 2:** Variances are calculated. For every line item more than 10% off plan, a variance note is written. The note should have two components: what caused the variance, and what the expected trajectory is. "Revenue GBP 180k vs. GBP 220k plan: two enterprise deals slipped to April due to extended procurement timelines; pipeline coverage for April is strong" is a variance note. "Revenue below plan" is not. **Week 2:** The forward forecast is updated if material assumptions have changed. If Q2 is looking weaker than planned based on pipeline data, update the forecast. Do not maintain an obviously outdated projection just to avoid a difficult conversation at the board meeting. **Week 3 (pre-board):** The board report is compiled. This is not the full model. It is the output of the model: headline metrics versus plan, P&L summary, cash and runway update, variance analysis, and updated forward forecast. Reviewed with the finance lead or CEO before distribution to the board.

The Board Report Structure That Works

Board members receive a lot of information. A board report that runs to 40 pages will not be read in full. A board report that runs to 8 to

12 pages on the right topics will be read completely and will drive a productive meeting. **Section 1: Headline Metrics vs. Plan (1 page)**

The most important KPIs shown as actuals versus plan. For a SaaS business this typically includes: MRR or ARR, net new MRR, churn rate, CAC, LTV, gross margin, and cash runway. Present each as a number, a plan number, and a variance indicator. Green for at or above plan. Amber for 5 to 10% below plan. Red for more than 10% below plan. No narrative in this section. Just the numbers. This section should be the first thing the board sees. Starting a board meeting with 20 minutes of strategic discussion before revealing that revenue is 25% below plan is a structurally poor approach to governance. Put the numbers first. **Section 2: P&L Summary (1 to 2 pages)**

Revenue, COGS, gross profit, each major OpEx category, EBITDA, net income. Three columns: current period actuals, current period plan, prior period actuals. The comparison to prior period is as useful as the comparison to plan because it shows trajectory. A company that was 15% below plan last month and is 8% below plan this month is improving, even if it is still behind. A company that was 5% below plan and is now 20% below plan is deteriorating. **Section 3: Cash and Runway Update (half page)**

Current cash balance. Current monthly burn (from the cash flow statement, not the P&L). Forward runway at current burn. Runway at revised forecast burn. Any material change from the last board meeting and its explanation. This section should never be a surprise. Board members who discover unexpected cash deterioration at the board meeting, rather than in a proactive update between meetings, lose confidence in management's ability to manage the business. Material runway changes should be communicated between board meetings, not at them. **Section 4: Variance Analysis (1 to 2 pages)**

For every variance more than 10% off plan, one to two sentences: the cause and the expected resolution. Organized by P&L category. Not a narrative essay. Specific, factual, forward-looking. The most important discipline in variance analysis is honesty about whether a variance is temporary or structural. A one-month slip in enterprise deals is likely temporary. A persistent 15% CAC overrun across four months is

structural and needs to change the forward forecast. **Section 5: Forward Forecast (1 page)**

The next three to six months, updated for current actuals and near-term plan changes. This is where the operational model's forward-looking capability demonstrates its value. A board that has only historical actuals is looking backward. A board that has actuals plus a current forecast can have a useful forward-looking conversation.

What Happens When the Model Goes Stale

Founders who close a round and then stop maintaining the model experience a predictable pattern. Board meetings become narrative-only. Without current numbers, there is no analytical foundation for the discussion. Variances are described verbally and approximately rather than quantified precisely. The board loses confidence in management's ability to forecast. Runway surprises emerge. Cash is a lagging indicator. By the time the bank balance is worryingly low, the model should have been signalling the problem for three to four months. A stale model means that signal is absent. The next fundraise is harder. At the next round, investors will ask to see actuals versus the plan from the previous raise. A founder who can present a clean track record of actuals against forecast, including honest variance analysis, is in a materially stronger position than one who cannot. The ability to say "here is what we said we would do, here is what we actually did, and here is what we learned" is one of the strongest credibility signals in a Series A process. Investor relationships deteriorate. Board members who funded the raise based on a financial model expect to see progress tracked against it. A company that reports only qualitative updates signals that the quantitative discipline is not there.

Building the Finance Function Around the Model

Post-raise, most companies are adding their first finance hire or contracting with a fractional CFO or finance partner. The operational model should be the core output that this function maintains. The

finance function's responsibilities relative to the model are: importing actuals monthly, calculating variances, writing variance notes in collaboration with the CEO and functional leads, updating the forward forecast, and compiling the board report. This takes approximately 8 to 12 hours per month for an early-stage company with a straightforward business model. It takes 15 to 20 hours for a company with multiple revenue lines, international operations, or complex cost structures. If you are not yet at the point where you have a finance hire or a fractional CFO, the model maintenance is the CEO's responsibility. The eight to twelve hours per month is not optional. It is the cost of running the business in a way that surfaces problems before they become crises and demonstrates to investors that their capital is being managed with discipline.

Chapter Exercise

Take the financial model you used for your most recent fundraise and add the actuals layer for the last three months. For each month, create the four-column structure: plan, actual, variance absolute, variance percentage. Write a one to two sentence variance note for every line more than 10% off plan. Then build the five-section board report template using the structure described in this article, populated with the most recent month's data. Share the board report with your existing investors or board members and ask for feedback on what information they found most and least useful.

PART

Part Three: Unit Economics

Unit economics are the foundation of every credible financial model. This part covers the metrics that matter most at each stage, how to calculate them correctly, and how to benchmark them against comparable businesses. LTV, CAC, NRR, gross margin, and payback period: each has a right way to be calculated and several wrong ways that are surprisingly common.

041 CAC, LTV, and Payback Period: Calculated Correctly for Once

What Each Metric Is Actually Measuring

CAC (Customer Acquisition Cost): The total cost of acquiring one new paying customer. This includes all sales and marketing spend that contributed to that acquisition, divided by the number of new customers acquired in the same period.

LTV (Lifetime Value): The total gross profit contribution expected from a customer over the entire duration of their relationship with the business. Not revenue. Gross profit.

Payback period: The number of months required for a customer to generate enough gross profit to recover the CAC invested to acquire them.

The relationship between them:

Metric	Formula	Common Error
CAC	Total S&M spend ÷ New customers acquired	Using ad spend only, not fully loaded S&M
LTV	Avg monthly gross profit per customer ÷ Monthly churn rate	Using revenue instead of gross profit
LTV:CAC	LTV ÷ CAC	Both inputs wrong = misleading ratio
Payback period	CAC ÷ Avg monthly gross profit per customer	Using revenue instead of gross profit

How to Calculate CAC Correctly

The most common CAC error is including only paid advertising spend. The correct denominator includes every cost that contributed to acquiring that customer:

All paid media (search, social, display)

Sales team salaries, commissions, and on-target earnings

Sales tools and CRM costs (allocated to acquisition)

Marketing team salaries (for marketing-led growth models)

Events, conferences, and trade show costs

Content and SEO investment (if it drives direct acquisition) Agency fees

The period of calculation matters too. A company that hires a sales team in month one and closes the first customers in month three should not calculate CAC only in month three (which would look artificially low)

cost.

For early-stage businesses with small numbers, use a rolling 3-6 month CAC to smooth out timing mismatches between spend and customers acquired.

CAC formula:

CAC = Total Sales & Marketing Spend (period) ÷ New Customers Acquired (period)

How to Calculate LTV Correctly

The two most common LTV errors: using revenue instead of gross profit, and using an overly optimistic churn rate.

The gross profit distinction:

LTV is the value the business retains from each customer, not the revenue the customer pays. A customer paying £1,000 per month in a business with 70% gross margins contributes £700 per month in gross profit. Using revenue overstates LTV by the inverse of the gross margin. At 70% gross margin, revenue-based LTV is 43% higher than the correct gross-profit-based LTV.

The churn rate input:

For a subscription business, the simplest LTV model is:

LTV = Average Monthly Gross Profit per Customer ÷ Monthly Churn Rate This formula assumes constant gross profit per customer and constant churn, which is an approximation. It is a starting point, not a precise model.

For a business with significant churn seasonality, cohort-specific churn rates, or a known expansion revenue component, build a more detailed LTV model using cohort analysis rather than the simple formula. **The churn error:**

Using an annual churn rate in a monthly formula (or vice versa) without converting. Monthly churn of 2% is not the same as annual churn of 24%

annual churn.

How to Calculate Payback Period Correctly

Payback period is the number of months to recover CAC through gross profit contribution.

Payback Period (months) = CAC ÷ Average Monthly Gross Profit per Customer

Using revenue instead of gross profit understates the payback period — which is the same as overstating capital efficiency. This error is common in pitches because the shorter payback number looks better. **Payback period benchmarks by business model:**

Business Model	Strong	Acceptable	Concerning
SaaS	< 12 months	12–18 months 18 months Marketplace	>18months

The LTV:CAC Ratio: What It Actually Measures

The LTV:CAC ratio measures the return on customer acquisition investment. A ratio of 3:1 means every pound invested in acquiring a customer returns three pounds in lifetime gross profit. Industry convention treats 3:1 as a minimum threshold for a sustainable SaaS business.

But the ratio is only as meaningful as its inputs. LTV:CAC calculated with revenue-based LTV and ad-spend-only CAC can look like 5:1 when the correct calculation is 2:1. The ratio is a useful benchmark when both components are calculated correctly.

What the ratio does not tell you:

How long it takes to realise the LTV (a 3:1 ratio over 3 years is different from 3:1 over 8 years)

Whether the churn assumption is realistic

Whether CAC is trending in the right direction at scale

The payback period addresses the first problem. Cohort analysis addresses the second. CAC by channel and acquisition cohort addresses the third.

> *Key insight: The LTV:CAC ratio and payback period are companion metrics, not substitutes. A business with a great LTV:CAC ratio but a 36-month payback is capital-intensive — it requires a lot of cash before the acquisition investment is recovered. A business with a 3:1 ratio and 8-month payback is capital-efficient. Both metrics together tell the full story.*

CAC must include all sales and marketing spend, not just ad spend. LTV must be calculated on gross profit, not revenue, with a realistic churn rate. Payback period must use gross profit per customer. The LTV:CAC ratio is only meaningful when both inputs are correct. The payback period and LTV:CAC ratio together tell the full capital efficiency story: the ratio shows the return, the payback period shows how long it takes to realise it. Get these three metrics right before any investor conversation that references unit economics.

042 Why Blended CAC Is a Lie (And How to Model It by Channel)

What Blended CAC Hides

A company with £200k monthly S&M spend and 400 new customers has a blended CAC of £500. That number tells a clean story for the deck. Now split the same month by channel:

Organic/SEO: 200 customers, £0 direct acquisition cost (CAC ~£0 marginal)

Paid search: 100 customers, £100k spend (CAC £1,000)

Outbound sales: 100 customers, £100k in sales salaries and tooling (CAC £1,000)

Blended CAC: £500. Channel-level CAC: £1,000 for every customer the business is actually paying to acquire.

If the business doubles paid acquisition spend to scale, the blended CAC increases toward £1,000 as organic's dilutive effect shrinks as a proportion of total customers. This dynamic — where organic performance artificially suppresses blended CAC and makes the business look more capital-efficient than it is — is one of the most common unit economics traps in growth stage companies.

The Three Channel Categories That Matter

Organic / earned channels (Brand, SEO, referral, word of mouth): These channels produce customers with low or zero marginal acquisition cost. They are valuable and should be grown. But they should not be used to hide the true cost of paid channels in a blended calculation. For modelling: organic channels have a cost — the content team, SEO investment, brand spend, referral programme — but it is often partially shared with brand-building objectives and lower than the cost of the customers it produces. Calculate a true organic CAC that includes direct investment in these channels.

Paid performance channels (Paid search, paid social, display): These channels have a direct, measurable relationship between spend and customers acquired. The marginal CAC for paid channels is clear and controllable within limits. Beyond those limits (audience saturation, keyword competition), paid CAC increases rapidly.

For modelling: paid CAC should be calculated separately from organic and tracked over time. Rising paid CAC is an early warning signal about channel efficiency that blended CAC masks.

Sales-led channels (Outbound, SDR/AE, events, partnerships): These channels have a people cost at their core: the fully loaded cost of the sales team plus all associated tooling, events, and support. Sales CAC is often higher than paid digital but produces higher ACV customers with different LTV profiles.

For modelling: sales-led CAC should be calculated as the fully loaded cost of the sales function divided by customers closed through that function. Separating this from marketing-led CAC reveals whether the sales team is efficient at the current quota and pipeline coverage.

How to Build a Channel-Level CAC Model

Step 1: Attribute every customer to an acquisition source. This requires a consistent attribution model in CRM and analytics. First-touch, last-touch, and multi-touch attribution each produce different channel CAC numbers. Pick one model and apply it consistently. **Step 2: Attribute every S&M cost to a channel.**

Some costs are easy to attribute: paid search spend goes to the paid search channel. Sales team salaries go to the sales channel. Some costs are shared: the marketing team's time is split across multiple channels. Build a simple allocation that assigns shared costs to channels based on estimated time or effort.

Step 3: Calculate channel CAC for each acquisition source. Channel CAC = Channel-Attributed Costs (period) ÷ Channel-Attributed New Customers (period)

Step 4: Track channel CAC over time.

The trend matters more than the absolute level. A paid search CAC that is increasing by 15% month-over-month is a signal worth investigating. A sales CAC that is declining as the team matures and improves conversion rates is a positive signal.

Step 5: Build a channel mix model.

Model the forward acquisition plan as a mix of channels, each with its own CAC and capacity constraint. This produces a blended CAC forecast that shows how blended CAC changes as the channel mix shifts . specifically, what happens to blended CAC as paid acquisition scales faster than organic.

What Investors Are Actually Asking When They Ask About CAC

When a VC asks "what is your CAC?" they are often really asking several questions at once:

Is organic traffic subsidising the paid acquisition cost?

What happens to CAC if you try to 2x or 5x customer acquisition? Is the sales team efficient, or is it taking 18 months to close deals that should close in 6?

Which channels are scalable beyond current volume?

Blended CAC answers none of these. Channel-level CAC with a channel mix model answers all of them.

> *Key insight: The founder who can say "our blended CAC is £500 but our paid digital CAC is £900 and our outbound CAC is £1,200, and organic currently accounts for 40% of new customers" is giving investors the information they need to assess scalability. The founder who only knows the blended number is giving investors a reason to dig.*

Common Mistakes in Channel-Level CAC Modelling

Allocating the entire S&M budget to paid acquisition. If marketing does content, brand, and events in addition to running paid campaigns, attributing 100% of the marketing budget to the paid channel overstates paid CAC and understates the true organic channel economics. **Ignoring the cost of the free tier in PLG models.** In product-led growth businesses, the free tier infrastructure, support, and product investment is part of the cost of acquiring paid customers. These costs belong in CAC, not in COGS.

Not tracking attribution model changes. Switching from last-touch to first-touch attribution changes every channel's CAC without any change in underlying performance. When attribution models change, restate the historical data under the new model before drawing trend conclusions.

Blended CAC is a headline metric that obscures which channels are profitable, whether organic performance is subsidising paid channels, and what happens to unit economics at scale. Building channel-level CAC requires attributing both customers and costs by source, tracking the trend over time, and building a channel mix model that shows how blended CAC evolves as acquisition scales. This is the analysis investors are implicitly requesting when they ask about CAC, and it is the analysis that distinguishes founders who understand their growth economics from those who know the headline but not the substance.

043 Net Revenue Retention: The Number That Predicts Everything

What NRR Is and How to Calculate It

NRR measures what happens to a cohort of customers' revenue over a period, typically 12 months.

Formula:

NRR = (Starting Revenue + Expansion Revenue - Churned Revenue - Contracted Revenue) ÷ Starting Revenue × 100

Where:

Starting revenue: MRR or ARR from the cohort at the start of the period

Expansion revenue: Additional revenue from the same customers (upsells, seat additions, usage growth)

Churned revenue: Revenue lost from customers who cancelled
Contracted revenue: Revenue lost from customers who downgraded but did not cancel

Why NRR Above 100% Is a Fundamental Business Quality Signal

A business with NRR above 100% has an inherent compounding mechanism: existing customers generate more revenue over time. New customer acquisition adds to a growing base, not a leaky one.

A business with NRR below 100% has a leaky bucket. New customers partially offset the revenue lost from existing ones. Growth requires continuous high-volume acquisition just to maintain the revenue run rate.

These are structurally different businesses. NRR is the cleanest single metric that distinguishes them.

The compounding effect is significant over time:

A business with 110% NRR doubles its existing revenue base in approximately 7 years from existing customers alone

A business with 90% NRR loses 65% of its existing revenue base over the same period if no new customers are acquired

This is why Series A and growth stage investors focus heavily on NRR. It is not just a retention metric. It is a prediction of the business's capital efficiency at scale.

How to Model NRR in a Financial Model

Step 1: Build a cohort retention schedule.

Group customers by acquisition cohort (month or quarter). For each cohort, model the retention rate and expansion rate over time. Use actual data where available; use benchmarks or assumptions where not. **Step 2: Separate gross retention from net retention.**

Gross Revenue Retention (GRR): What percentage of starting revenue is retained (ignoring expansion). GRR can only be 100% or below.

Net Revenue Retention (NRR): GRR plus expansion revenue. NRR can exceed 100%.

Both metrics matter. High NRR driven by massive expansion that masks poor GRR is a different risk profile than high NRR driven by both strong retention and moderate expansion.

Step 3: Apply realistic churn curves.

Early-month churn is typically higher than steady-state churn. Customers who survive the first 3-6 months tend to churn at lower rates thereafter. A flat monthly churn rate assumption overstates churn for mature cohorts and understates the early-stage churn risk. **Step 4: Model expansion explicitly.**

Expansion revenue from upsells, seat additions, or usage growth should be modelled as a separate line with its own assumption, not blended into new customer MRR. This keeps the NRR calculation clean and makes it easier to track expansion as a separate driver.

> *Key insight: The difference between a model that shows NRR and one that does not is the difference between a model that can be stress-tested on retention and one that cannot. Build NRR into the model from the start, not as a retrospective calculation.*

The NRR Questions Investors Ask

"Walk me through your NRR calculation."

The answer should demonstrate understanding of the components: what revenue is included in the starting base, how expansion is defined and tracked, and what drives churn. Vague answers signal that NRR has been calculated but not understood.

"What is your NRR excluding your top 3 customers?"

Concentration risk. If NRR is 115% but three customers account for 60% of the expansion, the underlying NRR is significantly lower. Know this number.

"How has NRR trended over the last 12 months?"

Investors want to see NRR stable or improving. Declining NRR with a "we've fixed the retention problem" explanation is a red flag if the data does not support it yet.

"What would NRR look like if the expansion rate drops to zero?" This isolates GRR from expansion. A business with 115% NRR but 85% GRR is heavily dependent on expansion to offset churn. The GRR tells the retention story; expansion tells the growth story.

NRR is the metric that reveals whether a subscription business is building on solid foundations or constantly replacing what it loses. NRR above 100% creates a compounding revenue base where new customers are additive. NRR below 100% creates a replacement treadmill. Model it at the cohort level, separate GRR from NRR, and track expansion revenue as a distinct driver. Know the concentration-adjusted NRR as well as the headline number. The investors who focus on NRR are not being pedantic

whether the business grows efficiently at scale.

044 How to Model LTV for a Marketplace Business

The Two-Sided LTV Problem

A SaaS business has one customer type. A marketplace has two: the demand side (buyers, employers, clients) and the supply side (workers, vendors, service providers). Both sides generate value, both sides can churn, and the economics of each differ.

The first modelling decision is which side to calculate LTV for. The answer depends on which side the marketplace monetises.

Monetisation-side LTV: Calculate LTV for whichever side pays the marketplace (demand side in most marketplace models). This is the LTV that maps directly to CAC, because CAC is typically the cost of acquiring the paying side.

Supply-side value: Even if the supply side does not pay, supply-side retention drives demand-side LTV. A marketplace where workers or vendors churn frequently forces the demand side to re-onboard new supply, reducing quality and driving demand-side churn. Model supply retention as a driver of demand-side LTV even if supply-side LTV is not calculated directly.

LTV Components for a Marketplace

For a transaction-fee marketplace (e.g. staffing, freelance, logistics):

LTV = Average Transaction Value × Take Rate × Gross Margin on Take Rate ×

Average Monthly Transactions × (1 ÷ Monthly Demand-Side Churn Rate) This is equivalent to:

LTV = Average Monthly Gross Profit per Customer ÷ Monthly Churn Rate \...where "gross profit" is the marketplace's net take after payment processing and direct fulfilment costs, not the full transaction GMV. For

a subscription marketplace (e.g. platforms where the buyer pays a subscription):

The LTV model is closer to SaaS: monthly subscription revenue × gross margin ÷ churn rate. But transaction data should be layered in if the subscription price is linked to usage volume.

The GMV vs. Net Revenue Distinction

The most common LTV error for marketplaces: using GMV (gross merchandise value, the full transaction value passing through the platform) as the revenue basis for LTV rather than net revenue (the take rate portion the marketplace actually keeps).

A staffing marketplace that places a worker at £15/hour and charges the employer £18/hour has:

GMV: £18/hour × hours worked

Net revenue (take): £3/hour × hours worked

Gross profit on the take: £3/hour × hours worked × gross margin on fulfilment

LTV should be calculated on gross profit from net revenue, not GMV. Using GMV as the LTV basis overstates LTV by the inverse of the take rate — at a typical 15-20% take rate, this is a 5-7x overstatement.

> *Key insight: When a marketplace founder says their LTV is £X, the first clarifying question is always: is that gross profit on net revenue, or is it GMV? The answer changes the number dramatically and determines whether the LTV:CAC ratio reflects a viable business or an accounting artefact.*

Modelling Transaction Frequency Over Time

One of the key differentiators in marketplace LTV is transaction frequency: how often does a demand-side customer transact, and does frequency grow or shrink over time?

Cohort analysis of transaction frequency is the most valuable input to marketplace LTV modelling:

In the first 1-3 months after acquisition, transaction frequency is often low as the customer tests the platform

In months 3-12, frequency typically increases as trust and workflow integration build

In months 12+, frequency stabilises at a steady state or continues to grow for deeply embedded customers

Building LTV as a constant (average monthly transactions × take × margin) misses this ramp-up dynamic. A more accurate model uses cohort-observed frequency curves to build a lifetime transaction schedule.

For early-stage marketplaces without 12+ months of cohort data, use a conservative frequency assumption (no ramp-up, steady-state from month one) as the base case and model the ramp-up as upside.

Supply-Side Retention as a LTV Driver

Supply-side churn creates demand-side LTV risk that does not appear in a buyer-only LTV model. If the supply side churns and quality or availability on the platform deteriorates, demand-side churn follows. For marketplace models, include a supply-side retention assumption and model the feedback loop:

The exact relationship varies by marketplace type, but acknowledging this linkage in the model notes is important for any marketplace with operational supply-demand matching.

Worked Example: Staffing Marketplace LTV

Assumptions:

Average employer (demand-side customer) places 3 workers per month Each worker placed generates £250 in net revenue (take) to the marketplace

Gross margin on fulfilment: 55%

Monthly gross profit per customer: 3 × £250 × 55% = £412.50 Monthly demand-side churn rate: 3%

LTV = £412.50 ÷ 0.03 = £13,750

If this were calculated on GMV instead (employer rate of, say, £1,500 per worker per month):

Incorrect LTV = (3 × £1,500 × 55%) ÷ 0.03 = £82,500

The GMV-based LTV is 6x the correct number. This is not an edge case — it is the most common marketplace LTV error.

Marketplace LTV is calculated on gross profit from net revenue (the take rate portion), not on GMV. Transaction frequency evolves over time and should be modelled using cohort data where available. Supply-side retention drives demand-side LTV and should be modelled as a linked assumption even if supply-side LTV is not calculated directly. The most common error — using GMV as the LTV basis — overstates LTV by the inverse of the take rate, which at typical marketplace take rates is a 5-7x error. Get the LTV denominator right before any investor conversation about unit economics.

045 Unit Economics Across Multiple Markets: UK vs US vs UAE

Why Unit Economics Differ by Market

The same product in the UK, US, and UAE faces different competitive landscapes, different customer behaviour, different pricing norms, and different employment cost structures. Each of these affects unit economics independently.

The drivers of CAC variation across markets:

Driver	UK	US	UAE
Paid media competition	Moderate–high	High	Low–moderate
Sales cycle length	Medium	Varies by segment	Short
Brand recognition leverage	Depends on category	Larger market, high if anchored	High if anchored
Average contract value	Moderate	Highest in most SaaS categories	Moderate–high
Churn patterns	Similar to US	Benchmark churn rates typically set here	Often lower in enterprise
Expansion potential	Moderate	Highest expansion rates	Lower expansion potential
Currency	GBP-denominated	USD-denominated	AED (pegged to USD)

The Market-Level Unit Economics Framework

Building market-level unit economics requires separating two things that are often combined in a blended model: the cost of acquiring customers in each market and the revenue and retention profile of those customers. **Step 1: Separate S&M costs by market.**

If the company has dedicated sales or marketing teams per market, the allocation is straightforward. If costs are shared (a central marketing team running campaigns across markets), build an allocation based on time spent or budget deployed in each market.

Step 2: Attribute customers by market.

This should come from CRM data. Every new customer should have a market tag that enables market-level acquisition volume tracking. **Step 3: Calculate market-level CAC.**

Market CAC = Market S&M Costs (period) ÷ New Customers Acquired in Market (period)

Step 4: Build market-level LTV.

Use market-specific churn rates, average contract values, and expansion assumptions. Markets with higher ACV often have lower churn and different expansion curves — do not apply the global average to every market.

Step 5: Calculate market-level payback period and LTV:CAC ratio. Compare these across markets. The insight is almost always that markets are not equally efficient, and that capital should be allocated toward the most efficient markets at early scale, not spread evenly.

What Typically Differs Between UK, US, and UAE

UK vs. US:

For many B2B products, the US has both higher CAC (more competitive market, longer enterprise sales cycles in some segments) and higher LTV (larger deal sizes, more expansion opportunity, stronger NRR benchmarks for SaaS). Whether the US is more or less capital-efficient depends on whether the higher LTV more than compensates for the higher CAC. This is not always true — many

UK-first startups discover that the US requires 2-3x the CAC per customer with only 1.5x the ACV, making the US less efficient in early expansion.

UK vs. UAE:

UAE enterprise markets often have shorter sales cycles and higher relationship-dependency. A founder with credibility and existing network in the UAE can achieve much lower CAC through direct outreach than through paid digital. The trade-off is scalability: relationship-driven acquisition does not scale the same way a repeatable paid or inbound motion does. UAE churn patterns in early enterprise relationships are often lower than UK SMB churn, but the customer base is smaller and more concentrated, creating different LTV risk profiles.

Currency considerations:

UK customers generate GBP-denominated revenue. US customers generate USD. UAE customers generate AED (pegged to USD). For a USD-functional company, UK LTV is subject to GBP/USD FX risk. Model this explicitly rather than converting at spot rate and assuming stability.

The Market Expansion Decision Framework

Unit economics at market level answer the core expansion question: which market should receive the next increment of capital, and why? The framework:

1. Calculate market-level payback period for each current market 2. Identify the market with the shortest payback and highest LTV:CAC ratio
3. Assess whether additional capital in that market can acquire more customers at similar unit economics (capacity test) or whether the market is approaching saturation
4. If capacity exists: concentrate capital in the most efficient market before expanding into less efficient ones

5. If near saturation: evaluate whether the next market has comparable or better unit economics at the scale being considered

This is a more capital-efficient expansion strategy than geographic diversification for its own sake. Growing a profitable UK market to maturity before scaling a less efficient US operation is often the right call, even if the US looks more attractive on an absolute basis.

> *Key insight: The startups that have the best Series A conversations about market expansion are the ones that can show market-level unit economics and explain exactly why they are choosing to expand in a specific direction and with a specific capital allocation. "We're expanding to the US because it's a bigger market" is a narrative. "UK payback is 8 months, US pilot shows 14 months at current CAC with 1.4x ACV premium — we need to prove the US CAC can come down through channel optimisation before we scale capital there" is an analysis.*

Unit economics differ systematically across markets due to differences in competitive intensity, pricing norms, deal sizes, and churn behaviour. Building market-level unit economics requires separating S&M costs by market, attributing customers by market, and calculating CAC, LTV, and payback period for each. The insight from this analysis — which markets are capital-efficient and which are subsidised — is the foundation for evidence-based expansion decisions. Blended global unit economics flatten these differences and produce a metric that is neither accurate for any single market nor useful for capital allocation.

046 The LTV:CAC Benchmarks by Stage and Business Model

Why There Is No Universal LTV:CAC Benchmark

The 3:1 rule — often stated as "SaaS businesses need at least a 3:1 LTV:CAC ratio to be viable" — is useful as a starting point and misleading as a universal standard.

Three reasons why:

Business model differences. Marketplace businesses with thin take rates need lower payback periods to compensate for lower absolute gross profit per transaction. E-commerce businesses can operate profitably at lower LTV:CAC ratios because payback periods are shorter. Professional services businesses often show lower ratios because LTV calculations are harder to apply to project-based relationships.

Stage differences. At pre-seed and seed, LTV:CAC is a directional signal built on limited data. At Series A, investors expect demonstrably calculated ratios with enough cohort history to make them meaningful. At growth stage, the ratio is benchmarked against public company comparables.

Input calculation differences. A ratio calculated with revenue-based LTV and ad-spend-only CAC is not the same metric as one calculated with gross-profit-based LTV and fully loaded CAC. Comparing across companies without standardising inputs produces comparisons that mean nothing.

LTV:CAC Benchmarks by Business Model

Business Model	Minimum Viable	Strong	Exceptional
SaaS (subscription)	3:1	4–5:1	*6:1 Market place (take rate)*

strong NRR and significant expansion revenue when the product is deeply embedded. The lower threshold for e-commerce reflects shorter payback periods and higher transaction frequency.

LTV:CAC Expectations by Funding Stage

Pre-seed:

LTV:CAC at pre-seed is almost always directional. There are rarely enough customers or enough cohort history to calculate it accurately. What investors want to see at pre-seed is a credible model with reasonable assumptions and an explanation of how LTV:CAC will be tracked and validated with the first cohorts. A directional calculation is acceptable; precision is not expected.

Seed:

At seed, investors expect demonstrated unit economics with early cohort data. The ratio should be calculated on actual customers, not solely on projections. If the cohort is small (fewer than 20 customers), the ratio is indicative and should be presented as such. Seed investors will look at the direction — is the ratio improving as more cohort data comes in? — as much as the absolute level.

Series A:

Series A is where LTV:CAC is scrutinised seriously. Investors at this stage expect the ratio to be calculated correctly, to be based on meaningful cohort data, and to be benchmarkable against comparable

companies at the same stage and business model. The typical Series A expectation for SaaS is 3:1 or better, with a trajectory toward 4:1 or above as scale increases.

Series B and growth stage:

At Series B and beyond, the ratio is compared to public company benchmarks and is used to assess whether the business has the unit economics to support the growth rate being targeted. Investors at this stage may also look at payback period as the primary capital efficiency metric, particularly if NRR is high and LTV:CAC is less informative than the compounding effect of the existing customer base.

What the LTV:CAC Ratio Cannot Tell You

Capital intensity: A 4:1 LTV:CAC ratio with a 36-month payback is far more capital-intensive than a 3:1 ratio with a 9-month payback. The ratio tells you the return; it does not tell you when the return is realised or how much capital is required before recovery.

Concentration risk: A 5:1 ratio driven by two large customers looks very different from a 5:1 ratio distributed across 200 customers. Know the concentration-adjusted ratio.

Trend: A 3:1 ratio that has declined from 5:1 over the last 12 months is a different story from a 3:1 ratio that has improved from 1.5:1. The direction matters as much as the current level. **Accuracy of inputs:** As discussed, the ratio is only as meaningful as its inputs. An apparently strong LTV:CAC may be an artefact of input errors.

> *Key insight: The LTV:CAC ratio is a starting point for an investor conversation, not the end of it. Investors who understand unit economics will ask: how is it calculated, what is the payback period, how has it trended, and what is the ratio excluding the top customers? Have answers to all four before the conversation starts.*

The LTV:CAC + Payback Period Dashboard

The most useful unit economics view combines LTV:CAC and payback period together, because each answers a different question:

A business that targets 3:1 LTV:CAC with an 8-month payback is in a fundamentally different capital position from one that targets the same ratio with a 30-month payback. Present both metrics together.

LTV:CAC benchmarks depend on the business model and funding stage, not on a universal threshold. Calculate both inputs correctly (gross-profit LTV, fully loaded CAC) before comparing to benchmarks. Know the payback period alongside the ratio — they answer different questions and together give a complete capital efficiency picture. Know the trend, the concentration adjustment, and the expected investor questions at the stage you are raising. The ratio is a starting point for the conversation; the surrounding context is what closes it.

047 Gross Margin by Business Model: What's Good, What's Bad, What's a Problem

Why Gross Margin Benchmarks Differ So Much

Gross margin reflects how much revenue remains after the direct cost of delivering the product or service. Because delivery costs differ fundamentally across business models, the benchmark ranges are wide. A pure SaaS business delivers software through infrastructure that costs a fraction of the revenue it generates. A marketplace with fulfilment operations delivers a service that requires significant people, logistics, or compliance costs with every transaction. A professional services firm delivers through human labour that is expensive relative to the revenue billed.

The benchmark is not about which model is better. It is about understanding what the unit economics look like for each model at scale, and whether the current gross margin is on a trajectory toward those benchmarks.

What Drives the Range Within Each Model

For SaaS:

The primary variable is the ratio of product-only revenue to professional services revenue. A SaaS business with 20% of revenue from implementation and customisation services will show lower gross margin than a pure product business, because services delivery is more labour-intensive. The trend matters too — is gross margin improving as the product matures and services become a smaller proportion of revenue? **For marketplaces:**

The primary variables are take rate and fulfilment intensity. A marketplace that takes 20% of transaction value on a fully automated platform has very different gross margin potential from one that takes

15% but employs an operations team to manage supply-side compliance and matching. The distinction between a "thin" platform and an "operated" marketplace is one of the most important structural choices for long-term gross margin.

For professional services:

The primary variables are utilisation rate and leverage ratio (junior staff per senior). Higher utilisation and higher leverage typically improve gross margin. The ceiling is lower than in product businesses because the cost of delivery scales directly with revenue. **For hardware + software:**

The blended gross margin depends on the relative weight of hardware and software revenue. Hardware gross margins are often 30-50%; software gross margins are 70-85%. As the software attach rate increases, blended margins improve significantly.

When High Gross Margin Is a Red Flag

Above-benchmark gross margin is not always a positive signal. It triggers the question: has COGS been correctly classified? The most common COGS misclassifications that inflate gross margin: Engineering time building new features classified as COGS instead of R&D

Customer success team entirely in Sales and Marketing rather than partially in COGS

Payment processing fees in OpEx rather than COGS

Infrastructure costs allocated to G&A rather than COGS

A SaaS business showing 90% gross margin when the benchmark is 70-80% will be asked these questions in diligence. If the answer is that the model is correctly classified and the business genuinely has very low infrastructure costs, that is worth explaining. If the answer is a COGS misclassification, fix it before the model is shared.

When Low Gross Margin Is a Red Flag

Below-benchmark gross margin triggers a different set of questions: is there a credible path to the benchmark, and what is the mechanism? The three legitimate explanations for below-benchmark gross margin: **Early-stage scale effects.** Fixed COGS are high relative to revenue at low scale. As revenue grows, the fixed cost component dilutes and gross margin improves. This is the most common explanation and is acceptable if the model shows the trajectory clearly.

Business model complexity. Some businesses operate at below-typical gross margins because their model involves more direct fulfilment than typical peers. A marketplace that takes more operational risk to guarantee quality may show lower gross margins than a pure-platform peer but generates more defensible supply relationships as a result. **Intentional investment.** A startup deliberately keeping delivery quality high (even at elevated cost) during the early customer acquisition phase, with a plan to reduce unit delivery costs at scale, should explain this explicitly.

What is not a legitimate explanation: "we expect gross margin to improve as we scale" without a model that shows the mechanism for that improvement.

The Gross Margin Trajectory Test

Investors reviewing a multi-year financial model will check whether gross margin improves over the forecast period and whether the mechanism is explained.

The test: if gross margin in year one is 45% and year five is 65%, the model should be able to explain:

Which COGS lines are fixed vs. variable?

What happens to fixed COGS per unit as revenue scales?

What negotiated rate improvements occur at volume?

What operational efficiencies reduce per-transaction costs? Each of these should be a named assumption in the model, not an implicit improvement baked into a declining cost percentage.

Gross margin benchmarks are model-specific and trajectory-dependent. Know the benchmark for your business model. If gross margin is above benchmark, be ready to explain why COGS classification is correct. If it is below benchmark, be ready to explain the mechanism for improvement and show it in the model. The gross margin question in investor diligence is not just "what is your margin" — it is "do you understand what drives it, why it is where it is, and where it is going?"

048 How to Build a Cohort Analysis That Actually Tells You Something

What Cohort Analysis Is and Why It Beats Aggregate Metrics

An aggregate retention metric — "we retain 85% of customers annually" — describes the current cross-section of the customer base. It does not tell you whether retention is improving or deteriorating, whether early cohorts behave differently from recent ones, or whether the 15% who churn do so in month one or month eighteen.

Cohort analysis tracks the same group of customers from their start date forward. This reveals:

When churn happens (is it front-loaded or distributed across the lifetime?)

Whether newer cohorts retain better or worse than older ones How revenue per cohort evolves (does expansion offset churn? At what point?)

What the actual LTV trajectory looks like from observed data rather than assumptions

The Structure of a Revenue Cohort Analysis

The standard revenue cohort table shows, for each acquisition cohort, the revenue generated in each subsequent period as a percentage of the revenue generated in the first period.

Example structure (simplified):

Cohort	Month 0	Month 1	Month 3	Month 6	Month 12
Jan 2024	100%	92%	88%	85%	90%
Feb 2024	100%	90%	86%	84%	89%
Mar 2024	100%	91%	87%	85%	—

What a Good vs. Problematic Cohort Chart Looks Like

Healthy patterns:

Revenue curve flattens after initial churn and stabilises, cohorts reach steady-state retention

Later months show values above 100%, expansion revenue drives NRR above 100%

Newer cohorts show similar or better retention curves than older ones, product-market fit is consistent

Warning patterns:

Revenue continues declining through month 12+ without stabilising ongoing churn with no floor

Newer cohorts show worse retention than older ones — potential product-market fit erosion or quality of acquisition declining

All growth is in month 0 with rapid decay — new customer acquisition driving the top line, not retention

Only a few cohorts are large enough to read — data concentration risk Key insight: Investors reviewing cohort data are specifically looking for the "smile curve" — a revenue cohort that dips in early months due to initial churn, then recovers and climbs above 100% as expansion revenue kicks in. Businesses with smile curves on their cohort data have materially stronger unit economics conversations than those without.

The Three Cohort Analyses Every Startup Should Have

1. Revenue retention cohort

Tracks revenue from each cohort over time, as a percentage of first-period revenue. This is the primary cohort analysis for any subscription or recurring revenue business. Shows NRR, GRR, and the expansion pattern all in one view.

2. Customer count cohort (logo retention)

Tracks the number of customers from each cohort that are still active in subsequent periods. This is different from revenue retention — a cohort might show improving revenue retention but declining logo retention if a few large customers are growing while many small ones churn. Both dimensions tell part of the story.

3. Usage or engagement cohort

For PLG businesses or products with usage-based pricing, cohort analysis on product usage (active sessions, feature adoption, API calls) is often the leading indicator of future revenue retention. Declining engagement cohorts predict revenue churn before it appears in the financials.

How to Build Cohort Analysis in Excel

The structure:

Rows: acquisition cohorts (each row is a group of customers who started in the same month)

Columns: periods since acquisition (month 0, month 1, month 2\...) Values: revenue (or customers) from that cohort in that period The calculation for each cell is: revenue from cohort X in period N ÷ revenue from cohort X in period 0.

Important: normalise to cohort start revenue, not absolute revenue. Absolute revenue comparisons across cohorts are misleading because later cohorts are typically larger. Percentage of starting revenue puts all cohorts on the same scale.

For startups with fewer than 12 months of data, the cohort table will have a triangular shape — early cohorts have more periods of data than recent ones. This is normal. Plot what you have and fill in the rest as data accumulates.

What Investors Ask When They See Cohort Data

"Show me the cohort where retention stabilises."

This asks for the month at which the decay curve flattens. Businesses where retention never stabilises have a churn problem that compounds over time. The stabilisation point also determines the appropriate LTV calculation window.

"Are your newer cohorts better or worse than your older ones?" Improving cohorts signal that retention is getting better as the product matures. Deteriorating cohorts signal that quality of acquisition may be declining (less targeted growth) or that the product is regressing. **"What drives the expansion in cohorts that show revenue above 100%?"**

Investors want to understand whether expansion is from organic usage growth, from proactive upselling, from seat additions, or from price increases. Each has different implications for sustainability and scalability.

Cohort analysis reveals what aggregate metrics hide: when churn happens, whether newer cohorts retain better than older ones, and how revenue evolves over the lifetime of a customer relationship. The revenue retention cohort table — showing each cohort's revenue as a percentage of its first-period revenue over time — is the primary tool. A healthy cohort chart shows a decay that stabilises and then recovers above 100% (the smile curve). Build all three cohort analyses (revenue, logo, engagement), keep the data monthly internally, and present it in investor materials as evidence of the retention dynamics the financial model is based on.

049 The Metrics That Matter at Pre-Seed vs. Seed vs. Series A

The Core Principle: Metrics Match Evidence Available

At each stage, investors expect to see metrics that are appropriate for the amount of data the company has. Presenting precise LTV:CAC calculations at pre-seed when there are no customers yet does not signal sophistication — it signals that the founder is confusing projections with evidence.

The reverse error is also common: presenting pre-seed-level metrics ("we have product-market fit signals from user interviews") at Series A, where investors expect quantified, cohort-validated evidence. Match the metrics to the evidence. Show what you have. Explain clearly what you do not have yet and how the current round will fund the data needed to answer those questions.

Pre-Seed: What Investors Are Actually Evaluating

At pre-seed, most companies have no revenue, no customers, and no retention data. The investor is evaluating:

Is the team capable of building this?

Is the problem real and the market large enough?

Does the founder understand how the business will work economically? **The metrics that matter at pre-seed:**

Metric	What it Shows
Total addressable market (TAM)	Is the opportunity large enough for a venture return?
Customer discovery interview	Has the founder validated the problem count

Number of letters of intent or	Is there early demand signal beyond pilot commitments
Time to first revenue	What does the founder believe the path (modelled)

rate, NRR. These should appear as projected assumptions in the model, clearly labelled as projections, not as measured metrics.

Seed: Early Evidence of Product-Market Fit

At seed, the company has typically launched, has early customers, and is beginning to understand whether the product solves the problem in a way that drives retention.

Investors at seed are asking: is there early evidence that this product has product-market fit, and is the unit economics hypothesis from the pre-seed model holding up in early data?

The metrics that matter at seed:

Metric	Threshold (directional)	What it Shows
MRR / ARR	Growing month over month	Revenue is real and recurring
MoM revenue growth	10–20%+ for early-stage	Trajectory of growth rate
Customer count	Enough for pattern	Not one or two recognition (20+)
Early NRR (even if 90% on early cohorts		Retention signal indicative)
CAC (directional)	Within reasonable range for Acquisition is not model	broken

Gross margin	In line with model for the	Unit economics are business type
Burn multiple	< 2 (ideally < 1.5)	Capital efficiency

useful seed-stage metric because it normalises capital efficiency to growth. A business burning £100k per month and adding £50k in net new ARR has a burn multiple of 2. Under 1.5 is considered capital-efficient for seed stage.

At seed, depth of retention data matters more than breadth of revenue. Two or three cohorts of 20+ customers each showing stable retention is more compelling than £500k ARR from 5 customers with no retention history.

Series A: Demonstrated Unit Economics and Scalable Growth Model

Series A is where the metrics conversation becomes rigorous. Investors at this stage are evaluating whether the business has repeatable, scalable growth with unit economics that improve (or hold) at scale.

The Transition Signals Between Stages

The clearest signals that a company is ready for the next stage:
Pre-seed to seed: First paying customers, early retention data (even thin), validation that the unit economics hypothesis is directionally correct from real transactions.

Seed to Series A: Cohort retention data showing NRR trending toward or above 100%, CAC calculation based on at least 6 months of actual acquisition data, gross margin in the model-appropriate range, ARR at or above £1M with a repeatable growth motion.

Series A to Series B: Magic number above 0.75 consistently, NRR above 110%, clear evidence of improving unit economics at scale (CAC

payback declining as the sales motion matures), ARR in the £5–15M range with strong year-over-year growth.

The metrics that matter change at each stage because the questions investors are asking change. Pre-seed: team and hypothesis. Seed: early signals of product-market fit with directional unit economics. Series A: demonstrated, cohort-validated unit economics with a repeatable growth model. Match the metrics to the evidence available. Do not present projections as measurements. Do not present seed-stage metrics at Series A. Know which metrics carry the most weight at your stage and make sure the model and the narrative address them directly.

050 How to Model Network Effects in a Marketplace

What Network Effects Are (and Are Not)

A network effect exists when the value of the product increases as more users join. There are two relevant types for marketplaces: **Direct network effects:** Each additional user makes the product more valuable for all existing users. Example: a professional network where more connections means more value for every member.

Indirect (two-sided) network effects: More supply makes the platform more valuable for demand, and more demand makes the platform more valuable for supply. This is the standard marketplace dynamic. More workers on a staffing platform means employers find staff faster. More employers means workers get more shifts.

What is not a network effect:

Growing faster because you have more marketing budget (scale economics, not network effects)

Improving the product with engineering resources as the company grows (product development, not network effects)

Lower CAC from brand recognition building over time (brand effects, not network effects)

The test: if you removed half the users from the platform tomorrow, would the product be significantly worse for the remaining users? If yes: network effects. If no: something else.

The Marketplace Flywheel Model

The two-sided marketplace flywheel works as follows:

More demand-side supply

The flywheel does not spin automatically. It requires sufficient density on both sides in a specific geography or category. A marketplace

that has broad coverage but thin density — many geographies, few workers or products in each — does not have network effects. It has fragmented supply that does not improve the demand-side experience.

The critical modelling question: at what supply density does the flywheel become self-reinforcing?

How to Model Network Effects in a Financial Model

Most financial models do not explicitly model network effects — they model the outcome (declining CAC, improving NRR, rising LTV) without attributing it to the mechanism. A model that explicitly models network effects requires three things:

1. Define the density threshold.

What is the minimum supply density in a given market segment or geography for the network effect to be active? For a local staffing marketplace, this might be: enough workers to fill 90% of open shifts within 4 hours. Below this threshold, the network effect is not yet active. Above it, each additional worker improves fill rate on the margin.

2. Model the supply-demand balance by market.

Build supply and demand counts by geography or category. When the supply:demand ratio in a specific market passes the density threshold, begin modelling the reduced CAC, improved NRR, or higher transaction frequency that the network effect produces.

3. Separate network effect benefits from organic growth.

The model should show the baseline unit economics without network effects and then show the incremental improvement as density passes the threshold. This makes the network effect assumption explicit and testable rather than baked into an improving growth rate curve with no attribution.

What Network Effect Claims Investors Scrutinise

Investors hear "network effects" in most marketplace pitches. The claims that hold up under scrutiny have evidence. The ones that do not have narrative.

Claims that hold up:

"Fill rate in markets where we have 200+ workers is 94%. In markets where we have under 50 workers, fill rate is 68%. This is the density effect."

"Customer NRR in our top 3 markets (highest density) is 118%. In markets we entered in the last 6 months, NRR is 95%. We attribute the gap primarily to density."

"Organic word-of-mouth CAC in mature markets is 40% of paid CAC in new markets. This is consistent with a network-driven referral effect at density."

Claims that do not hold up:

"Our platform gets better as more people join." (What does better mean, specifically?)

"We expect network effects to reduce CAC over time." (At what density? What is the data so far?)

"Our NRR improves as we scale." (Does it improve in specific markets as density increases, or is it improving globally as the product matures?)

The evidence standard for network effects is market-level data showing the relationship between density and the unit economic metric it is supposed to improve.

The Cold Start Problem: Modelling the Threshold

The cold start problem is the primary challenge for marketplace network effects: the flywheel does not spin until sufficient density is reached, but reaching that density requires investing before the flywheel produces returns.

Modelling the cold start requires:

Supply subsidy economics: What does it cost to seed supply density in a new market? For a staffing marketplace, this might be the cost of guaranteed minimum hours or premium rates for early workers. Model this as a market entry cost with an explicit density target and timeline. **Time to threshold:** How long does it take to reach network effect density in a new market? This determines the payback horizon for market expansion. If it takes 6 months and £50k in supply subsidy to reach threshold density in a new geography, that is the effective CAC for the market, and it should appear in the model.

Threshold sensitivity: What happens if the density threshold is higher than modelled? A market entry cost model that assumes 6 months to threshold should show what happens if it takes 9 months. This is the most important sensitivity for marketplace expansion models.

Network effects in marketplaces are real, powerful, and frequently overclaimed. Modelling them requires defining the density threshold at which they activate, building supply-demand balance by market, and separating the network effect contribution from organic growth. The evidence standard for investor conversations is market-level data showing the relationship between supply density and the unit economic outcome (fill rate, NRR, CAC) that the network effect is supposed to improve. Claims without this data are narratives. Claims with it are analyses. Be the founder with the analysis.

Segmented Cohort Analysis: The Full Framework

Layer 1: Structural Segmentation

Start by splitting cohorts along the dimensions that reflect genuinely different economic relationships with your product.

By customer type:

Segment	Why it Matters Analytically
SME	Higher churn, lower ACV, faster sales cycle, more price-sensitive, often self-serve
Mid-market	Mixed characteristics, often the most variable segment Enterprise

reveals whether a healthy blended NRR is masking a problematic SME segment being subsidised by strong enterprise retention. This is extremely common and almost invisible in aggregate data.

Layer 2: Health Status Classification

This is where it gets operationally powerful. Rather than just tracking whether a customer is active or churned, classify every active customer into a health state at each period.

The six health states:

Growing — MRR or usage is expanding month over month. Seat additions, upsell, usage growth. These are candidates for further expansion plays.

Stable — Revenue flat, engagement consistent. Healthy baseline. Lower near-term risk but not a growth driver.

Shrinking — MRR or usage declining but customer has not cancelled. Contraction revenue. Often the leading indicator of churn 60 to 90 days out. The most important state to catch early.

At-risk / Inscribed — Your terminology "inscribed" is

interesting; if you mean customers showing the specific behavioural precursors to churn (declining logins, unresolved tickets, champion has left the account), this is the intervention window. Action here prevents the churn.

Churned — Cancelled. But segmenting churned customers by *why* they left (price, product fit, competitor, company closure, champion departure) is what turns churn data into product and commercial intelligence.

Revived — Previously churned, reactivated. Tracking revived cohorts separately is valuable because their second-lifetime retention profile is often different from their first. They also tell you something about whether the product improved enough to be worth returning to.

Layer 3: The Signal Stack

The health state classification becomes far more precise when you pair it with multiple signal types:

Quantitative product signals (your Uber example is perfect): For Uber: wait time, price per mile, driver availability. For a SaaS product: DAU/MAU ratio, feature depth (how many modules actively used vs. licensed), API call volume trend, time-to-value in onboarding, support ticket frequency and resolution time.

The key is identifying which metrics have *statistically validated* correlation with churn in your specific product rather than using generic proxies. The methodology:

1. Take all customers who churned in the last 12 months
2. Pull their metric values for the 30, 60, 90 days before churn 3. Compare to the same metrics for customers who renewed
4. Identify which metrics diverge most clearly between the two groups
5. Those are your leading churn indicators — build them into the health score

Qualitative signals:

NPS and CSAT are the most structured. But their value multiplies when they are routed back to the health state classification:

A Stable customer who gives an NPS of 6 is a different risk from a Stable customer who gives an NPS of 9

A Shrinking customer who gives an NPS of 8 may be shrinking due to budget constraints rather than product dissatisfaction — a fundamentally different intervention

A Revived customer with an NPS of 7 is still a watch item despite reactivation

Account manager qualitative input is underused in most analytics systems. AMs often know things the data does not capture yet: the procurement freeze, the champion who is leaving next month, the pilot of a competitor's product. A structured "AM health override" field in the CRM, with a required reason code, integrates this into the classification systematically rather than leaving it in heads and emails.

Layer 4: The Combined View

When you pull all of this together, a single cell in the cohort table goes from a number (85% retention) to a structured picture: Enterprise cohort, Q1 2024: 82% revenue retained

- 40% of cohort in Growing state (driven by seat expansion in accounts using 3+ modules)
- 35% Stable
- 15% Shrinking (mostly single-module users with declining login frequency)
- 10% Churned (3 of 4 churned accounts: champion departure as the primary trigger)
- NPS for this cohort: 42 average; Shrinking sub-segment NPS: 28 - AM flags: 2 accounts in Shrinking have confirmed competitor pilots in progress

That is a completely different level of analytical richness than "82% retention." It tells you exactly where to focus, what the intervention looks like, and which risk factors are product versus commercial versus relationship-driven.

The Practical Build

To implement this you need three things working together:

A customer health score model — built in SQL or your BI tool, pulling product usage, billing, and support data, producing a health state classification for every customer every month. The classification logic should be version-controlled and auditable.

A structured qualitative data collection process — NPS with follow-up reason codes, CSAT tied to specific interaction types, AM health override with mandatory reason codes in CRM. None of these are useful if the data collection is inconsistent.

A cohort table that can be filtered by segment and health state — the simplest implementation is a pivot table in Excel or Google Sheets pulling from the SQL output. The more scalable version is a BI tool (Metabase, Mode, Looker) where the segment and health state filters are UI elements rather than manual rebuilds.

The insight this generates is not just analytical. It becomes the briefing document for customer success, the evidence base for product prioritisation, and the risk register that surfaces in board reporting. That is where cohort analysis stops being a finance metric and becomes a cross-functional operating tool.

051 Churn Rate: Every Way to Calculate It (And Why the Method You Choose Changes the Story You Tell)

Why Churn Has So Many Definitions

Churn measures how much you are losing. But "losing" means different things depending on what you count. Losing a customer (logo churn) is different from losing revenue (revenue churn), which is different from losing revenue net of expansion (net revenue churn). A company can have 5% logo churn and negative revenue churn if expanding customers more than offset the departures.

This is not academic. During the Indeed Flex fundraising processes, investors asked for churn in at least three different ways. Creandum wanted to see cohort-level revenue retention. Profounders asked for monthly logo churn. B2Ventures wanted annual net revenue retention. If we had only prepared one version, two of those three conversations would have stalled.

The Six Churn Metrics That Matter

1. Monthly Logo Churn

Formula: Customers lost in the month / customers at the start of the month.

Example: You start January with 200 customers. 8 cancel. Monthly logo churn = 8 / 200 = 4%.

This is the simplest and most commonly reported churn metric. It tells you what percentage of your customers leave each month, regardless of how much they were paying. Its weakness: it treats a $50/month customer and a $5,000/month customer as equal. A

company losing its smallest accounts has a very different problem than one losing its largest.

2. Monthly Revenue Churn (Gross)

Formula: MRR lost from churned and contracted customers / MRR at start of month.

Example: You start the month with $100K MRR. $3K churns completely. $1K downgrades. Gross revenue churn = ($3K + $1K) / $100K = 4%. This is more informative than logo churn because it weights losses by revenue. If your highest-paying customers are the ones leaving, revenue churn will be higher than logo churn, which is a much more urgent signal. If only your smallest accounts churn, revenue churn will be lower than logo churn, which is uncomfortable but less dangerous.

3. Net Revenue Churn (Net MRR Churn)

Formula: (MRR lost from churn and contraction - MRR gained from expansion) / MRR at start of month.

Example: You lose $4K in churn and contraction but gain $6K from upsells and seat expansion. Net revenue churn = ($4K - $6K) / $100K = -2%.

Negative net revenue churn means your existing customer base is growing without any new customer acquisition. This is the holy grail metric for SaaS and recurring revenue businesses. Companies like Slack, Twilio, and Datadog famously had negative net churn for years, meaning the business would grow even if they stopped acquiring new customers entirely.

4. Net Revenue Retention (NRR)

Formula: (Starting MRR - churn - contraction + expansion) / Starting MRR. Or equivalently: 1 minus net revenue churn.

Example: From the above, NRR = ($100K - $4K + $6K) / $100K = 102%. NRR above 100% means your existing customers are worth more this month than last month. This is the metric Series A and growth investors focus on more than any other. Bessemer's State of the Cloud report shows that top-quartile SaaS companies have NRR of 120%+.

Median is around 105-110%. Below 100% means your customer base is eroding, and you are on a treadmill where new acquisition must outrun existing loss.

5. Annual Churn Rate

Formula: There are two approaches, and they produce different numbers. Approach A (simple annualization): 1 - (1 - monthly churn rate)\^12. With 4% monthly churn: 1 - (0.96)\^12 = 39.2% annual churn. This compounds the monthly rate and is mathematically correct.

Approach B (direct annual measurement): Customers lost over 12 months / customers at start of the period. This can differ from Approach A if churn rates are not constant across months (seasonality, product changes, pricing changes).

Use Approach A for projections in your financial model. Use Approach B for reporting historical performance. Always specify which method you are using, because 4% monthly sounds manageable while 39% annual sounds alarming, even though they are the same number.

6. Cohort-Based Churn

Formula: For each acquisition cohort, track the percentage of original revenue (or customers) retained in each subsequent month.

This is the most powerful churn analysis because it captures time-dependent patterns that aggregate metrics miss. Early-stage churn (months 1-3) is almost always higher than steady-state churn (months 6+). A company with 8% churn in Month 1 and 2% churn in Month 12 has a very different retention story than one with a constant 4%, even though the blended average might look similar.

At Indeed Flex, cohort-based churn was the metric that mattered most during the exit. The acquirer wanted to see that newer cohorts retained better than older ones (proving the product was improving) and that mature cohorts stabilized (proving there was a retention floor). Both patterns were visible in the cohort data and invisible in aggregate churn.

> **Key insight: The churn metric you report should match the story you are telling. Reporting logo churn when your revenue*

> *churn is worse looks like you are hiding something. Reporting gross churn when your net churn is negative means you are underselling your business. Know all six numbers and present the ones that are most relevant to the conversation, with the others available when asked.**

How Churn Differs by Business Model

Business Model	Typical Churn Characteristics
B2B SaaS (enterprise)	1-2% monthly logo churn, NRR 110-130%
B2B SaaS (SMB)	3-7% monthly logo churn, NRR 90-110%
B2C subscription	5-10% monthly logo churn, expansion limited
Marketplace (demand side)	3-8% monthly, varies by segment Marketplace (supply side)

higher, contracts are longer, and the product is embedded in workflows. SMB SaaS has higher churn because smaller companies go out of business, have less budget commitment, and switch products more easily. Marketplaces have asymmetric churn: the paying side (usually demand) behaves more like SaaS, while the supply side has irregular activity patterns that look like churn but may be seasonal dormancy.

The Churn Traps That Mislead Founders

Averaging churn across segments. Your enterprise churn is 1% monthly and your SMB churn is 8% monthly. Blended, it looks like 4%. But these are two completely different businesses with different dynamics. Report and model churn by segment, not blended.

Ignoring contraction as churn. A customer who downgrades from $5K/month to $1K/month has not churned by logo count, but you lost $4K in MRR. Revenue churn captures this. Logo churn does not. **Not distinguishing voluntary from involuntary churn.** Involuntary churn (failed payment, expired card) is a billing operations problem with known solutions. Voluntary churn (customer actively cancels) is a product or value problem. Mixing them hides the root cause. **Measuring churn on too short a window.** One month of churn data is noise. Three months is a trend. Six months is a pattern. Do not make strategic decisions on a single month's churn spike, which might be driven by one large customer or a billing cycle anomaly.

Investor Perspective: Which Churn Metric at Which Stage

Churn is not one metric. It is six related metrics that each tell a different story about your customer base. Logo churn counts customers lost. Revenue churn weights by value. Net revenue churn includes expansion. NRR expresses the health of your existing base as a single number. Annual churn compounds the impact over time. Cohort churn reveals time-dependent patterns invisible in aggregates. Know all six, model the ones that matter for your business, and present the right metric for the right conversation. The founders who understand their churn dynamics at this level of granularity are the ones investors trust with growth capital.

Part Four: Finance Operations & Analytics

Running a funded company requires a different kind of financial discipline than raising one. This part covers the metrics that matter in practice: what to track, how to build dashboards that get used, and how to read the leading indicators that give you three months of warning before a problem reaches your revenue.

067 The 10 KPIs Every Founder Should Track (And the 20 They Should Ignore)

The Criteria for a KPI That Matters

A KPI is only useful if it passes three tests:

Actionable. If the number changes, you know what to do. If MRR drops, you investigate churn and acquisition. If "number of Slack messages sent" drops, you do nothing, because it does not connect to a business decision.

Connected to the model. Every KPI should map to an assumption or output in your financial model. If it does not appear in the model, it is either a vanity metric or you have a gap in your model. **Understandable to investors.** Your KPIs should be communicable in one sentence. If you need a paragraph to explain what a metric means and why it matters, it is not a KPI. It is an internal diagnostic.

The 10 KPIs

1. Monthly Recurring Revenue (MRR) or Monthly GMV

What it measures: the total recurring revenue generated each month (SaaS) or the total transaction value facilitated (marketplace). This is the top-line heartbeat of the business. Every investor conversation starts here.

Why it matters: MRR trend tells you whether the business is growing, flat, or shrinking. Decompose it into new MRR, expansion MRR, churned MRR, and contraction MRR for the full picture.

2. MRR Growth Rate (Month-over-Month)

What it measures: the percentage change in MRR from last month. Early-stage startups are evaluated on growth trajectory more than absolute revenue. A 15% MoM growth rate from $20K MRR is a very

different story from 3% MoM from $500K MRR, and both are relevant. Benchmark: pre-Series A, top-quartile SaaS companies grow 15-25% MoM. Post-Series A, 8-15% MoM is strong. Anything below 5% MoM for more than 3 consecutive months at early stage is a warning signal.

3. Burn Rate and Runway

What it measures: how much cash you spend per month (burn) and how many months of cash remain (runway). This is the existential KPI. Everything else is secondary if you are running out of money.

Track net burn (total cash out minus total cash in) rather than gross burn (total cash out) because net burn reflects the revenue you are generating. Runway = current cash / net burn.

4. Customer Acquisition Cost (CAC) by Channel

What it measures: how much it costs to acquire one customer, broken down by marketing channel. Blended CAC hides which channels work and which destroy value. Investors at Series A will ask for channel-level data. Model implication: CAC by channel feeds directly into your marketing budget allocation and your revenue projections (more efficient channels should get more spend).

5. LTV:CAC Ratio

What it measures: the relationship between the lifetime value of a customer and the cost to acquire them. The standard benchmark is 3:1 minimum for a healthy business. Below 3:1, you are likely losing money on each customer when fully loaded. Above 10:1, you may be underinvesting in growth.

Calculate this by segment and channel, not just as a blended number. Your enterprise segment might be 8:1 while your SMB segment is 1.5:1. Those require completely different strategies.

6. Net Revenue Retention (NRR)

What it measures: how much revenue from your existing customer base grows or shrinks over time, independent of new customer acquisition. NRR above 100% means your existing customers are becoming more valuable. Below 100% means you are on a treadmill.

This is the single most valued metric at Series A and beyond. Bessemer and OpenView both rank it as the number one predictor of long-term enterprise value for SaaS businesses.

7. Monthly Churn Rate (Logo and Revenue)

What it measures: the rate at which customers and revenue are leaving. Track both logo churn (percentage of customers lost) and revenue churn (percentage of MRR lost), because they tell different stories. High logo churn with low revenue churn means you are losing small accounts. Low logo churn with high revenue churn means your big customers are downgrading.

8. Gross Margin

What it measures: the percentage of revenue remaining after cost of goods sold. This determines how much of every dollar you can invest in growth, operations, and profitability. SaaS companies should target 70-85%. Marketplaces on net revenue should target 50-75%.

Gross margin trend matters as much as the absolute number. Improving margins at scale is a strong signal. Declining margins suggests structural delivery cost problems.

9. Cash Conversion Cycle

What it measures: how many days it takes for cash spent to convert back into cash received. This combines days sales outstanding (how long customers take to pay), days payable outstanding (how long you take to pay vendors), and inventory days (if applicable).

For SaaS companies with upfront annual billing, the cash conversion cycle can be negative (you receive cash before you deliver the service), which is a significant advantage. For companies with net-60 enterprise payment terms, the cycle can be 60-90 days, which means you need 2-3 months of working capital that is invisible in a P&L-only view. **10. Payback Period**

What it measures: how many months until a new customer has paid back their acquisition cost. CAC / (monthly revenue per customer x gross margin). Best-in-class SaaS companies target under 12 months. Most early-stage startups run 18-24 months.

Payback period is the bridge between unit economics and cash flow. A 24-month payback means you need 24 months of cash to fund each new customer's acquisition before they become profitable. If you are growing fast, that cash requirement compounds quickly.

The Metrics You Should Ignore (For Now)

Metric	Why It Is Noise at Early Stage
Total registered users	Vanity. Active users is what matters.
Page views / sessions	Vanity unless conversion rate is attached.
Social media followers	Zero correlation with revenue. NPS score alone
Employee satisfaction score	Important for culture, not for the investor model.

> *Key insight: The best KPI set is the smallest set that covers the entire value chain: acquisition (CAC), monetization (MRR, ARPA), retention (churn, NRR), economics (gross margin, LTV:CAC, payback), and survival (burn, runway). If you can recite these 10 numbers from memory and explain the trend in each, you understand your business better than 90% of founders at your stage.*

Track 10 KPIs that cover the entire value chain of your business: MRR, MRR growth, burn and runway, CAC by channel, LTV:CAC, NRR, churn (logo and revenue), gross margin, cash conversion cycle, and payback period. Each connects directly to your financial model and to the questions investors will ask. Ignore vanity metrics that do not drive decisions. Review weekly and monthly. Report quarterly to the board. The discipline of a focused KPI set produces better decisions, faster fundraises, and a founder who can defend every number in the model because they live with the data every week.

068 How to Build a Metrics Dashboard That You Will Actually Use Every Week

Why Most Dashboards Fail

The pattern is predictable. A founder decides they need a dashboard. They spend a week building an elaborate 30-metric view with charts, filters, and drill-downs. They check it daily for two weeks. Then a fundraise starts, or a product launch absorbs their attention, and the dashboard goes stale because half the metrics require manual data entry. The failure is not in the tools. It is in the design. A dashboard that requires manual input is a report that will eventually stop being produced. A dashboard with 30 metrics is a data dump, not a decision tool. A dashboard that does not show variance from plan is a rearview mirror without a speedometer.

The Three-Layer Dashboard Architecture

Build three views, not one. Each serves a different purpose and audience.

Layer 1: CEO Daily View (5 metrics)

This is the view you check every morning. It should load in under 3 seconds and answer one question: is anything on fire?

The 5 metrics: today's cash balance, trailing 7-day new customers (or new revenue), trailing 7-day churn count, current MRR, and days of runway remaining. These five numbers tell you whether the business is alive and whether any trend requires immediate action.

This layer should be automated with zero manual input. Connect it directly to your database, your payment processor (Stripe, for example, has excellent APIs), and your banking data (via Plaid or manual daily update if needed).

Layer 2: Team Weekly View (10-12 metrics)

This is the view the leadership team reviews every Monday. It covers the full KPI set: MRR and growth rate, CAC by top 3 channels, new customers this week and month, churn this month, gross margin, burn rate, pipeline value (if applicable), and the top 3 milestones with status. The weekly view should include a comparison to plan. For each metric, show: actual, plan, and variance (both absolute and percentage). Color-code: green if within 10% of plan, yellow if 10-20% off, red if more than 20% off. This visual system lets the team scan the dashboard in 2 minutes and focus discussion on red items.

Layer 3: Investor Monthly View (10-12 metrics with trends) This is the view you share with investors in your monthly update. It includes the same metrics as the weekly view but with 6-12 months of historical trend data and narrative context.

The investor view should be exportable as a PDF or built into your monthly update template. It should show each KPI with a small trend chart (sparkline), the current value, the plan value, and a one-line note on any material variance. This becomes the financial section of your investor update.

Tool Selection: What to Use at Each Stage

Stage	Recommended Tool
Pre-seed / early seed	Google Sheets with automated data imports
Seed with revenue	Metabase (free, open-source) or Notion dashboards
Post-seed with engineering support	Metabase, Looker, or Mode Analytics Series A+

combination of Google Sheets (for the financial model and plan comparison) and a simple BI tool like Metabase (for automated metric

tracking from the database). This costs $0 and takes a day to set up.

Do not buy a $2K/month BI platform at seed stage. The data volume does not justify it.

The Build Process

Step 1: Define the metric list

Use the 10 KPIs from the previous article. Remove any that you cannot currently measure. Add any that are business-specific and material (for a marketplace: fill rate, supply density. For e-commerce: average order value, return rate).

Step 2: Identify data sources

For each metric, identify where the data lives. MRR: payment processor or database. CAC: marketing platform + database. Churn: database. Cash: bank account or accounting software. Map each metric to a source and determine whether it can be automated or requires manual entry. Maximize automation.

Step 3: Build the plan baseline

A dashboard without a plan is just a collection of numbers. Export your financial model's monthly projections for each KPI and load them as the comparison baseline. Every actual metric should display alongside its planned value. This turns the dashboard from a monitoring tool into a performance management tool.

Step 4: Set alerts

Configure alerts for the metrics that require immediate action. Cash below a threshold: alert. Weekly churn above 2x the average: alert. MRR declining for 2 consecutive weeks: alert. These alerts should go to the CEO's phone or Slack, not buried in an email digest.

Step 5: Establish the review cadence

Daily: CEO checks the 5-metric view. Weekly: leadership team reviews the full dashboard in a 30-minute meeting. Monthly: CEO compiles the investor view and sends the update. Quarterly: full variance analysis against the financial model with model recalibration.

> **Key insight: The best dashboard I have worked with was at Indeed Flex after we installed an automated daily metrics email. Six numbers, every morning at 8am, color-coded against plan. No login required, no dashboard to open. The CEO read it on his phone during his commute. When something was red, we discussed it by 10am. When everything was green, we moved on. Simplicity is the feature that makes a dashboard sustainable.**

A useful dashboard has three layers: a daily CEO view (5 survival metrics), a weekly team view (10-12 KPIs with plan comparison), and a monthly investor view (trends and narrative). Automate as much as possible. Show actuals versus plan, not just actuals. Set alerts for critical thresholds. Establish a review cadence that ties the dashboard to decisions. Use simple tools: Google Sheets and Metabase will cover 90% of needs through Series A. The goal is not a beautiful dashboard. It is a dashboard you check every day and that changes how you run the company.

069 Leading vs. Lagging Indicators: How to See Problems Three Months Before They Hit Your Revenue

The Difference Explained Simply

A lagging indicator measures an outcome. A leading indicator measures a driver of that outcome. Revenue is lagging; it is the result of sales activity that happened weeks or months ago. Pipeline value is leading; it predicts what revenue will be in the coming months. Churn is lagging; the customer already left. Product usage decline is leading; it predicts that the customer is about to leave.

The distinction matters because you can only act on leading indicators. By the time a lagging indicator changes, the opportunity to intervene has passed. A customer who churned last month cannot be unchurned. But a customer whose product usage dropped 50% this month can be saved with a proactive outreach before they reach the cancellation decision.

Lagging Indicator (What	Leading Indicator (What Will
Happened)	Happen)

The Leading Indicators That Predict Revenue

1. Qualified Pipeline Value

For any business with a sales process, the pipeline is the most predictive leading indicator of future revenue. If your average sales cycle is 45 days and your current pipeline is $200K in qualified opportunities, you have a reasonable forecast of what next month and the month after will look like.

Track pipeline by stage and apply historical conversion rates: if 30% of qualified opportunities close, $200K in pipeline predicts roughly

$60K in new revenue over the next 1-2 months. When pipeline drops, revenue will follow 30-60 days later. That lag is your intervention window. **2. Demo Booking Rate / Trial Activation Rate**

For product-led or demo-driven businesses, the rate at which visitors convert to demos or activate trials predicts customer acquisition 30-90 days out. A declining demo rate with stable traffic means your messaging or landing page is degrading. A declining demo rate with declining traffic means your top-of-funnel is weakening.

At Indeed Flex, the leading indicator for employer acquisition was the rate at which new employer sign-ups posted their first shift within 7 days. If that activation rate dropped, we knew new employer revenue would soften 4-6 weeks later, giving us time to investigate and respond.

3. Product Usage and Engagement

For SaaS and platform businesses, how actively customers use the product predicts retention. Customers who log in daily churn at a fraction of the rate of customers who log in weekly, who in turn churn less than customers who log in monthly. The usage pattern is the earliest warning of future churn.

Define your activation and engagement thresholds: what does a "healthy" usage pattern look like? Then track the percentage of customers who meet that threshold. A declining percentage predicts rising churn 60-90 days later.

4. NPS or Customer Satisfaction

Net Promoter Score is a crude but useful leading indicator of retention. Customers who rate you 9-10 (promoters) rarely churn. Customers who rate you 0-6 (detractors) churn at 3-5x the rate. Monitoring NPS shifts, particularly a growing detractor segment, gives you 30-90 days of warning before churn spikes.

NPS alone is insufficient because it does not explain why. Always pair NPS with a follow-up question that identifies the specific issue. "Why did you give that score?" produces actionable data. A number alone does not.

5. Hiring Pipeline Quality

This is the leading indicator that most founders overlook. If your model assumes hiring 4 engineers in Q2 and your recruiting pipeline has 2 qualified candidates, you will miss your hiring targets, which delays product milestones, which delays revenue targets. Recruiting pipeline is a leading indicator for operational capacity.

How to Use Leading Indicators in Your Financial Model

The best financial models incorporate leading indicators as input assumptions that update monthly. Instead of projecting "50 new customers next month" as a static assumption, the model should project: current pipeline ($300K) x historical close rate (20%) / average ACV ($6K) = 10 new customers. This makes the model responsive to pipeline changes.

Similarly, churn in the model can be linked to leading indicators. If product usage data shows 15% of customers dropped below the engagement threshold this month, and historically 40% of disengaged customers churn within 90 days, the model projects 6% incremental churn over the next quarter.

This approach turns the model from a static projection into a dynamic forecast that updates as leading indicators move. It also gives you a structured way to explain forecast changes to investors: "We revised Q3 revenue down by 8% because pipeline entering Q2 was 20% below plan, and our historical conversion rate suggests that flows through to 8% fewer closes." *Key insight: The founders who manage by leading indicators make smaller, earlier course corrections. The founders who manage by lagging indicators make larger, later ones. In a startup with 18 months of runway, the difference between catching a problem in Month 3 versus Month 9 is the difference between a strategic adjustment and a crisis.*

Building a Leading Indicator System

Step 1: For each lagging KPI, identify 2-3 leading drivers
Revenue is driven by pipeline, conversion, and ACV. Churn is driven by usage, satisfaction, and support interactions. CAC is driven by traffic

volume, conversion rate, and spend level. Map these relationships explicitly.

Step 2: Establish the time lag

How far ahead does each leading indicator predict? Pipeline predicts revenue 30-60 days ahead. Usage decline predicts churn 60-90 days ahead. Demo rate predicts customer acquisition 30-45 days ahead. Knowing the lag tells you how much intervention time you have.

Step 3: Set thresholds and alerts

When a leading indicator drops below a threshold, trigger an investigation. Pipeline drops 20% week-over-week: sales team reviews sourcing. Activation rate drops below 50%: product team reviews onboarding. Feature usage for a key segment drops 15%: customer success team reaches out.

Step 4: Track predictive accuracy

Over time, measure how accurately each leading indicator predicts the lagging outcome. If pipeline value predicts revenue within 15% accuracy, it is a reliable indicator. If NPS has no observable correlation with churn in your business, it is not useful as a leading indicator and should be replaced with something that is.

Leading indicators predict what your lagging KPIs will look like in 30-90 days. Pipeline predicts revenue. Usage predicts churn. Demo rates predict acquisition. Engagement predicts retention. Build a system that maps each lagging KPI to its leading drivers, establishes the time lag, sets alert thresholds, and tracks predictive accuracy. Integrate leading indicators into your financial model so your forecast updates as real-time signals move. The competitive advantage is time: founders who see problems 3 months early have 3 months to fix them. Founders who see problems when they hit revenue have zero months, and they are the ones who end up raising emergency bridges.

076 You Do Not Need to Be a Genius With Excel Anymore. AI Does That Now.

How Is AI Changing the Job of the Financial Analyst?

Category	Detail
Tasks most impacted by AI	Formula auditing, scenario building, narrative generation
Tasks least impacted by AI	Assumption judgment, business understanding, investor relationships
Time saving on model construction	40–70% on mechanical tasks (observed, not benchmarked)
Biggest unlock	Breadth: doing things in hours that previously took days
Risk of over-reliance	AI confidently produces wrong numbers without flagging it

The Tasks Where AI Actually Helps (And Is Not Just a Gimmick)

There is a lot of noise about AI in finance. Here is what is actually useful right now, based on applying these tools to real models across multiple companies:

Formula auditing and error checking.

Ask an AI to review a financial model and identify potential circular references, broken links, or logic errors. This used to take a modeller hours of careful tab-by-tab work. With the right prompt and a pasted section of model structure, it takes minutes. The AI is not always right, but it surfaces candidates for review faster than manual auditing.

Scenario and sensitivity analysis construction.

Building a clean scenario toggle that switches between conservative, base, and aggressive assumptions across a full model is fiddly and error-prone when done manually. AI can generate the structural logic, the named ranges, the switch formulas, and the summary tab in a fraction of the time. The analyst still needs to define the scenarios. The mechanics become much faster.

Translating model outputs into investor narratives.

One of the most time-consuming tasks in fundraising prep is writing the commentary that explains the numbers: why the model makes the assumptions it does, what the key sensitivities are, how the business dynamics drive the forecast. AI can generate a strong first draft of this commentary from a structured model input. The analyst edits and adds judgment. The blank page problem goes away.

Market research and comparable analysis.

Building a comparable company set for benchmarking unit economics or valuation multiples previously required hours of manual data collection. AI tools with web access can produce a first-pass comparable set in minutes, which the analyst then refines and validates.

Due diligence preparation.

Generating a list of 30 investor diligence questions mapped to specific model tabs, then drafting the answer framework for each one: this is a task AI does well and that previously required experienced senior finance time. A founder using AI for this arrives at investor meetings better prepared than one who did not.

> *Key insight: The value of AI in financial analysis is not about replacing the analyst. It is about eliminating the ceiling on what one analyst can do alone. The limiting factor shifts from hours in the day to the quality of judgment about which questions to ask.*

How to Actually Use AI in a Financial Model Workflow

The difference between using AI and using AI well is prompt quality. Here is the workflow that produces consistent results:

Step 1: Structure the model first, use AI second.

AI is a powerful collaborator once the model structure is defined. It is less useful when asked to design the structure from scratch, because it does not know your business. Define your drivers, your revenue logic, your cohort approach. Then use AI to build the mechanics faster. **Step 2: Use AI for formula generation, not for assumption setting.** Tell AI: "Here is my revenue model structure. Generate the Excel formulas for a monthly cohort calculation that outputs NRR by period." Do not tell AI: "What should my churn assumption be?" The first question has a mechanical answer. The second requires business judgment that AI cannot have.

Step 3: Use AI to audit what you built.

After building a section of the model, paste the logic into an AI tool and ask: "Are there any logical errors in this revenue calculation? What edge cases might break this?" This is faster and often more thorough than self-auditing.

Step 4: Use AI for the investor narrative.

Once the model is final, use AI to generate the first draft of the assumptions commentary: "Here are my key model assumptions and the rationale behind each. Generate an investor-facing explanation of this model's logic and key sensitivities in 500 words." Edit heavily. Do not publish first drafts.

What AI Cannot Do (And Where the Judgment Still Lives)

There is a specific failure mode in using AI for financial analysis that is worth naming directly: AI produces outputs that look correct and are wrong. A formula that is syntactically valid but logically backwards. A market sizing calculation that uses the right structure but misunderstands what the numbers represent.

This is not a reason to avoid AI. It is a reason to apply judgment to everything it produces.

The things that still require human judgment:

Task	Why AI Cannot Own It
Which assumptions drive the	Requires knowing the business from business
Whether a forecast is credible to	Requires knowing this specific investors
How to defend a number under	Requires owning the reasoning, not pressure
Identifying what an acquirer will	Requires M&A pattern recognition focus on

judgment to the outputs is significantly more valuable than the analyst who refuses to use AI, or the one who uses AI and does not check the work.

Specific AI Tools That Are Useful for Finance Work Right Now

This is not a sponsored list. These are tools that are actually useful based on direct use:

Claude and ChatGPT: For narrative generation, formula logic, model structure review, and diligence question preparation. The quality of output depends almost entirely on the quality of the prompt. **Perplexity:** For market research and comparable data. Better than a standard search for structured financial queries because it aggregates and cites sources.

Notion AI / Coda AI: For turning model commentary and board update drafts into polished documents faster. Useful for the reporting cycle. **Excel Copilot (Microsoft 365):** For formula generation and

data cleaning directly in the spreadsheet environment. Still early but improving quickly.

The analyst who has built a personal library of prompts that work is operating at a structurally different level than one who is prompting from scratch each time. Treat prompt development as a skill worth investing in.

The Question Worth Asking

The most useful framing for a financial analyst or founder thinking about AI is not "will AI replace me?" It is: "What could I do if I had 40 extra hours this week?"

Because that is roughly what the mechanical automation of model tasks creates. The question is whether those hours go toward more analysis, more investor preparation, more business understanding, or back into operational noise.

The analysts and finance leaders who use this time to go deeper on the judgment work are the ones who become harder to replace. The ones who use AI to do the same work faster without going deeper are only delaying the question.

AI has changed the ceiling on what a solo financial analyst can produce. Formula auditing, scenario construction, narrative generation, and comparable analysis are all materially faster with AI than without. The tasks that AI cannot own are the judgment calls: which assumptions are right, which numbers will be challenged, how to defend the model under pressure. The analysts who use AI for the mechanical work and apply sharp judgment to everything it produces are operating at a level previously only possible with a larger team. The tool is here. The question is whether you are using it.

WAR STORIES

PART

Part Five: War Stories

Frameworks are useful until the call comes on a Friday evening and someone says: we have a problem. This section is about what happens when the model has to work in real conditions, under real pressure, with real consequences.

091 How I Saved a more than 9 figures Exit on a Saturday Morning (While My Girlfriend Waited)

What Is Client Cohort Value and Why Did It Save an Exit?

Client cohort value is the projected lifetime revenue extractable from a specific group of customers acquired in the same period, adjusted for observed spending growth within that cohort over time. When cohorts show increasing spend per client as they mature, standard LTV calculations understate the true value of the customer base. That understatement can be corrected, and in the right context, it changes a valuation conversation entirely.

Company	Syft (acquired by Indeed, now Indeed Flex)
Exit size	more than 9 figures+
Revenue multiple	Undisclosed
Diligence duration	~6 months
The problem valuation cut	Actuals below forecast; acquirer proposed 20%
The fix	Growth-adjusted cohort valuation reframe

The Setup: Six Months of Due Diligence

Nobody tells you how long a sale process actually takes. Six months sounds manageable until you are living it, knowing that at any point the acquirer can walk away, find something they do not like, or simply decide the price is wrong.

During the sale of Syft to Indeed, I was the finance lead responsible for the model that underpinned the valuation. We had been through seed rounds, a Series A, venture debt, and a convertible note together. The model had survived multiple investor diligence processes. I was, if not confident, at least not afraid of questions.

Then the call came.

The Oh-Sh*t Moment

The acquirer had been running forensic diligence on the financials for months. Standard process. Then they came back with a finding: actual gross profit performance was below the forecasted figures in the model. Their logic was clean and, honestly, fair: if the company performed X% below forecast, the valuation should be reduced by X%. In a world where you are being valued on a revenue multiple, a performance shortfall translates directly into a lower price. No drama, no bad faith. Just maths.

What they did not frame it as was: the forecast was always aspirational. Any founder or finance professional knows that early-stage forecasts are directional. The rule of thumb across every raise I have worked on is that 80% of forecast is considered strong execution. Missing by a moderate amount is not failure. It is a startup behaving like a startup. But in an M&A process, "it's a startup, these things happen" is not an argument. You need a model.

I was not in the room when this landed. The founders heard it on a call. The now-CEO rang me that Friday evening, and I remember exactly where I was. Walking across a bridge in London. He said: "Hey, how are you doing. We have a problem. Can you come in tomorrow?"

No urgency. Very calm. The kind of calm that means it is serious. I said I could come in that evening.

He said tomorrow was fine.

It was not fine. But I went home, and I spent the evening thinking.

The Saturday Morning Fix

I arrived at the Holborn office the next morning. He told me what had happened and said three words: "Fix it."

Not: "See if there's anything we can do."

Not: "Take a look and let me know."

"Fix it."

What followed was a few hours of going back through the model looking for something defensible. Not something invented. Something real that had been underweighted in the original valuation framework. What I found was this: when you look at client cohorts individually rather than in aggregate, the spending behaviour of mature clients was increasing over time. Clients were not just staying. They were spending more. And the original valuation had not captured that growth trajectory in the way it deserved to be captured.

The parallel I reached for was the PEG ratio: the

Price/Earnings-to-Growth ratio used in public equity analysis. The standard P/E ratio tells you what the market is paying for current earnings. PEG adjusts for growth, because a company growing at 30% per year deserves a higher multiple than one growing at 5%, even if their current earnings look identical.

The same logic applied here. A client base where cohort spending is growing 25% year over year is worth more than a client base at flat spending, even if the current revenue number is the same. The standard LTV calculation was not capturing that growth premium.

So I built the case: our clients were not just loyal, they were accelerating. That changed the effective LTV. Which changed the revenue the acquirer could expect to extract from the existing client base. Which meant the valuation should be higher, not lower, relative to the standard framework.

The numbers worked out to roughly a 25% uplift in expected client value. Not a rescue. A correction.

> *Key insight: In M&A diligence, the acquirer sets the framing by default. The only way to change the outcome is to show up with a different, defensible framework before the conversation*

> *closes. A financial model that has no capacity for reframing is a liability in the negotiating room.*

Why This Worked (And When It Would Not)

This reframe worked for three reasons, and it is worth being honest about all of them.

1. The underlying data was real.

The cohort spending growth was an observed pattern in actual transaction data, not an assumption. If I had tried to argue for higher LTV on the basis of forecasts or estimates, it would not have landed. The model has to point to something real.

2. The logic had a credible precedent.

Using growth-adjusted valuation frameworks is standard in public equity analysis. The PEG parallel gave the argument a respectable intellectual foundation. It was not invented for the occasion.

3. The acquirer was motivated to close.

This matters. A motivated acquirer wants to find a way to make the numbers work. An acquirer who is looking for an exit will use any gap to walk. Reading the room correctly is part of the job.

What would not have worked:

Arguing that the forecast was "aspirational" and should not be taken literally. That is true, but saying it concedes the framing. You are essentially agreeing that the variance is a problem and asking for sympathy. That is not a negotiating position.

What This Means for Your Model

Most founders stop thinking about their financial model once the round is closed or the LOI is signed. That is understandable. It is also wrong.

The model that raises your round is the same model that gets interrogated during M&A diligence. It is the document that will be used to argue for or against your valuation when there is real money on the

table and someone on the other side is paid specifically to find reasons to pay less.

Building a model with this in mind means a few things:

Cohort data should be tracked from day one. Not just aggregate revenue, but revenue by customer cohort, with enough history to identify spending trends. The pattern I used to save the Syft exit was only available because we had good cohort-level data going back years. **Variance should be explainable, not just reported.** Every material miss against forecast should have a documented explanation: a market factor, a timing shift, an assumption that turned out to be wrong and why. Unexplained variance in M&A diligence reads as a management problem.

The model should have flex. Not in the sense of being manipulable, but in the sense of being reframeable from multiple angles. Multiple scenarios. Sensitivity analysis pre-built. Alternative valuation frameworks ready to be applied to the same underlying data.

The Part Nobody Writes About

The CEO liked me before this moment. After it, something changed. He has told the story of that Saturday morning in Holborn to other people. Not because of the technical reframe. Because when the problem arrived, there was someone whose job it was to find the answer, who went to find it.

That is what the finance function is supposed to be in a startup. Not the person who reports the numbers. The person who knows what the numbers mean and can defend what they mean when the stakes are highest. The company sold. People who had worked very hard for years became wealthy. That is what the model was for.

It does not take a genius to do the work. It takes someone willing to do the work before they are under pressure.

A variance between forecast and actuals is not automatically a crisis. In M&A diligence, it is a negotiating position the acquirer takes by default, and the only way to counter it is with a better analytical

framework applied to real data. In the case of the Syft exit, that framework was growth-adjusted cohort value: the observation that a customer base with increasing spend per cohort over time is worth more than a static LTV calculation suggests. The fix worked because the data was real, the logic was credible, and the work was done before the conversation closed. Build your model to survive the questions that come after the LOI, not just the ones that come before it.

Appendix C: Templates

The following templates are starting points, not final documents. Every company is different. These are the four most commonly requested formats from founders preparing to raise: the investor update email, a financial model structure guide, a fundraising CRM framework, and a data room checklist. Adapt them to your context.

Template 1: Monthly Investor Update Email

Send this once per month to all investors, advisors, and active angels. Short is better than comprehensive. Consistency is more important than length.

Subject line: [Company Name] Update -- [Month Year]

Opening line: one sentence on how the month went overall. No vagueness. ”Strong month on acquisition, slower on expansion” is better than “making good progress”.

Metrics section (include every month, same format):

MRR: [current] (vs [last month], vs [plan])

Customers: [number] ([new], [churned])

Gross margin: [%]

Burn: [monthly], runway: [months]

NRR (if applicable): [%]

Highlights (two to three bullet points): what went well and why.

Lowlights (one to two bullet points): what did not go to plan and what you are doing about it. This section is the most important. Investors who never hear about problems assume you are hiding them.

Asks (optional but high value): specific requests from your network. Introductions to a named target customer, a candidate for a specific

role, or a question you are trying to answer. Investors who cannot help with money can often help with this.

Next month focus: one to two sentences on the priority area for the next 30 days.

Template 2: Financial Model Structure Guide

This is a guide to what each tab of an investor-facing startup financial model should contain. The model should be navigable in under five minutes by someone seeing it for the first time.

Tab 1 -- Summary (start here): high-level outputs. Revenue, gross margin, headcount, cash, burn, runway across the forecast period. No inputs on this tab; every number is a formula pulling from other sheets.

Tab 2 -- Assumptions: every key input in one place. Time horizon, pricing assumptions, growth rates, hiring timeline, COGS drivers, sales efficiency metrics. Every cell is an input. Everything on every other tab references this tab.

Tab 3 -- Revenue: driver-based revenue calculation. Customers times ACV, or users times conversion times ARPU, or GMV times take rate. The mechanics of the revenue model live here, pulling from Assumptions.

Tab 4 -- Headcount: role-by-role hiring plan with loaded costs (salary plus employer contributions plus benefits). Tie the hire dates to the Assumptions tab so scenarios can be adjusted centrally.

Tab 5 -- P&L: revenue minus COGS equals gross profit, minus OpEx equals EBITDA. Each line should trace to either the Revenue tab, Headcount tab, or Assumptions tab. No hardcoded numbers.

Tab 6 -- Cash Flow: operating cash flow, investing activities, financing activities. The cash balance at the bottom of this tab should drive the runway calculation on the Summary tab.

Template 3: Fundraising CRM Tracker

Track every investor in a single spreadsheet throughout the process. The minimum columns are:

Investor name. Firm name. Fund size (approximate). Stage focus. Check size range. Current stage (not contacted / warm intro / deck sent / first meeting / second meeting / partner meeting / due diligence / term sheet / closed / passed). Date of last contact. Next action. Owner (which co-founder manages this relationship). Notes from conversations.

Review this tracker weekly with your co-founder. Move anyone who has been silent for two weeks to a re-engagement action. The investors who pass rarely volunteer feedback; you have to ask specifically and accept that most will still not give it.

Keep a separate tab for warm introduction paths. For every investor you want to reach, identify two to three people in your network who might know them. The quality of your warm intro is one of the largest determinants of whether you get a first meeting.

Template 4: Data Room Checklist

Have this ready before your first partner meeting. Do not wait to be asked for it. A complete data room at the right time signals operational readiness and accelerates due diligence.

Corporate and legal:

Certificate of incorporation. Articles of association or equivalent. Shareholders agreement. Current cap table (fully diluted). Any existing term sheets or convertible notes. Existing investor agreements.

Financial:

Last 12 months of management accounts (P&L and balance sheet). Current financial model. Last three months of bank statements. Any existing debt facilities or credit agreements.

Commercial:

Three to five anonymised customer contracts (or template contract). Pipeline summary. Customer churn data. NPS or customer satisfaction data if available.

Product and technology:

Product roadmap. Architecture overview (one page). Any IP registrations, patents, or pending applications. Key technology dependencies and vendor agreements.

Team:

Organisation chart. Employment contracts for all full-time employees. Advisor agreements. Option pool grants and schedule. Any equity agreements with contractors.

Market:

Competitive landscape summary. Market sizing methodology. Any customer research, surveys, or interview notes.

One rule about the data room: label every document clearly. A folder called ”Legal Docs” containing 40 unlabelled PDFs is not a data room. It is an obstacle course. An investor who cannot find what they need will assume it does not exist.

Appendix A: Chapter Exercises

Each exercise is designed to build one component of your financial model or fundraising toolkit. Completed in sequence, they produce a complete investor-ready financial model and fundraising preparation package.

Fundraising Pillar (Articles 001-020)

Exercise 001 - You Have 93 Seconds

Pull up your current pitch deck and set a timer for 93 seconds. Flip through it at the pace an investor would skim, one to two seconds per slide. Write down the three things that were legible in that time. Then write down the three things that were not. Restructure the first three slides so that team, market, and problem are all immediately visible without reading a single word of body copy. Ask a founder who has never seen the deck to do the same test and tell you what they understood.

Exercise 002 - Stop Polishing Your Deck

Open your pitch deck and go through each slide. For every design element (colour, animation, icon, illustration), ask one question: does this make the idea clearer, or does it make the slide look better? Remove everything in the second category. Time how long the redesign took. That time is your opportunity cost benchmark for future deck iterations.

Exercise 003 - Running Out of Cash During a Fundraise

Calculate your real fundraising timeline runway. Take your current cash balance and divide it by your honest monthly cash burn (from your cash flow statement, not your P&L). Then subtract 20 weeks from today. If the resulting date is less than four months out, you are raising too close to the edge. Write down the three decisions you would make this month if you knew the fundraise would take six months to close instead of three.

Exercise 004 - Raise When You Don't Need To

Map your fundraising trigger points. For your current business, identify: the metric level at which you will start the next raise, the amount of runway you need to have at that point, and the date by which you need to reach that metric to hit the runway requirement. Work backward from those numbers to build a milestone plan that creates the conditions for raising from strength rather than necessity.

Exercise 005 - Venture Debt

Model the venture debt scenario for your business. Assuming you have recently closed an equity round, calculate what 20 to 30% of that round amount would look like as venture debt at a 9 to 12% interest rate over a 36-month term. Build the monthly principal and interest repayment schedule in your financial model. Identify the specific use of proceeds (hiring, inventory, marketing) and the revenue milestone the venture debt is designed to reach. Then model what happens to your runway and equity dilution with and without the debt.

Exercise 006 - Convertible Notes

List all outstanding convertible instruments on your cap table: SAFEs, convertible notes, any informal agreements. For each one, document: the principal amount, the valuation cap, the discount rate, and the maturity date if applicable. Then model what each instrument converts to at three different priced round valuations: your floor, your target, and your stretch. Identify which instruments create the most dilution and at which valuation that dilution is most concentrated.

Exercise 007 - Picking the Right VC

Build a VC targeting matrix with five columns: fund name, stage focus, sector focus, typical check size, and a portfolio founder you can reference. Populate it with 20 funds relevant to your stage and sector. Score each one out of five on: stage fit, sector depth, and network value to your specific business. Remove any fund where you score below three on stage fit. The remaining list is your Tier 1 and Tier 2 outreach targets.

Exercise 008 - What Investors Look For in Your Model

Open your financial model and locate the following five elements. If any are missing, this is where you build them this week. First, a revenue model built from drivers rather than a growth percentage. Second, an assumptions tab where every key input lives in a single cell with a note explaining its source. Third, a gross margin calculation that separates COGS from OpEx correctly. Fourth, monthly cash flow with runway clearly stated. Fifth, a conservative scenario that shows the company still viable if the top two assumptions are wrong.

Exercise 009 - The 30 Diligence Questions

Print the 30 diligence questions from this article and answer each one in writing, in two to three sentences, without opening your financial model. Mark any question where your answer was vague or where you had to say "I would need to check." Those marked questions are your model gaps. For each gap, identify which tab in the model is missing or incomplete and build it this week.

Exercise 010 - Seed vs. Series A Model

Place your current financial model on a spectrum between the two columns in this article: seed model (hypothesis, benchmark-based, directional) and Series A model (data-driven, actuals-supported, precise). For each element where you are still in seed-model territory, write down what data you would need to collect over the next 12 months to make it Series A ready. This becomes your data collection roadmap.

Exercise 011 - Competitive Fundraising Process

Build your investor outreach tracker. Create a spreadsheet with the following columns: investor name, fund, stage fit score (1-5), sector fit score (1-5), warm intro available (yes/no), intro source, outreach date, first meeting date, follow-up status, and current stage. Populate Tier 1 (10 funds), Tier 2 (15 funds), and Tier 3 (10 funds). Identify the three people who can provide the highest-quality introductions to your Tier 1 list and draft the specific ask you will send each of them.

Exercise 012 - Term Sheet Red Flags

Take any term sheet you have received or can access (many are published publicly by funds like Index and Balderton). Go through each clause and identify: the liquidation preference structure and what it

means at a 1x, 2x, and 5x exit. The anti-dilution provision type and what it means in a down round scenario. The board composition and what control it gives investors. The pro-rata rights and what they mean for future rounds. For each clause, calculate the founder outcome at three exit values: 1x, 3x, and 10x on the invested capital.

Exercise 013 - Bridge Rounds

Model the bridge round scenario for your business. Start with your current cash runway. Add a hypothetical bridge of 20% of your last round at a 20% discount to the last round valuation. Calculate: the new runway, the dilution to existing shareholders, and the milestone you are committing to reach with the bridge capital. Then model the alternative: a significant cost reduction that extends runway without a bridge. Compare the two paths on founder dilution, runway extension, and milestone achievability.

Exercise 014 - Valuation Negotiation

Build your valuation defence document. Calculate your implied valuation using three methods: revenue multiple (take your forward ARR and multiply by the median SaaS multiple for your stage from recent comparable rounds), return math (work backward from what a 10x return on a typical fund check size implies about ownership percentage and therefore valuation), and comparable transactions (identify three publicly reported rounds in your sector and stage from the last 18 months). Present all three numbers and identify where your target valuation sits relative to each.

Exercise 015 - When Investors Challenge Your Numbers

Identify the three assumptions in your model that are most likely to be challenged. For each one, write a three-part response: the assumption and its exact value, the specific source or evidence behind it (not "industry research," but the specific data point), and the range of plausible values the assumption could take and the model output at each end of that range. Deliver this as a verbal rehearsal with a fellow founder who is asked to push back as hard as possible.

Exercise 016 - Data Room

Create your data room folder structure today using the seven-section template from this article: corporate and legal, capitalization, financial, metrics and KPIs, team and culture, product and technology, market and competitive. For each section, list every document that should be there and mark each one as: exists and is current, exists but needs updating, or does not exist yet. Build the "does not exist yet" list into your preparation schedule for the next fundraise.

Exercise 017 - Fundraising Timeline

Build your fundraising Gantt chart. Start from your target close date and work backward. Add: 30 days for documentation and closing after term sheet. 30 days for due diligence before term sheet. 60 days for partner meetings and second meetings. 30 days for warm-up conversations and first meetings. 30 days for materials preparation and outreach. Mark the resulting start date on your calendar. If it is sooner than you thought, adjust your runway planning accordingly.

Exercise 018 - Post-Raise First 90 Days

Build your 90-day post-close operating plan. Divide it into three 30-day sprints. Sprint one covers: setting up board reporting, calibrating the financial model with actuals, completing key hires in pipeline, establishing the monthly finance cadence. Sprint two covers: first full month of actuals versus plan, first board report, key product or commercial milestones from the round plan. Sprint three covers: first quarterly review, pipeline health check against the model, identification of any early plan variances and their implications. Assign a named owner and a completion date to every item.

Exercise 019 - Investor Updates

Write the first version of your investor update using the five-section format: headline metric, what went well, what did not go well, metrics table, and a specific ask. Use your real numbers from last month. Include one thing that did not go well. Send it to at least three investors or advisors who have expressed interest in your company, even if you are not yet fundraising. Collect feedback on what they found most useful.

Exercise 020 - The Cap Table

Build your current fully diluted cap table: founding team shares, any angel or pre-seed investors, outstanding SAFEs or convertible notes modelled as if they convert at your target seed valuation, and your current or planned option pool. Calculate the founder ownership percentage on both a basic and fully diluted basis. Then model what that percentage becomes if you raise your next round at three different valuations: your floor, your target, and your stretch. Save this as a tab in your financial model.

Financial Modeling Pillar (Articles 021-040)

Exercise 021 - How to Build a Financial Model

Build the skeleton of your financial model before adding any numbers. Create the following tabs in order: Assumptions, Revenue Model, Headcount Plan, COGS, OpEx, Cash Flow Statement, P&L Summary, Balance Sheet, and Scenarios. Label every tab. On the Assumptions tab, list the ten inputs you believe will most influence your revenue forecast. Do not fill in any numbers yet. The goal of this exercise is to build the architecture before the content.

Exercise 022 - Driver-Based Modeling

Replace the revenue line in your current model with a driver-based build. For a SaaS business: start from monthly new leads, apply a conversion rate to qualified opportunities, apply a close rate to contracts, multiply by average contract value, then apply a monthly churn rate to the existing base. For a marketplace: start from supply-side acquisition, model utilisation rate, multiply by average transaction value, apply your take rate. Every number in the revenue output must trace back to a named driver in the assumptions tab with a source note.

Exercise 023 - Cohort Revenue Forecasting

Build a cohort revenue model alongside your current revenue forecast. Take your last six months of customer data and organize customers by the month they first paid. For each cohort, track the revenue generated in each subsequent month. Calculate the average retention rate across cohorts at month 1, month 3, month 6, and month

12. Use these rates to project the forward revenue from your existing customer base. Compare the cohort-based forward revenue to your current revenue forecast. Identify where they diverge and what assumption is causing the divergence.

Exercise 024 - The Assumptions Tab

Audit your current financial model for hard-coded values. Open every tab and highlight every cell that contains a number rather than a formula. For each highlighted cell, decide: should this be in the assumptions tab as a named input, or is it genuinely a fixed constant that will never change? Move everything that should be a named input to the assumptions tab. Add a "Source" column and a "Last Updated" column. Fill in the source for every assumption. Any assumption without a source is either a placeholder or a guess, and both need to be addressed before the model goes to an investor.

Exercise 025 - Headcount Plan

Build your headcount plan role by role for the next 24 months. For each planned hire: role title, department, planned start month, gross annual salary, employer overhead percentage (look up the actual rate for the country you are hiring in), and a one-sentence business justification. Add a productivity ramp: month one at 25% productive output, month two at 50%, month three at 75%, month four at 100%. Connect this plan to your revenue model so that revenue-generating roles have a direct impact on the revenue forecast.

Exercise 026 - COGS vs. OpEx

Go through your current cost model and apply the delivery test to every line item: is this cost incurred to deliver the product or service to a paying customer? If yes, it belongs in COGS. If no, it belongs in OpEx. Reclassify any misallocated costs. Then recalculate your gross margin with the corrected classification. If your gross margin changes by more than two percentage points, the original model had a material classification error. Document the new classification rationale in a note on the COGS tab.

Exercise 027 - Cash Flow Statement

Build a direct cash flow statement for the last three months of your business using actual bank statements and your accounting system. Organize cash flows into three sections: operating activities (cash received from customers minus cash paid to suppliers and employees), investing activities (equipment, software purchases), and financing activities (capital raised, debt drawn or repaid). Reconcile the ending cash balance to your bank balance for each month. If they do not match, find the reconciling items. This is the most important financial literacy exercise in the book.

Exercise 028 - Three Scenarios

Build your three scenarios by varying the two most uncertain assumptions in your model, not by applying a percentage multiplier to the base case revenue line. For most early-stage companies, the most uncertain assumptions are CAC and time-to-revenue. Run: base case (assumptions as modelled), conservative case (CAC 30% higher, first revenue delayed two months), and aggressive case (CAC 20% lower, product-market fit accelerates in month ten). For each scenario, record runway at end of period, ARR at end of period, and headcount required. Present these side by side on a summary tab.

Exercise 029 - Two-Sided Marketplace

If you are building a marketplace, build the supply and demand models as separate modules before connecting them. Supply module: how many providers are acquired per month, at what cost, with what activation rate and what utilisation rate once active. Demand module: how many customers are acquired per month, at what cost, with what booking frequency and at what average transaction value. Connect the modules through a matching efficiency rate: what percentage of demand can be fulfilled given the current supply base. Model what happens when supply or demand is out of balance.

Exercise 030 - SaaS vs. Marketplace

Identify which financial model type applies to your business or, if your business has characteristics of both, which elements from each model apply. Write down the three financial metrics that are most important for your specific model type and the three that do not apply.

Then check whether your current model is built around those core metrics. If it is built around the wrong metrics for your business model, the model is telling the wrong story to investors.

Exercise 031 - Multi-Currency Model

If you operate in more than one currency, add a currency module to your financial model. For each operating entity, specify the functional currency. Add a single FX rate assumption table to the assumptions tab: one rate per currency pair, sourced from a current market rate and updated monthly. Connect all revenue and cost lines to the relevant functional currency, then translate to the presentation currency using the rate in the assumptions tab. Test the model by changing one FX rate by 10% and verifying the P&L impact flows correctly.

Exercise 032 - Gross Margin by Business Model

Find the gross margin benchmark for your specific business model from the reference table in this article. Then calculate your current gross margin in three ways: using your own COGS definition, using the most restrictive definition (all customer-related costs in COGS), and using the most generous definition (only direct delivery costs in COGS). You now have a range. Identify where your model's gross margin sits within that range and write a one-paragraph explanation of why. That paragraph is your gross margin defence for investor conversations.

Exercise 033 - Sensitivity Analysis

Build a sensitivity table in your financial model for the two most important assumptions: CAC and monthly churn. Create a 5x5 grid where CAC varies from 20% below your base case to 20% above it in 10% increments, and churn varies from 0.5% to 2.5% per month in 0.5% increments. The output in each cell should be 18-month runway. Identify the combinations that make the business unviable (runway below nine months) and the combinations that create the most attractive outcomes. This grid is your risk map.

Exercise 034 - Model Audit

Run the five-check audit on your current financial model before your next investor meeting. Check one: does the cash balance in the cash flow statement match the cash balance on the balance sheet in every

period? Check two: does every key assumption reference a single cell in the assumptions tab? Check three: does gross margin fall within the benchmark range for your business model? Check four: is there a conservative scenario? Check five: are there any formula errors (REF, VALUE, N/A) on any tab? Fix everything that fails before sending the model to anyone.

Exercise 035 - Excel vs. Google Sheets

If your model is currently in Google Sheets, spend 30 minutes exporting it to Excel and checking that all formulas translated correctly. Pay particular attention to: array formulas, any Google-specific functions (IMPORTRANGE, QUERY, GOOGLEFINANCE), and charts. If your model is in Excel, check that the cloud-sharing link works and that collaborators can access it without needing a local Excel installation. The goal is a model that can be accessed and reviewed by any investor on any platform without technical friction.

Exercise 036 - The Balance Sheet

Add a balance sheet tab to your financial model. Build the equity section first by pulling paid-in capital from your funding history and accumulated deficit from the cumulative net income total in your P&L. Then build the liabilities section, making sure deferred revenue is linked to your revenue recognition timing. Add current assets, ensuring cash ties exactly to your cash flow statement. Build a balance check cell. If the model does not balance, work backward through each section to find where the disconnect is.

Exercise 037 - Pre-Revenue Model

Build the three-section pre-revenue model for your business. Start with the cost model: list every hire, every tool, every fixed cost by month for the next 24 months. Then build the revenue model with an explicit assumptions tab where every input has a note explaining its basis. Finally, build the runway calculation under three scenarios. For each scenario, identify the single most important assumption that determines whether the company reaches the next fundraising milestone.

Exercise 038 - Red Flags Audit

Send your model to a founder or advisor who will ask hard questions. Ask them to probe the four most important assumptions: revenue ramp, gross margin, CAC, and headcount plan. For each assumption they challenge, write a two-sentence note in the assumptions tab explaining the basis and the range of comparable benchmarks. Repeat until you can answer every material assumption question without referring to the screen.

Exercise 039 - Mid-Fundraise Updates

Take your current model and identify every place where a key assumption appears in more than one tab. Consolidate each assumption into the assumptions tab and replace the other instances with cell references. Save the current version with the naming convention: CompanyName_Model_v[number]_[YYYYMMDD].xlsx. Create a versions folder and move all previous files into it. Write a one-paragraph change summary as if sending it to an investor. This surfaces whether your model is set up for safe iteration.

Exercise 040 - Board Report

Take the financial model from your most recent fundraise and add the actuals layer for the last three months. For each month, create four columns: plan, actual, variance absolute, variance percentage. Write a one to two sentence variance note for every line more than 10% off plan. Then build the five-section board report template (headline metrics, P&L summary, cash and runway, variance analysis, forward forecast) populated with the most recent month's data. Share with your existing investors or board members and ask what was most and least useful.

Unit Economics Pillar (Articles 041-051)

Exercise 041 - CAC, LTV, and Payback Period

Calculate all three metrics for your business using the correct definitions from this article. CAC: total sales and marketing spend in the period divided by new customers acquired, not just ad spend. LTV: average gross profit per customer per month divided by monthly churn rate, not average revenue. Payback period: CAC divided by average monthly gross profit per customer, not average monthly revenue per

customer. If any of your previous calculations used a different definition, recalculate and note by how much the number changed.

Exercise 042 - Blended vs. Channel CAC

Break your customer acquisition spend into at least three channels: paid (separated by platform), organic (SEO and content), and outbound (sales development, if applicable). For each channel, calculate the CAC separately using the spend and customers attributed to that channel over the last 90 days. Compare each channel CAC to your blended CAC. Identify which channels are subsidising which, and model what your blended CAC would look like if you scaled the most efficient channel by 2x at the expense of the least efficient.

Exercise 043 - Net Revenue Retention

Calculate your NRR for the last 12 months using the cohort method. Take the cohort of customers who were active 12 months ago. Calculate their total revenue last month (including any expansion, excluding any churned customers). Divide by their total revenue 12 months ago. The result is your NRR. If you cannot calculate this because you do not have the cohort data, that is the exercise: set up the data collection process that will allow you to calculate NRR in 90 days.

Exercise 044 - Marketplace LTV

If you run a marketplace, calculate your LTV correctly using the gross profit method from this article. Start from gross profit on net revenue (your take rate times GMV), not from GMV directly. Apply your supply-side and demand-side churn rates separately to understand which side of the market degrades faster. Calculate LTV:CAC for both sides independently. Identify which side of the market is your binding constraint for LTV expansion and what one action would most improve it.

Exercise 045 - Unit Economics by Market

If you operate in more than one geographic market, separate your unit economics by market. Calculate CAC, LTV, payback period, and gross margin for each market independently. Identify the market with the best unit economics and the market with the worst. For the worst-performing market, determine whether the poor economics are

structural (market characteristics) or operational (early stage, suboptimal channel mix). This analysis should feed directly into your capital allocation decisions for the next quarter.

Exercise 046 - LTV:CAC Benchmarks

Look up the median LTV:CAC ratio for three companies in your sector at your stage (public SaaS company filings and investor reports from Bessemer, Andreessen, and Sequoia are useful sources). Compare your own LTV:CAC to this benchmark. If you are above the benchmark, identify the specific reason (lower churn, higher expansion revenue, more efficient acquisition). If you are below it, identify the most leveraged change you could make to close the gap. Build the improved LTV:CAC scenario into your financial model.

Exercise 047 - Gross Margin Trajectory

Build a gross margin improvement roadmap for the next 24 months. Identify the three specific operational levers that will improve your gross margin: cost renegotiations with suppliers, automation of manual delivery steps, moving from a variable to a fixed infrastructure cost model as you scale. For each lever, estimate the gross margin impact in percentage points and the timeline. Build this trajectory into the operating model so that the gross margin in year two reflects the expected improvement from these specific actions.

Exercise 048 - Cohort Analysis

Build your first cohort retention chart. Take at least six months of customer data. Group customers by the month they became paying customers. For each cohort, calculate the percentage still active (and the revenue retained) in each subsequent month. Create a heatmap visualization: cohorts as rows, months since acquisition as columns, revenue retention percentage as the cell value. Identify: when is churn highest for most cohorts, whether newer cohorts are retaining better than older ones, and what is the long-run retention floor where the curve flattens.

Exercise 049 - Stage-Appropriate Metrics

Identify which stage your company is at: pre-seed, seed, or Series A. From the metrics framework in this article, select the five metrics that

are most important for your current stage. For each one, write down your current value, the benchmark range from this article, and whether you are above, within, or below the benchmark. The metrics where you are most above benchmark are your fundraising strengths. The metrics where you are most below benchmark are either your biggest weaknesses or your biggest opportunities for the next six months.

Exercise 050 - Network Effects

If your business has network effects, map them explicitly. Identify: the type of network effect (same-side or cross-side), the density threshold at which the effect activates in a specific geography or segment, and the number of supply-side and demand-side participants required to reach that threshold. Build this into your financial model as a variable that affects conversion rates or transaction frequency above the threshold. Model the company in two states: pre-threshold and post-threshold. Show what the unit economics look like in each state and the capital required to reach the threshold in your target market.

Appendix B: Recommended Reading

This list is organized by the topics the book covers. These are books that informed the thinking behind Raise Ready, and books that founders have found useful across fundraising, financial modeling, and building high-growth businesses.

On Venture Capital and Fundraising

Secrets of Sand Hill Road - Scott Kupor (2019)

The clearest explanation of how venture capital funds work, how GPs make decisions, and what founders misunderstand about term sheets. Required reading before negotiating any VC term sheet. **Venture Deals** - Brad Feld and Jason Mendelson (4th edition, 2019) The definitive reference for term sheet mechanics. Every clause, every provision, with plain-language explanations of economic and governance implications. Keep it open during any term sheet negotiation. **The Hard Thing About Hard Things** - Ben Horowitz (2014)

Not a fundraising manual. A frank account of what it is actually like to run a startup through adversity. Useful for calibrating the gap between how fundraising is described and what the company going through it actually looks like. **Founder vs. Investor** - Elizabeth Joy Zalman and Jerry Neumann (2023) Explores the structural tensions in the founder-investor relationship with unusual candor from

both sides. Particularly useful for understanding why investors behave the way they do in difficult situations. **Zero to One** - Peter Thiel with Blake Masters (2014)

Less a how-to guide and more a framework for thinking about defensibility and monopoly. Useful for building the strategic narrative that sits behind a financial model.

On Financial Modeling and Business Finance

Financial Intelligence for Entrepreneurs - Karen Berman and Joe Knight (2008) The most accessible introduction to financial statements for non-finance founders. Explains P&L, balance sheet, and cash flow concepts without MBA prerequisites. **The McKinsey Way** - Ethan Rasiel (1999)

Useful for understanding how analysts at top firms approach financial and strategic analysis. Relevant for the modeling discipline and the habit of working from facts rather than assumptions. **Valuation: Measuring and Managing the Value of Companies** - McKinsey and Company (7th edition, 2020) The academic and practitioner standard for valuation methodology. Chapters on DCF, multiples, and cost of capital are most relevant for founders preparing for Series A and growth rounds. **Financial Modeling** - Simon Benninga (4th edition, 2014)

The textbook used in most finance programs. Dense but comprehensive on model construction, scenario analysis, and the mechanics of three-statement modeling. **Principles of Corporate Finance** - Brealey, Myers, and Allen (13th edition, 2019) The standard finance textbook. More useful as a reference than as reading cover to cover.

The chapters on capital structure, working capital, and risk measurement are most relevant for startup founders.

On SaaS and Recurring Revenue Metrics

Winning by Design - Jacco van der Kooij and Fernando Pizarro (2018) The most rigorous treatment of SaaS revenue architecture and the metrics that drive it. Particularly useful for founders building a recurring revenue model for the first time. **From Impossible to Inevitable** - Aaron Ross and Jason Lemkin (2nd edition, 2019) The standard playbook for B2B SaaS sales scaling. Useful context for the CAC and headcount assumptions that go into a financial model. **The SaaS Playbook** - Rob Walling (2023)

A practical guide to building a SaaS business with a focus on metrics, pricing, and growth mechanics. More actionable than academic. **Subscribed** - Tien Tzuo and Gabe Weisert (2018)

Written by the CEO of Zuora. Explains the economics of the subscription model and the metrics that matter in recurring revenue businesses. Useful for grounding SaaS financial model assumptions.

On Startup Finance and Operations

Mastering the VC Game - Jeffrey Bussgang (2010)

Written by a VC who was also a founder. Provides perspective from inside the partnership meeting that most founders never see. **Traction** - Gabriel Weinberg and Justin Mares (2nd edition, 2022) Not directly about financial modeling, but the 19-channel framework for customer acquisition is useful for building channel-level

CAC models rather than blended averages. **The Lean Startup** - Eric Ries (2011)

The intellectual foundation for build-measure-learn cycles. Relevant for thinking about how to structure pre-revenue model assumptions as hypotheses rather than forecasts. **High Output Management** - Andy Grove (1983, reissued 2015)

The foundational text on operational management metrics. The concepts of output metrics and leverage map directly to the driver-based modeling approach taught throughout this book. **Measure What Matters** - John Doerr (2018)

The OKR framework. Directly applicable to the milestone planning section of pre-revenue models and the post-raise operational model structure.

On Unit Economics and Growth

Growth Hacker Marketing - Ryan Holiday (2014)

A short, practical read on finding growth levers. Useful background for the customer acquisition assumption sections of financial models. **The Cold Start Problem** - Andrew Chen (2021)

The most rigorous treatment of network effects and the cold start dynamic in marketplace and platform businesses. Directly relevant to Articles 029, 030, and 050 on marketplace modeling. **Obviously Awesome** - April Dunford (2019)

On product positioning. Relevant for understanding how positioning decisions feed into pricing assumptions in the financial model. **The Mom Test** - Rob Fitzpatrick (2013)

How to have honest customer conversations. Directly applicable to building revenue assumptions that are grounded in real customer willingness-to-pay rather than optimistic projections.

On Exits and M&A

Investment Banking - Rosenbaum and Pearl (4th edition, 2022) The standard reference for M&A valuation methodology. The comparable company and precedent transaction analysis chapters are directly relevant for founders approaching exit discussions. **Barbarians at the Gate** - Burrough and Helyar (1989)

A narrative account of a leveraged buyout. Less directly applicable but builds intuition for how deals are structured, how leverage works, and how valuation becomes a negotiating battleground. **The Art of M&A** - Stanley Foster Reed and Alexandra Reed Lajoux (5th edition, 2019) Reference-format guide to the mechanics of M&A processes. Most useful during or after a letter of intent, not as a how-to guide for founders who have never sold a company.

About the Author

Appendix D: Glossary

The following definitions cover the terms used most frequently in investor conversations, financial models, and due diligence processes. Where a term has a precise calculation, it is shown. Where usage varies by context, that is noted.

ARR (Annual Recurring Revenue)

The annualised value of a company's recurring subscription revenue. Calculated as MRR multiplied by 12. ARR excludes one-time fees, professional services, and non-recurring revenue. It is the primary top-line metric for SaaS businesses.

Blended CAC

Customer acquisition cost calculated across all acquisition channels combined. Blended CAC tends to be lower than channel-specific CAC because it averages organic and inbound channels (which have near-zero marginal cost) with paid channels. This is why blended CAC is often misleading: it makes paid acquisition look more efficient than it is.

Bridge Round

An additional financing between two priced equity rounds, typically structured as a convertible note or SAFE. Bridge rounds are used when a company needs more runway to reach the milestones required for the next priced round. They are not inherently negative, but they add a layer of complexity to the cap table.

Burn Rate

The rate at which a company spends its cash. Gross burn is total monthly cash outflow. Net burn is total monthly cash outflow minus total monthly cash inflow. Runway is calculated using net burn.

CAC (Customer Acquisition Cost)

The total cost of acquiring one new customer. Calculated as total sales and marketing spend in a period divided by the number of new customers acquired in that period. Payback period is CAC divided by monthly gross profit per customer.

Cap Table (Capitalisation Table)

A document that records who owns what percentage of a company, in what form (common shares, preferred shares, options, warrants, convertible notes), and at what cost basis. Every financing event changes the cap table.

Carry (Carried Interest)

The share of investment profits that a VC fund manager (general partner) receives, typically 20% of gains above the hurdle rate. Carry is the primary financial incentive for VC investors and only materialises on exits.

Cohort Analysis

A method of grouping customers by the period they were acquired and tracking their behaviour over time. Cohort analysis reveals retention curves, expansion patterns, and LTV trajectories that aggregate metrics conceal.

Convertible Note

A form of short-term debt that converts into equity at a future financing round. Key terms include the interest rate, maturity date, discount rate (a percentage off the next

round price), and valuation cap (the maximum price at which the note converts).

COGS (Cost of Goods Sold)

The direct costs required to deliver the product or service, as opposed to operating expenses (sales, marketing, G&A, R&D). Gross profit equals revenue minus COGS. Gross margin equals gross profit divided by revenue. Getting this distinction wrong distorts the entire P&L.

Dilution

The reduction in a shareholder's ownership percentage when new shares are issued. Every financing round, option grant, or warrant exercise dilutes existing shareholders. Fully diluted share count includes all issued shares plus all options, warrants, and convertible instruments.

Due Diligence

The process by which an investor investigates a company before investing. Due diligence typically covers the financial model, legal documents, customer contracts, team backgrounds, product code, and market assumptions. A complete data room accelerates this process.

EBITDA

Earnings before interest, taxes, depreciation, and amortisation. A measure of operating profitability that strips out capital structure and accounting treatment. Early-stage startups rarely have positive EBITDA, but investors use it to model the path to profitability.

EOR (Employer of Record)

A third-party service that employs workers on behalf of a company in a jurisdiction where the company does not

have a legal entity. Relevant for startups hiring internationally before establishing a subsidiary.

Fund of Funds

An investment fund that invests in other funds rather than directly in companies. Fund of funds are LPs in VC funds, and their return expectations influence what return targets VCs need to hit.

GP (General Partner)

The managing partner(s) of a venture capital fund who make investment decisions and manage the portfolio. GPs receive management fees (typically 2% of fund size annually) and carried interest.

Gross Margin

Revenue minus cost of goods sold, expressed as a percentage of revenue. Software businesses typically target 70-85% gross margin at scale. Marketplaces are typically 50-70%. Hardware-heavy businesses 30-50%.

LTV (Lifetime Value)

The total gross profit expected from a customer over their entire relationship with the company. Calculated as average monthly gross profit per customer divided by monthly churn rate. LTV is almost always higher when calculated on a cohort basis (because early cohorts have lower churn as weak accounts have already churned).

LP (Limited Partner)

The investors in a venture capital fund. LPs provide the capital that GPs deploy. LPs include pension funds, endowments, family offices, sovereign wealth funds, and high-net-worth individuals. LP return expectations drive GP fund economics.

Liquidation Preference

A provision in preferred stock that determines the order and amount of payout in a liquidation event (sale or dissolution). A 1x non-participating liquidation preference means preferred shareholders get their investment back first before common shareholders receive anything. Participating preferred gives preferred holders their money back AND a pro-rata share of remaining proceeds, creating a double-dip.

MRR (Monthly Recurring Revenue)

The normalised monthly value of all active subscription contracts. New MRR, expansion MRR, churned MRR, and net new MRR are the four components of MRR movement.

NRR (Net Revenue Retention)

The percentage of revenue retained from existing customers over a period, including expansion, contraction, and churn. NRR above 100% means the existing customer base is growing without any new customers. NRR is one of the most important single metrics for SaaS businesses because it determines whether the business has compounding growth built in.

Option Pool

A block of shares reserved for future employees, advisors, and contractors in the form of stock options. Option pools are typically 10-20% of the fully diluted share count. VCs often require an option pool refresh (increasing the pool) as a condition of investment, which dilutes founders.

Power Law

The statistical distribution that governs venture capital returns, where a very small number of investments

(typically 5-10% of a portfolio) generate the vast majority of returns. Power law dynamics mean VCs must optimise for outlier potential rather than average expected return.

Post-Money Valuation

The company's valuation immediately after a financing round. Calculated as pre-money valuation plus the amount raised. If a company raises at a $10M pre-money valuation with a $2M investment, the post-money valuation is $12M.

Pre-Money Valuation

The company's valuation before a financing round. This is the valuation at which the new investor's shares are priced. Pre-money valuation negotiations determine how much ownership the investor receives.

Pro-Rata Rights

The right of an existing investor to maintain their ownership percentage by participating in future financing rounds. Pro-rata rights are typically included in term sheets for lead investors. Super pro-rata rights allow an investor to take more than their proportional share.

Runway

The number of months a company can operate at its current net burn rate before running out of cash. Calculated as cash balance divided by monthly net burn. The fundraising process typically takes four to six months, so raising at six months of runway is dangerously late.

SAFE (Simple Agreement for Future Equity)

A financing instrument developed by Y Combinator as an alternative to convertible notes. SAFEs are not debt, have no interest rate, and have no maturity date. Like

convertible notes, they convert into equity at the next priced round with a discount and/or a valuation cap.

Term Sheet

A non-binding document that outlines the key terms of a proposed investment. Term sheets include valuation, investment amount, liquidation preferences, board composition, option pool requirements, and protective provisions. Once signed, the parties proceed to legal documentation and due diligence.

Unit Economics

The revenue and cost structure of a single unit of business: typically one customer, one transaction, or one order. Strong unit economics means the company makes money on each individual transaction after accounting for the direct costs of servicing it.

Vesting Schedule

The timeline over which equity is earned. Standard founder vesting is four years with a one-year cliff, meaning no equity vests for the first year and then 25% vests at the one-year anniversary, with monthly vesting thereafter. Vesting protects co-founders from each other and investors from founding teams that leave early.

Waterfall